

Overview



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Introduction

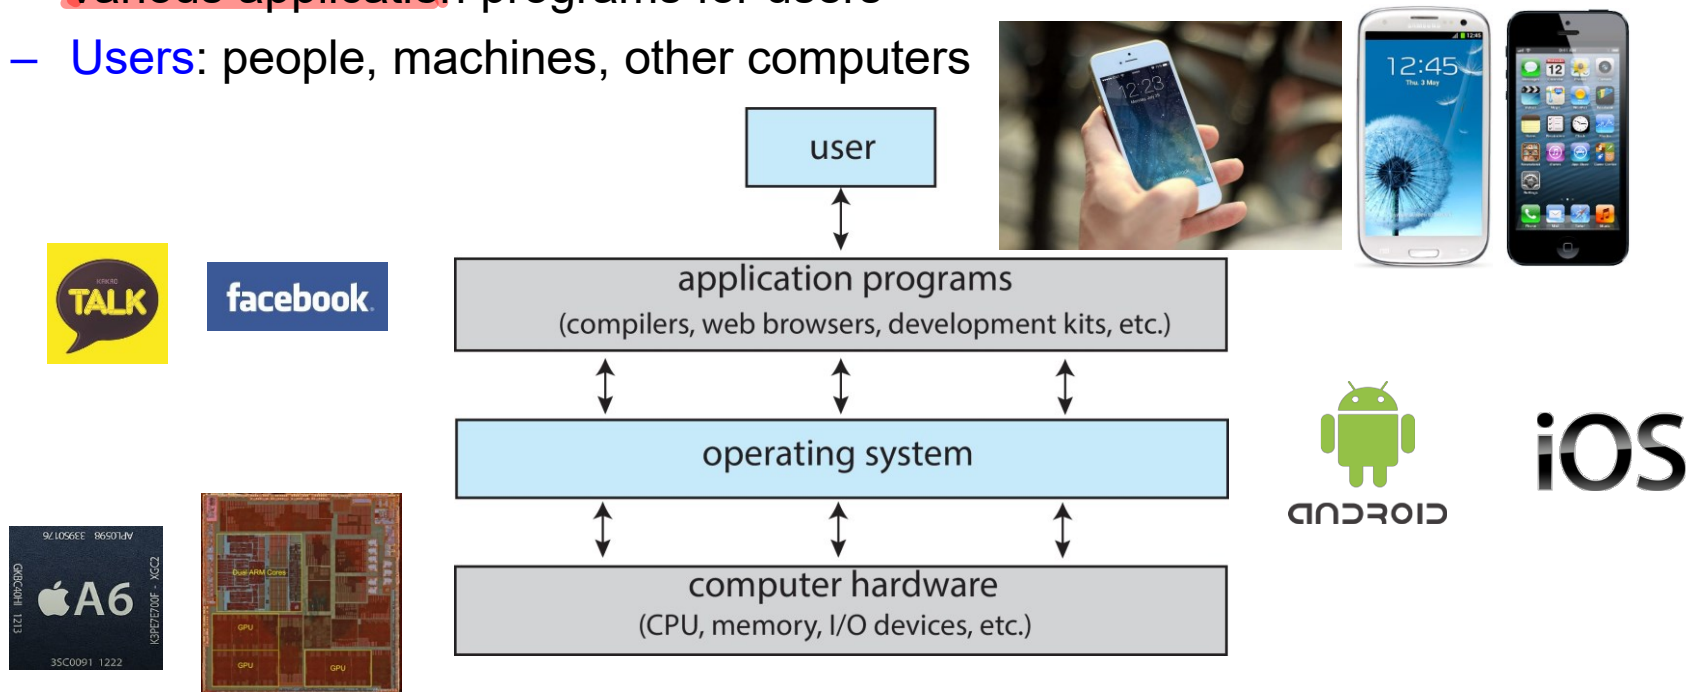
- **Operating system (OS)** is **software** that manages a computer's **hardware**
 - A program that acts as **an intermediary between the computer user and the computer hardware**
- **Operating systems are everywhere,**
 - from cars and home appliances that includes Internet of Things (IoT) devices to smart phones, personal/enterprise computers, and cloud computing
- **The goals of operating systems**
 - Execute user programs and make solving user problems easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner



Computer System

APP
OS
HW

- A computer system can be divided into four components: hardware, operating system, application programs, a user
 - **Hardware**: central processing unit (CPU), memory, input/output (I/O) devices
 - **Application programs**: define the ways where these resources are used to solve users' computing problems (e.g. word processors, web browsers, etc)
 - **Operating systems**: control the hardware and coordinate its use among various application programs for users
 - **Users**: people, machines, other computers



What Operating Systems Do

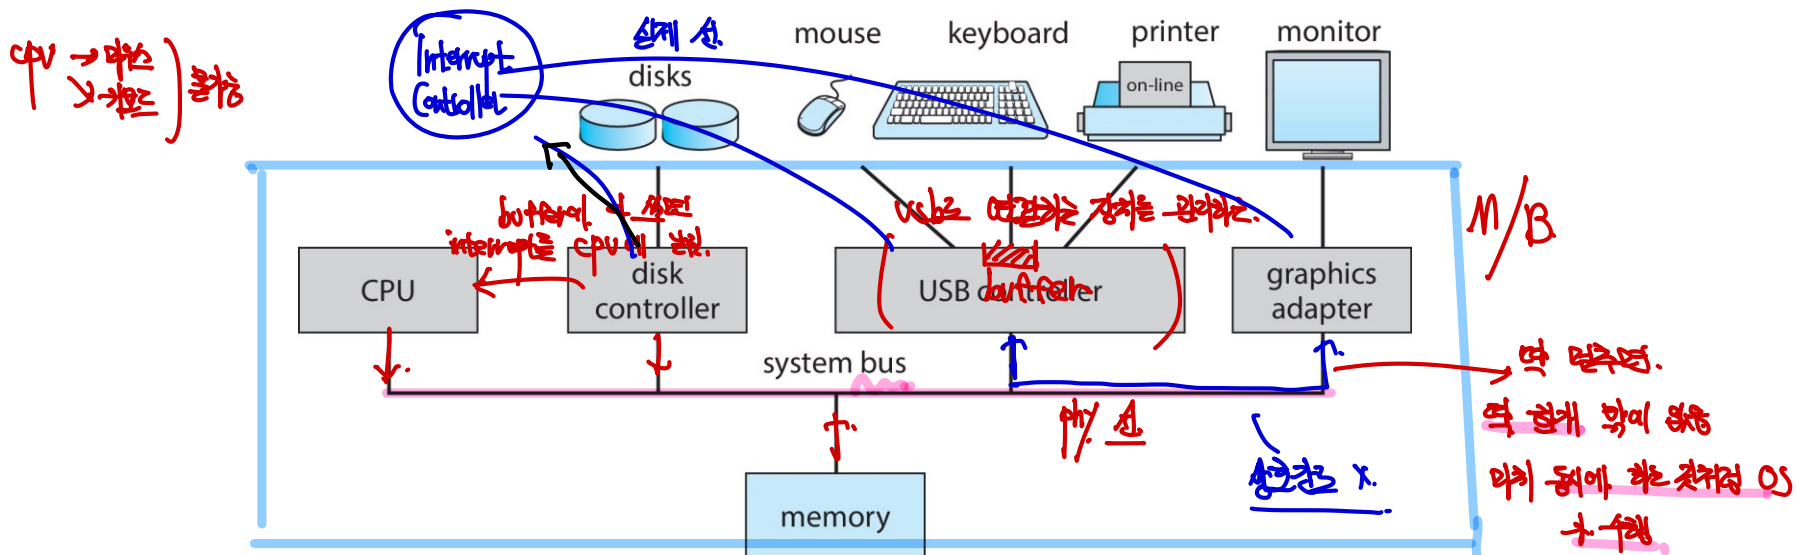
- **An operating system is similar to a government**
 - Like a government, it performs no useful function by itself
 - It provides an environment within which other programs can do useful work
- **In user view, the OS is designed for ease of use without considering resource utilization**
 - (– The user interface for mobile devices (e.g., smartphones) features touch screens, voice recognitions)
 - Some computers have little or no user view, such as embedded computers, and their OS is designed to run without user intervention
- **In system view, the OS can be considered as a resource allocator and a control program**
 - The OS acts as the manager of computing system resources and a control program that manages the execution of user programs to prevent errors

Defining Operating Systems

- **The OS covers many roles and functions and, in general, we have no completely adequate definition of an OS**
 - “Everything a vendor ships when you order an OS” can be a rough definition but this varies greatly across systems (less than a MB ~ GBs)
- **A more common definition is that “the one program running all time on the computer”, which is called kernel, part of OS**
- **Along with the kernel, there are two other types of program:**
 - 1) **System programs** that ships with OS but not part of the kernel
 - 2) **Application programs** that include all programs not associated with OS
- **Today’s OS for general purpose and mobile computing also include middleware**
 - A set of software frameworks that provide additional services to application developers, such as databases, multimedia, and graphics

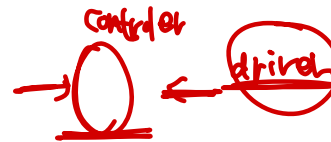
Computer System Organization

- A computer includes CPUs and device controllers connected through **common bus** providing access to **shared memory**
 - Each controller is in charge of a specific type of device and maintains some local **buffer** and special-purpose **registers** → CPU의 레지스터와는 다른. / CPU나 command를 인식하기 위한 것 = CPU와 소통하기 위한 것
 - The controller is responsible for moving data between devices and its buffer
- OSes have a **device driver** for each device controller/ which provides the rest of OS with uniform interface to device
 - CPU and controllers can execute in parallel, competing for memory cycles
 - device 컨트롤러를 위한 인터페이스 → 레지스터 사용 내기 경쟁



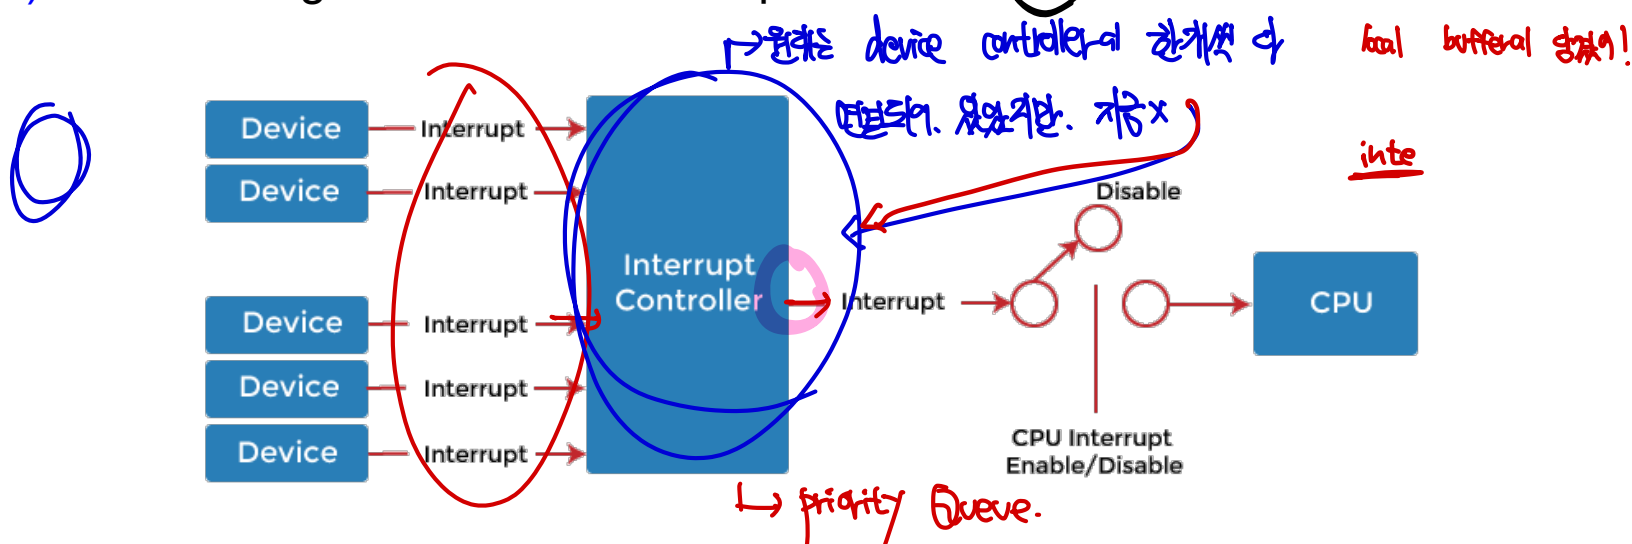
Interrupts

→ 원래는 polling 방식 ⇒ Interrupts



- The **device controller** informs the CPU that it has finished its operation by causing an **interrupt**.

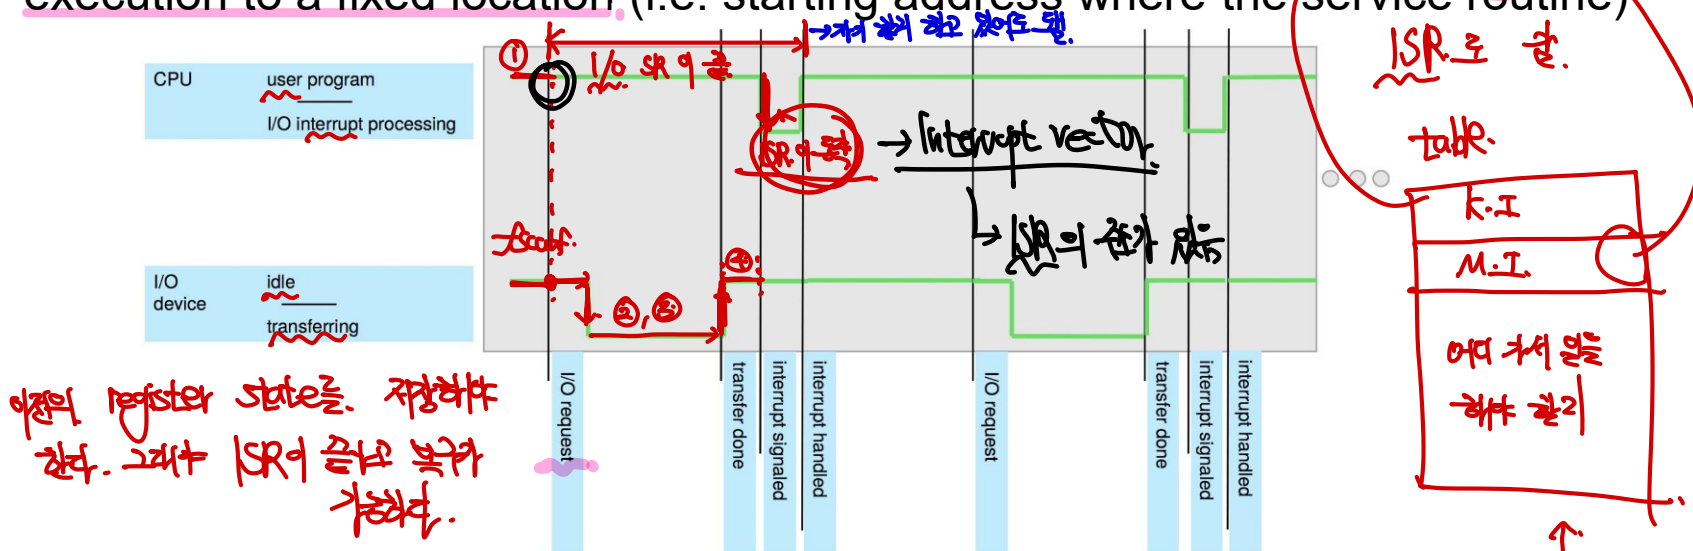
- 1) To start an I/O operation, the **device driver** loads the appropriate **registers** in the **device controller**의 레지스터를 로드.
- 2) The device **controller**, in turn, examines the contents of these **registers** to determine what action to take. → 해야 할지 결정
- 3) The controller starts **transfer of data** between the device and its local buffer
- 4) Once the transfer is complete, the device controller issues an **interrupts** (available)
- 5) The driver gives control to other parts of the **OS**



Interrupt Timeline

- Interrupts are a key part of how the OS and hardware interact, and the OS works in interrupt driven way

- When CPU is interrupted, it stops what it is doing and immediately transfers execution to a fixed location (i.e. starting address where the service routine)



- The interrupt must transfer control to an appropriate routine

- Interrupts must be handled quickly, as they occur very frequently
- The interrupt routine is called indirectly through a table called interrupt vector, which contains the addresses of all the interrupt service routine
- Interrupt architecture must save a state information so it can restore

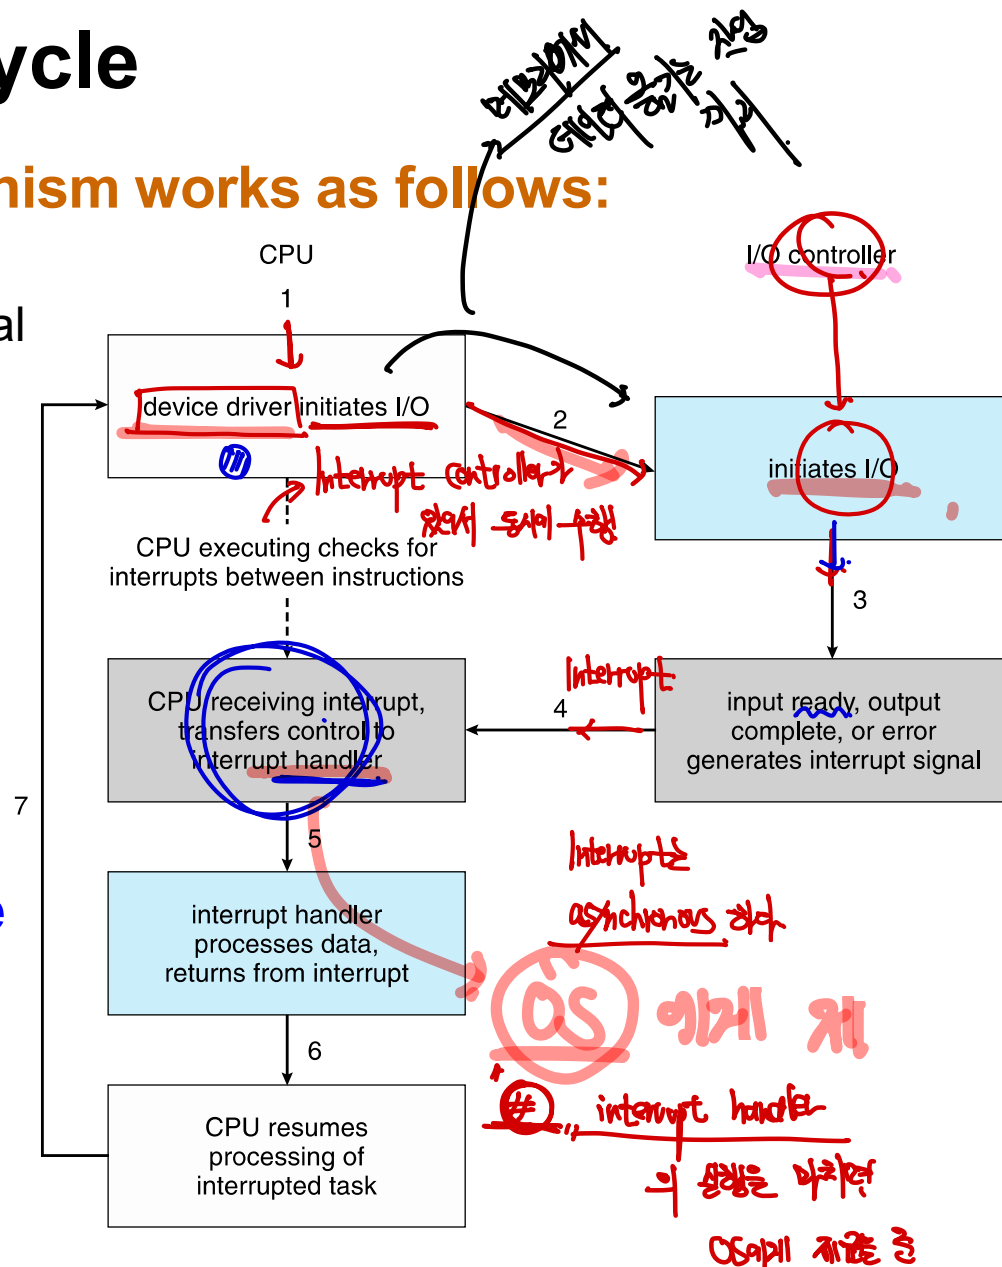
Interrupt Driven I/O Cycle

■ The basic interrupt mechanism works as follows:

- The device controller raises an interrupt by asserting a signal on interrupt request line (wire)
- The CPU catches the interrupt and dispatches it to the interrupt handler
- The handler clears the interrupt by servicing the device

■ Modern computer includes interrupt controller hardware to deal with sophisticated interrupt handling features

- e.g.) (interrupt prioritization)



device driver → OS의 명령을 하드웨어에 정확하게 전달

Storage Structure

■ The CPU can load instruction only from memory

- Main memory is common implemented in **dynamic random access memory (DRAM)**, which is a **volatile** memory that loses its contents when power off

휘발성

■ Computers use other forms of memory

→ 운영체제 부팅에 사용, 프로그램

(디스크에 있던 프로그램을 메모리로 로드)

- **Bootstrap program**, which is loaded at computer power-up/reboot, is typically stored in **electrically erasable programmable read only memory (EEPROM)**
- Generally known as **firmware** and initialize the system and load the OS kernel
- Also known as **basic input output system (BIOS)** in computers

여기에 저장됨

or

UEFI

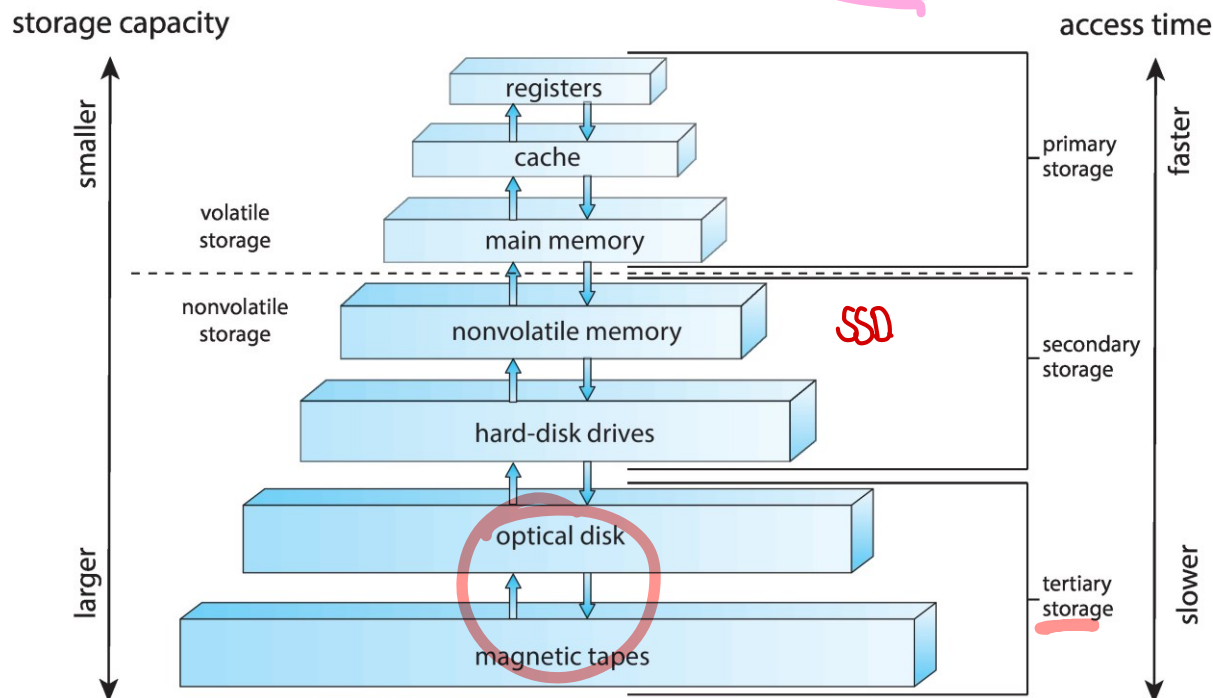
프로그램 부팅 (시스템 장치는 부팅)

■ Ideally, we want the programs and data to reside in main memory permanently, but not possible

- Main memory is usually **too small** to store all needed programs and data permanently
- Main memory is **volatile**

Secondary Storage

- **Most computer systems provides secondary storage as an extension of main memory**
 - Hard-disk drives (HDDs) and nonvolatile memory (NVM) devices are most common as the secondary storage
- **Tertiary storage is slow enough and large enough (e.g. blu-ray)**
 - They are used only for special purposes (e.g. backup)

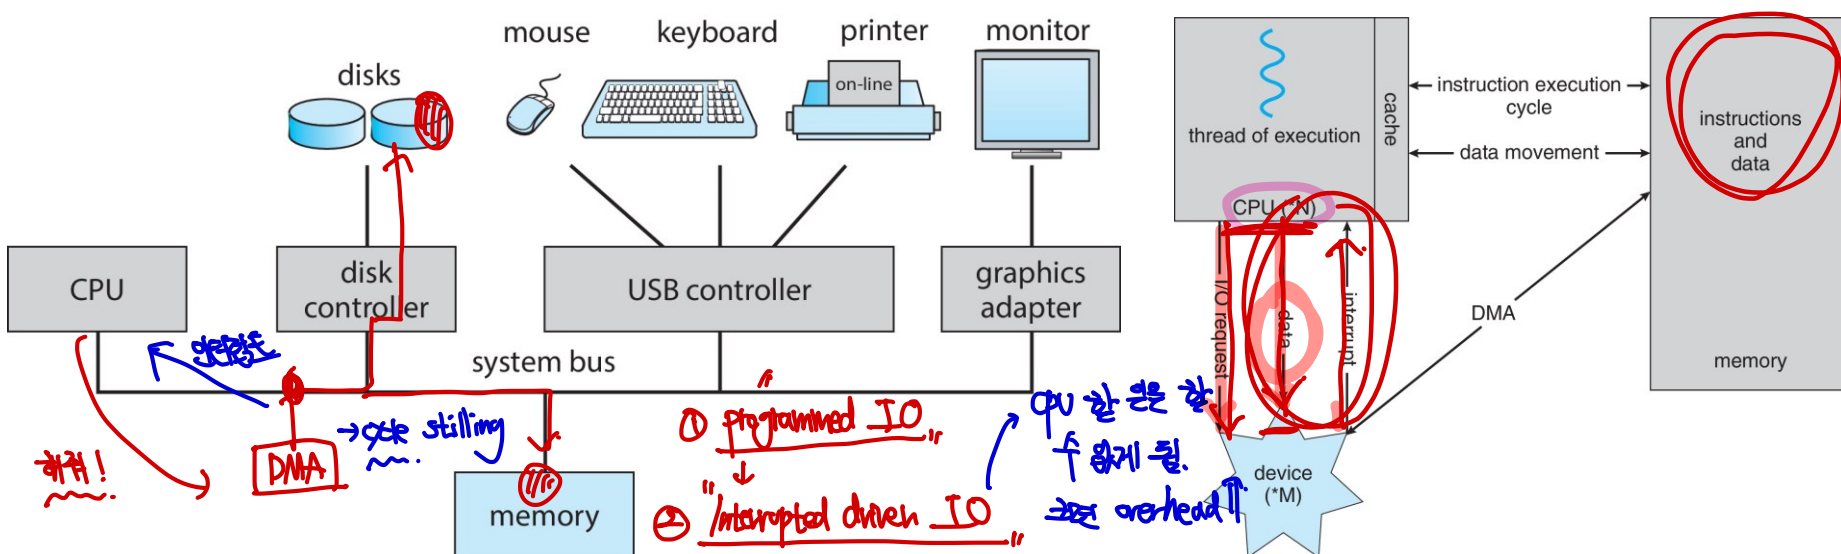


I/O Structure

cpu를 위해서 I/O장치와 Memory 사이의 데이터 전송을 담당하는 것
 → cpu의 처리량이 주변장치와 주변장치와의 데이터 전송에 의해

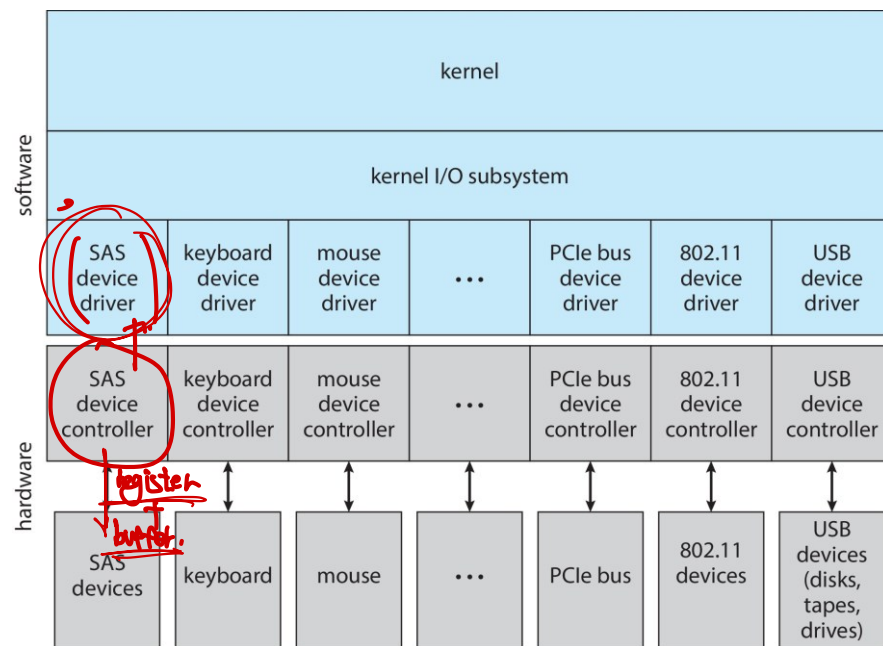
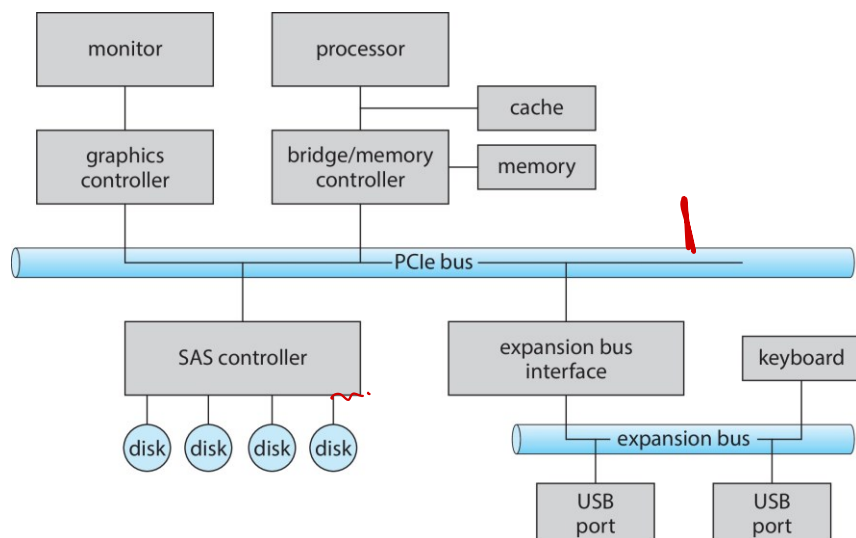
- The form of interrupt-driven I/O in a common bus is fine for moving small amounts of data
 - However, this can produce high overhead when used for bulk data movement
- To resolve this problem, **direct memory access (DMA)** is used
 - The device controller transfers data directly to/from the device and main memory with **no CPU intervention** by the DMA controller

→ 인터럽트의 발생 횟수를 최소화.
 - Only one interrupt is generated per block, rather than one interrupt per byte



Typical PC I/O Structure

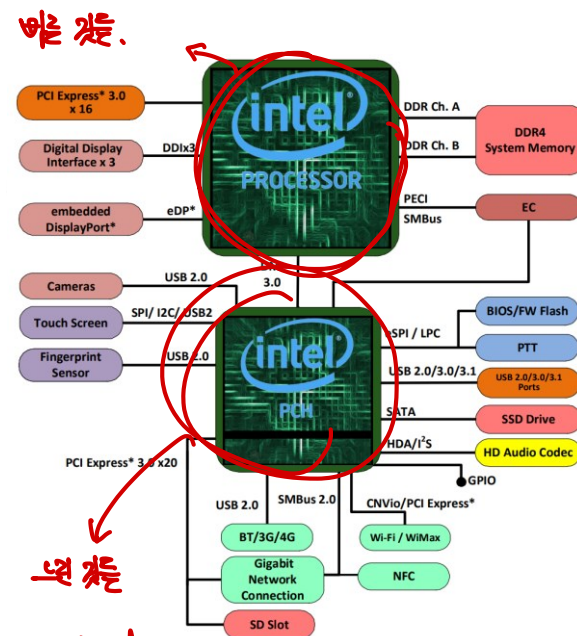
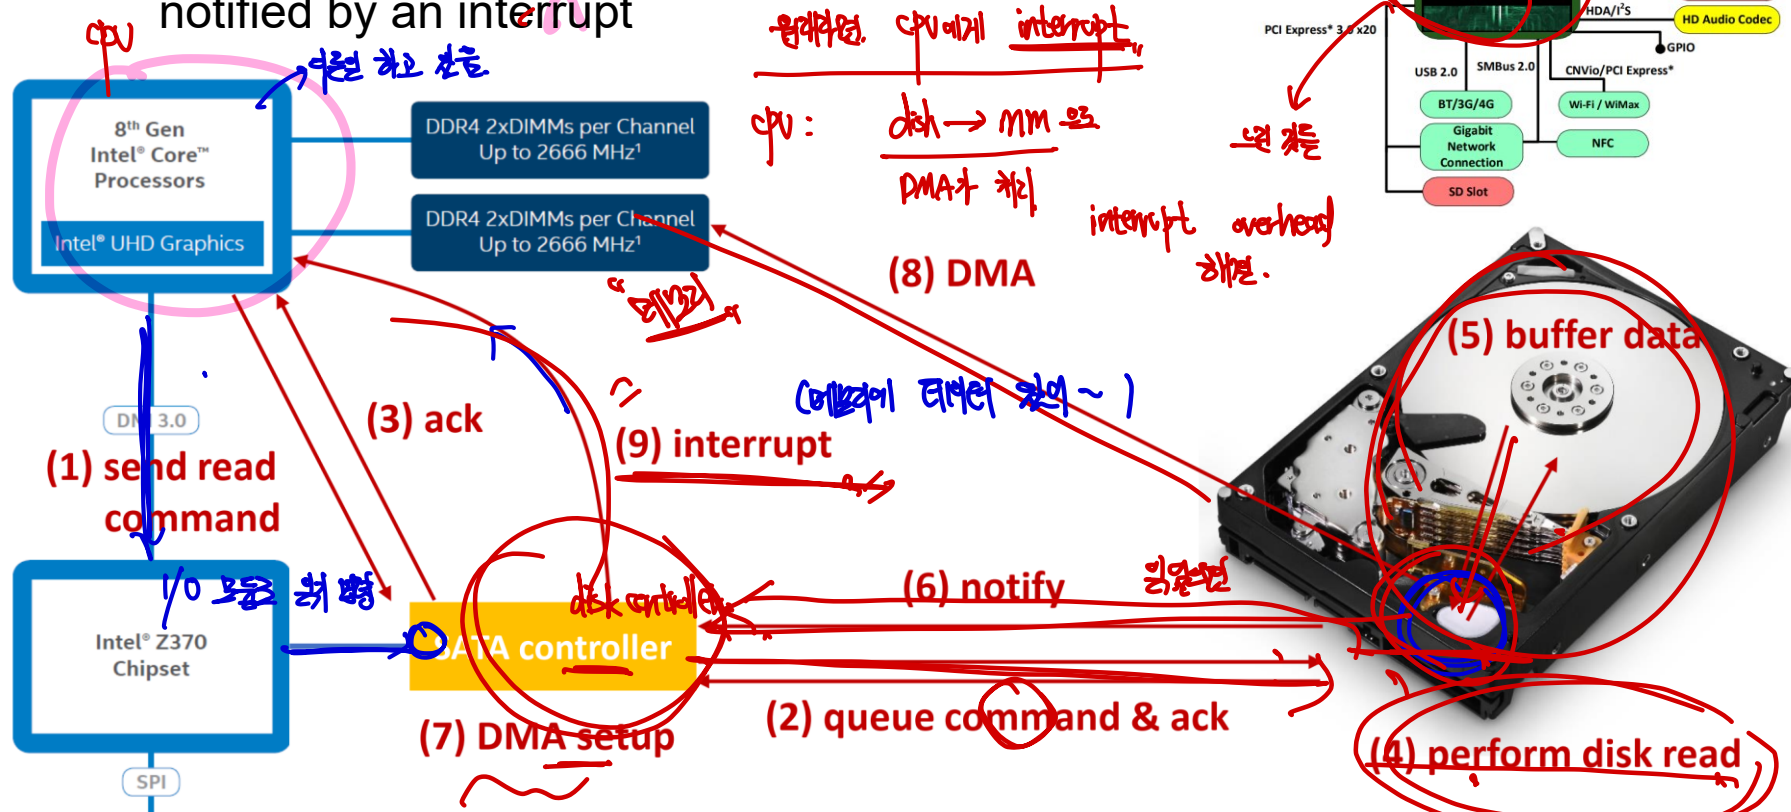
- **Computers operate a great many kinds of devices**
 - Storage devices, network devices, and human-interface devices
 - They are connected each other in a common bus (e.g. PCI bus)
- **Each OS has its own I/O subsystem structures and device driver frameworks**
 - The device drivers hide the details of device interactions from kernel



Disk I/O Example

■ SATA disk read operation using DMA

- Intel's processor platform point of view
- The CPU initiates a read command and forget
- Once the DMA transfer is done, the CPU gets notified by an interrupt



Courtesy of Prof. Jin-Soo Kim @ SNU