

Introduction to Operating Systems



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

What Happens when a Program Runs?

- **A running program does one very simple thing**
 - Executing instructions by von Neumann model of computing
 - The processor **fetches** an instruction from memory, **decode** and **executes**
 - Keep doing this until the program finally completes
- **There is a body of software that is responsible for making it easy to run programs, such as**
 - Allowing programs to share memory
 - Enabling programs to interact with devices
- **That body of software is called the **operating system (OS)****
 - It is in charge of making sure the system operates correctly and efficiently in an **easy-to-use** manner

Virtualization

→ 가상화.

추상화, 가상화.

■ The primary way the OS does this is through virtualization

- The OS takes a physical resource (i.e. processor, memory) and transforms it into a more general, powerful, and easy-to-use virtual form
- Thus, we sometimes refer to the OS as a virtual machine

■ A typical OS exports a few hundred system calls that are available to applications

운영체제 → (fopen)

- To allow users to tell the OS what to do and thus make use of the features of the virtual machine (i.e. running a program, allocating memory, accessing file)
- We sometime can say that the OS provides a standard library to applications

■ Virtualization allows many programs

- 1) to run (sharing the CPU), 2) to access devices (sharing the disks, etc), 3) to concurrently access their own instructions and data (sharing the memory)
- The OS is sometime known as a resource manager in which each of the CPU, memory, and disk is the resource of the system

↓
리소스관리

리소스들

가상화 해서 제공

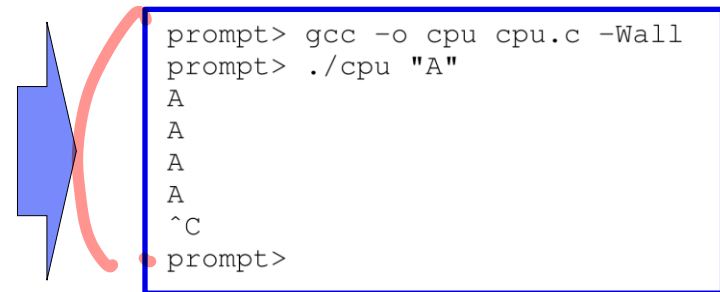
Virtualizing the CPU (1)

- Let's consider a simple example of code that loops and prints

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <assert.h>
#include "common.h"

int
main(int argc, char *argv[])
{
    if (argc != 2) {
        fprintf(stderr, "usage: cpu <string>\n");
        exit(1);
    }
    char *str = argv[1];
    while (1) {
        Spin(1);
        printf("%s\n", str);
    }
    return 0;
}
```

output of the program



```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
A
^C
prompt>
```

Spin() is a function that repeatedly checks the time and returns once it has run for a second

- When running it on a system with a single processor
 - Not too interesting of a run – repeatedly checks time until a sec. has elapsed
 - Run forever; by pressing “Control-C” to halt the program

Virtualizing the CPU (2)

- When running many programs at once on the same system,
 - Things are getting a little more interesting – somehow all four of these programs seems to be running at the same time


```

prompt> ./cpu A & ./cpu B & ./cpu C & ./cpu D &
[1] 7353
[2] 7354
[3] 7355
[4] 7356
A
B
D
C
A
B
D
C
A
...

```

서로 동시에

OS가 가상화된 CPU를 여러개
만들었다.



- The OS is in charge of this illusion that the system has a very large number of virtual CPUs.
 - Turning a single CPU into a seemingly infinite number of CPUs and thus allowing many programs to seemingly run at once
 - We call it virtualizing the CPU

Virtualizing Memory (1)

- **The physical memory model of modern machines is very simple**
 - Memory is just **an array of bytes** and is accessed using address and data
 - Memory is accessed all the time when a program is running (e.g. IF, LW, SW)
- **Let's look at a program that allocates some memory by malloc()**

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int
main(int argc, char *argv[])
{
    int *p = malloc(sizeof(int)); // a1
    assert(p != NULL);
    printf("(2134) address pointed to by p: %p\n", // a2
           getpid(), p); // a3
    *p = 0;
    while (1) {
        Spin(1);
        *p = *p + 1;
        printf("(2134) p: %d\n", getpid(), *p); // a4
    }
    return 0;
}
```



```
prompt> ./mem
(2134) address pointed to by p: 0x200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

- **a1)** allocate some memory, **a2)** print its address, **a3)** put zero into its first slot, **a4)** increase p by 1 and print process id (PID), which is unique per process
- In fact, this result is not too interesting

Virtualizing Memory (2)

- When running **multiple instances** of the same program,
 - The memory at the **same address (0x200000)** is allocated for two programs
 - Each program seems to be updating the value at 0x200000 **independently!**
 - It is as if each has **its own private memory**, instead of sharing memory

```

prompt> ./mem & ./mem
[1] 24113
[2] 24114
(24113) address pointed to by p: 0x200000
(24114) address pointed to by p: 0x200000
(24113) p: 1
(24114) p: 1
(24114) p: 2
(24113) p: 2
(24113) p: 3
(24114) p: 3
(24113) p: 4
(24114) p: 4
...

```

소프트웨어 마치 독립된
공간이 있다

→ 가상화된 메모리 공간. 실제 물리 공간.

→ 공유 공간

- This is what is happening here as the OS is **virtualizing memory**
 - Each process accesses its own private **virtual address space**, which the OS somehow maps onto the physical memory of the machine
 - A memory within a running program **does not affect the other's address space**
 - The reality is that physical memory is a shared resource, managed by the OS

Concurrency (1) → 동시성

- The OS is juggling many things at once by first running one process, then another, and so forth
 - Modern multi-threaded programs exhibit the problems of concurrency
- Let's demonstrate with an example of a multi-threaded program

- The main program creates two threads using `Pthread_create()`
- The thread is a function running within the same memory space
- Each thread calls `worker()` that increases a counter in a loop

- Output with `loops = 1000`

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value   : 2000
```



- When `loops = N`, the output is expected to be $2N$ since we have two threads

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"
#include "common_threads.h"

volatile int counter = 0;
int loops;

void *worker(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        counter++;
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: threads <value>\n");
        exit(1);
    }
    loops = atoi(argv[1]);
    pthread_t p1, p2;
    printf("Initial value : %d\n", counter);

    Pthread_create(&p1, NULL, worker, NULL);
    Pthread_create(&p2, NULL, worker, NULL);
    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("Final value   : %d\n", counter);
    return 0;
}
```


Concurrency (2)

- **Let's run the same, but with higher values for loops**

- When we gave an input of 100000, we first get 143012, instead of 200000
- We still get a wrong value of 137298 for the second try

```
prompt> ./thread 100000
Initial value : 0
Final value   : 143012    // huh??
prompt> ./thread 100000
Initial value : 0
Final value   : 137298    // what the??
```

- **The reason for these odd and unusual outcomes related to how instructions are executed, which is once at a time**

- Incrementing the shared counter takes three instructions: 1) load the counter value from memory into a register, 2) increment it, 3) store it back into memory
- These three instructions do not execute atomically, which causes the problem of concurrency

중요 문제

Persistence (1)

disk.
특성

- **Hardware and software are needed to store data persistently**
 - In system memory, data can be easily lost as devices, such as DRAM, store values in a volatile manner
 - The hardware comes in I/O devices, such as hard drive or solid-state drives
 - The software (OS) that manages the disk is called the file system, which is responsible for storing any files the user creates efficiently and reliably

- **Unlike the CPU and memory virtualization, the OS does not create a private, virtualized disk for each application**

- It is assumed that (users will want) to share information that is in files
- e.g.) google docs

↓
shared-disk

disk

데이터를 share 할 수 있어야 한다.

Persistence (2)

Let's look at some code for better understanding

```
#include <stdio.h>
#include <unistd.h>
#include <assert.h>
#include <fcntl.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    int fd = open("/tmp/file", O_WRONLY|O_CREAT|O_TRUNC,
                  S_IRWXU);
    assert(fd > -1);
    int rc = write(fd, "hello world\n", 13);
    assert(rc == 13);
    close(fd);
    return 0;
}
```

Handwritten notes: "file descriptor" with an arrow pointing to `fd`; "crash이 가능" with an arrow pointing to the `close(fd);` line.

- The program makes three calls into the OS: 1) `open()` to open the file, 2) `write()` to write some data to the file, 3) `close()` to close the file
- These system calls are routed to the part of the OS called the file system, which handles the requests

What the OS does in order to write to disk?


- Figure out where on disk this new data will reside
- Issue I/O request to the underlying storage device

The file system can handle system crashes during writes

Summary


■ Virtualization

- Processes
- Scheduling
- Virtual memory



■ Concurrency

- Threads
- Locks
- Semaphores



■ Persistence

- I/O devices
- File systems



Courtesy of Prof. Jin-Soo Kim @ SNU

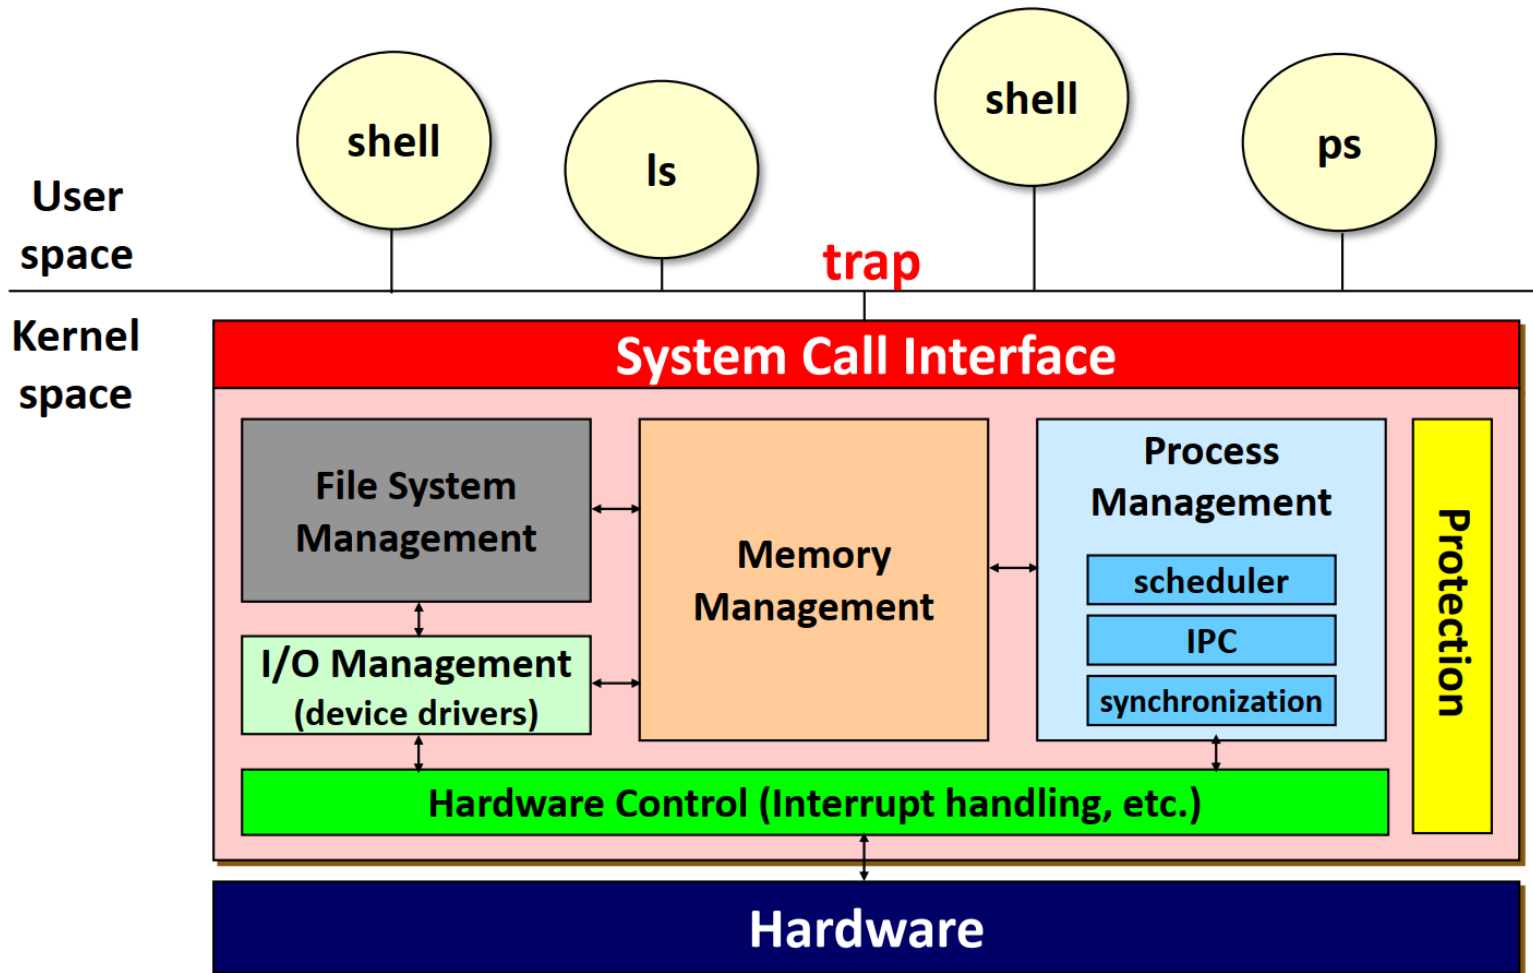
The Abstraction: The Process



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

OS Internals




Courtesy of Prof. Jin-Soo Kim @ SNU

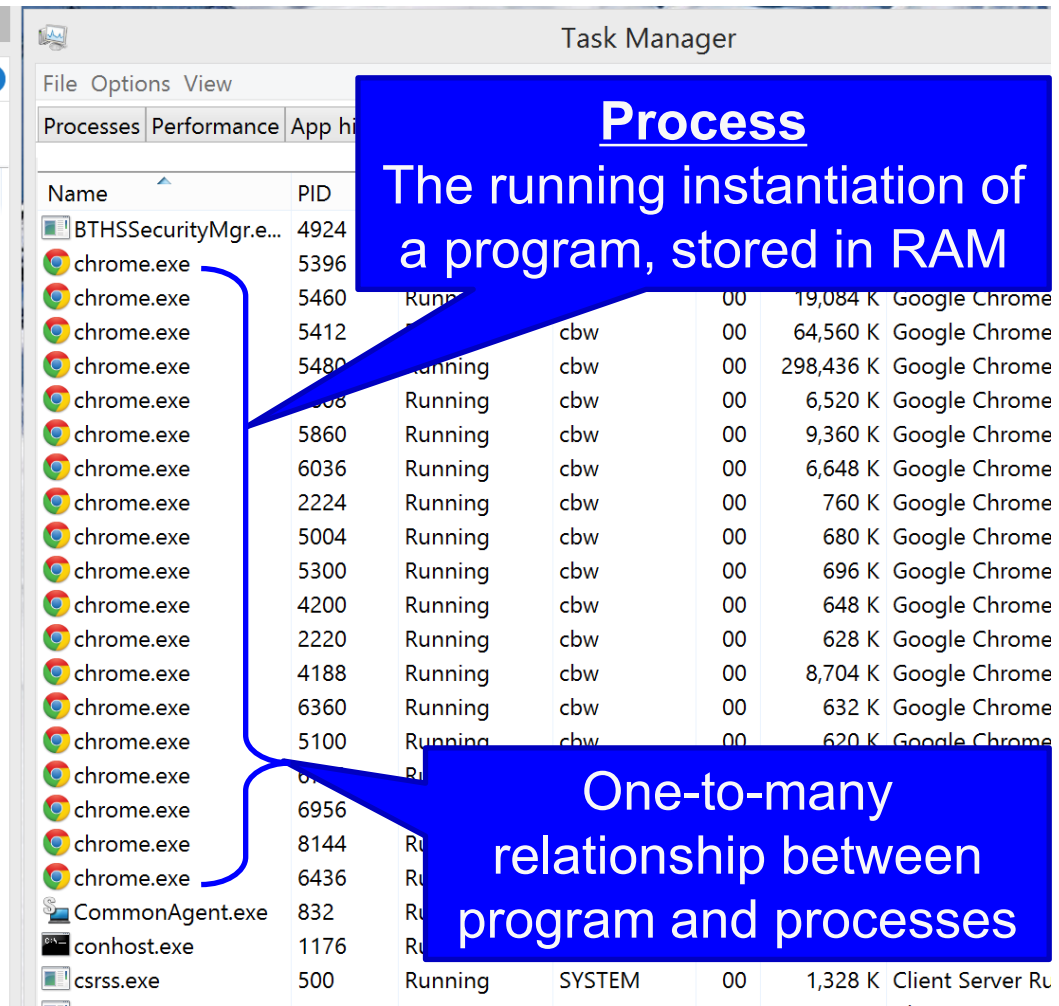
How to Provide the Illusion of Many CPUs?

- **The definition of a process, informally, is a running program**
 - The program itself is a lifeless thing; just sitting there on the disk, a bunch of instructions with some data
 - The OS gets them running, transforming it into something useful
 - We often want to run more than one program at once
- **The OS creates the illusion by virtualizing the CPU**
 - The OS can promote the illusion that many virtual CPUs exist when in fact there is only one physical CPU by **time sharing** of the CPU
 - Time sharing is implemented by running one process, then stopping it and running another, and so forth, and its potential **cost is performance**
- **To implement the virtualization, the OS will need both**
 - some low-level machinery: **mechanisms**; how to do something? (context switching)
 - some high-level intelligence: **policies**; what should be done? (scheduling)

Program vs Process



Program
An executable file in long-term storage



Process
The running instantiation of a program, stored in RAM

One-to-many relationship between program and processes

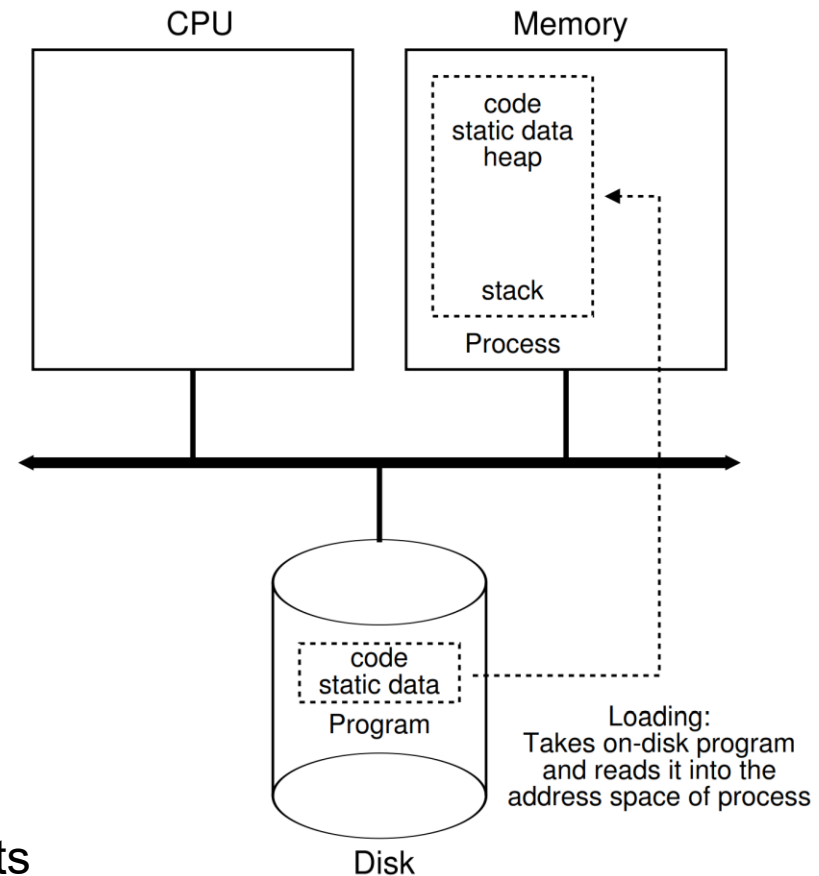
The Abstraction: A Process and API

- The abstraction provided by OS of a running program is **process**
- To understand what constitutes a process, we have to know its **machine state**
 - The machine state that comprises a process are:
 - 1) **Memory**: instructions and data lie in memory (address space)
 - 2) **Registers**: many instructions explicitly read or update registers, such as program counter, stack pointer, and frame pointer
 - 3) **Storage**: I/O information that the process is currently accessing (e.g. file)
- **Modern OS has an interface by providing APIs:**
 - **Create**: create a new process to run a program
 - **Destroy**: halt a runaway process (interface to destroy processes forcefully)
 - **Wait**: wait for a process to stop running
 - **Miscellaneous control**: other controls (e.g. suspend a process and resume it)
 - **Status**: get some status information about a process

Process Creation: A Little More Detail (1)

1) The first step is to load its code and any static data into memory, into the address space of the process

- Programs initially reside on disk in an **executable format**
- Early OSes load **eagerly** all at once before running the program
- Modern OSes perform the load process **lazily**, by loading pieces of code or data only as they are needed during program execution (swapping, paging)



2) Some memory is allocated for the program's run-time stack

- The stack for local variable, function parameters, and return addresses
- The OS initializes the stack with arguments (i.e. **argc** and **argv** of **main()**)

Process Creation: A Little More Detail (2)

3) The OS allocates some memory for the program's heap

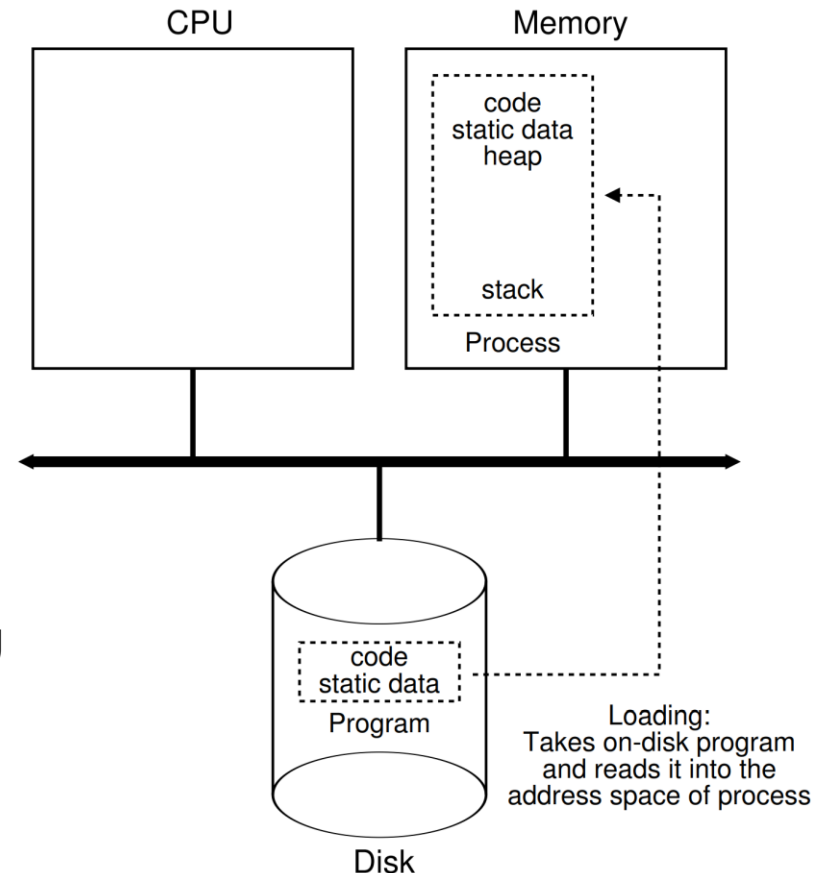
- The heap for explicitly requested dynamically-allocated data (e.g. `malloc()`)

4) The OS does some other initialization tasks, related to I/O

- e.g.) each process in UNIX systems has three open file descriptors (standard input, output, error)

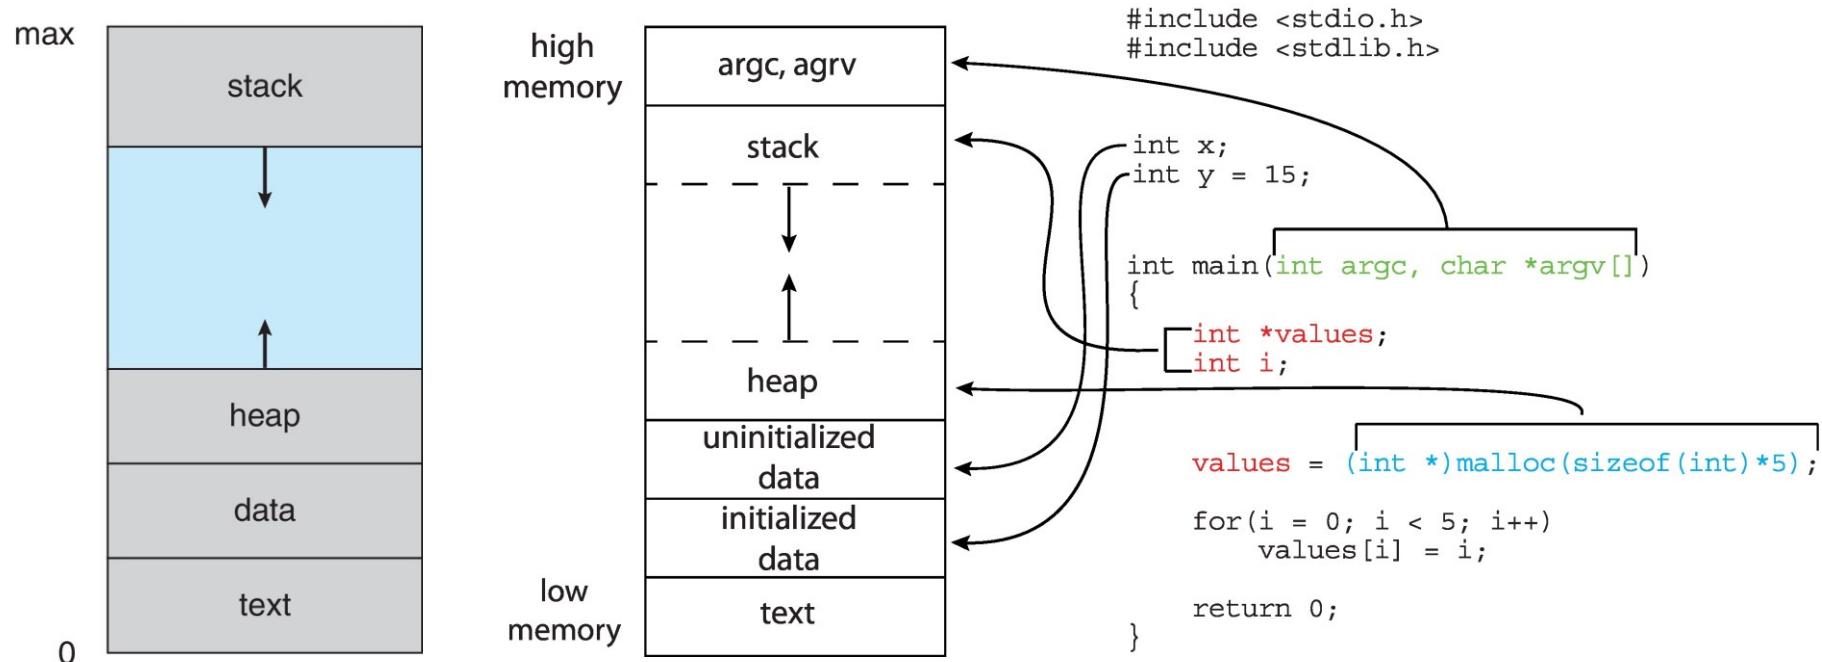
5) The OS starts the program running at the entry point, `main()`

- Then, the OS transfers control of the CPU to the newly-created process
- Thus, the program begins its execution



Memory Layout of a C Program

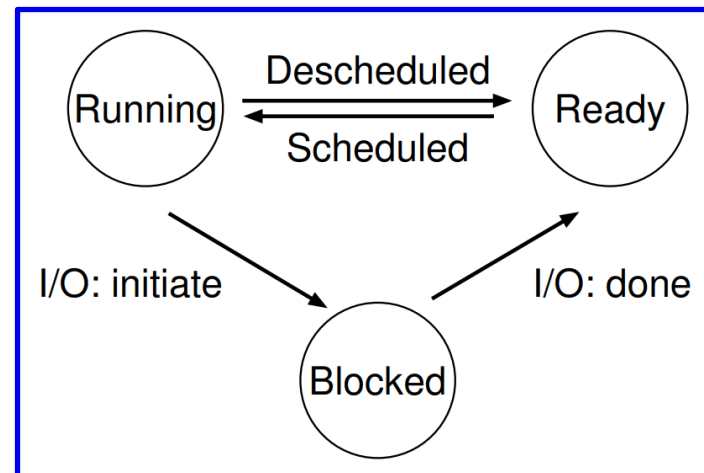
- **The memory layout of a process is typically divided into:**
 - 1) **text** (executable code), 2) **data** (global variable),
 - 3) **heap** (dynamically allocated memory), 4) **stack** (temporary data storage)
 - The data section is divided into (a) initialized data and (b) uninitialized data
 - A separate section is provided for the **argc** and **argv** parameters
- **Remember that each process has its own private address space**



Process States

- In a simplified view, a process can be in one of three states:

- Running**: a process is running on a CPU (executing instructions)
- Ready**: a process is ready to run, but the OS chosen not to run at this moment
- Blocked**: a process has performed some kind of operation that makes it not ready to run until some other event takes place (e.g. when a process requests an I/O)



Examples

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process ₁ now done

Tracing process states (CPU only)

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked, so Process ₁ runs
5	Blocked	Running	
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

Tracing process states (CPU and I/O)

Data Structures

- **The OS has some key data structures that track various relevant pieces of information:**
 - **Process list** that includes all of the ready, blocked, and running processes
 - **Register context** to hold the register contents for a stopped process, which will be used for **context switching** (save them when stop → restore them when resume)
- **Each process is represented by a process control block (PCB)**
 - C-structure that contains all the information about a process

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

xv6's register context and process states

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                                // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If !zero, sleeping on chan
    int killed;               // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;        // Current directory
    struct context context;    // Switch here to run process
    struct trapframe *tf;     // Trap frame for the
                                // current interrupt
};
```

xv6's PCB (process descriptor)