

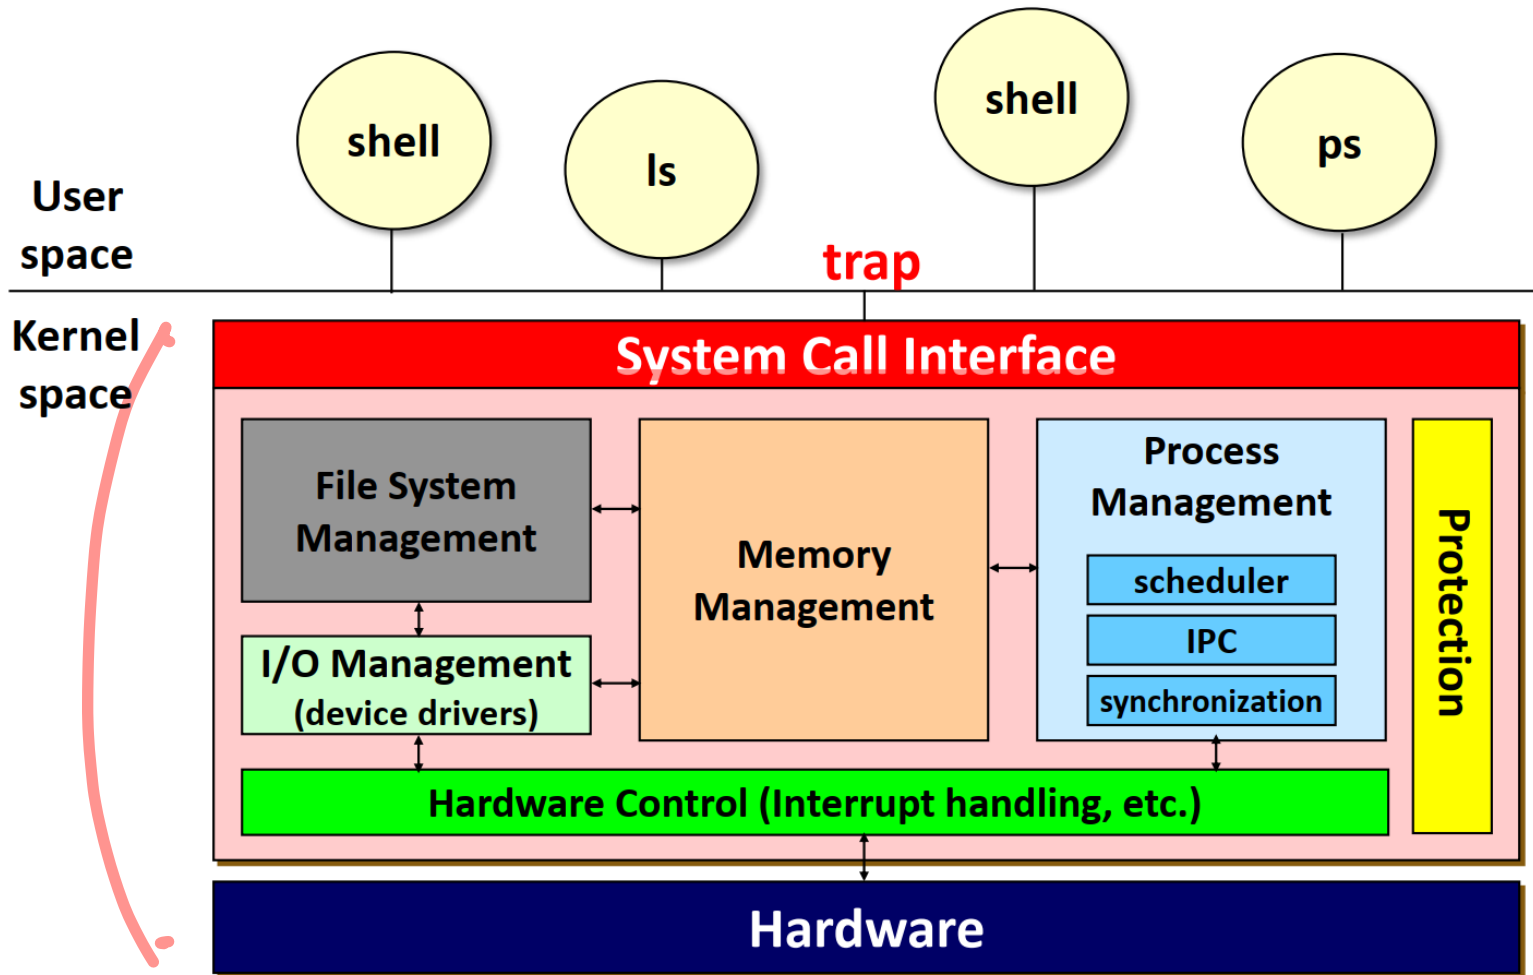
The Abstraction: The Process



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

OS Internals




Courtesy of Prof. Jin-Soo Kim @ SNU

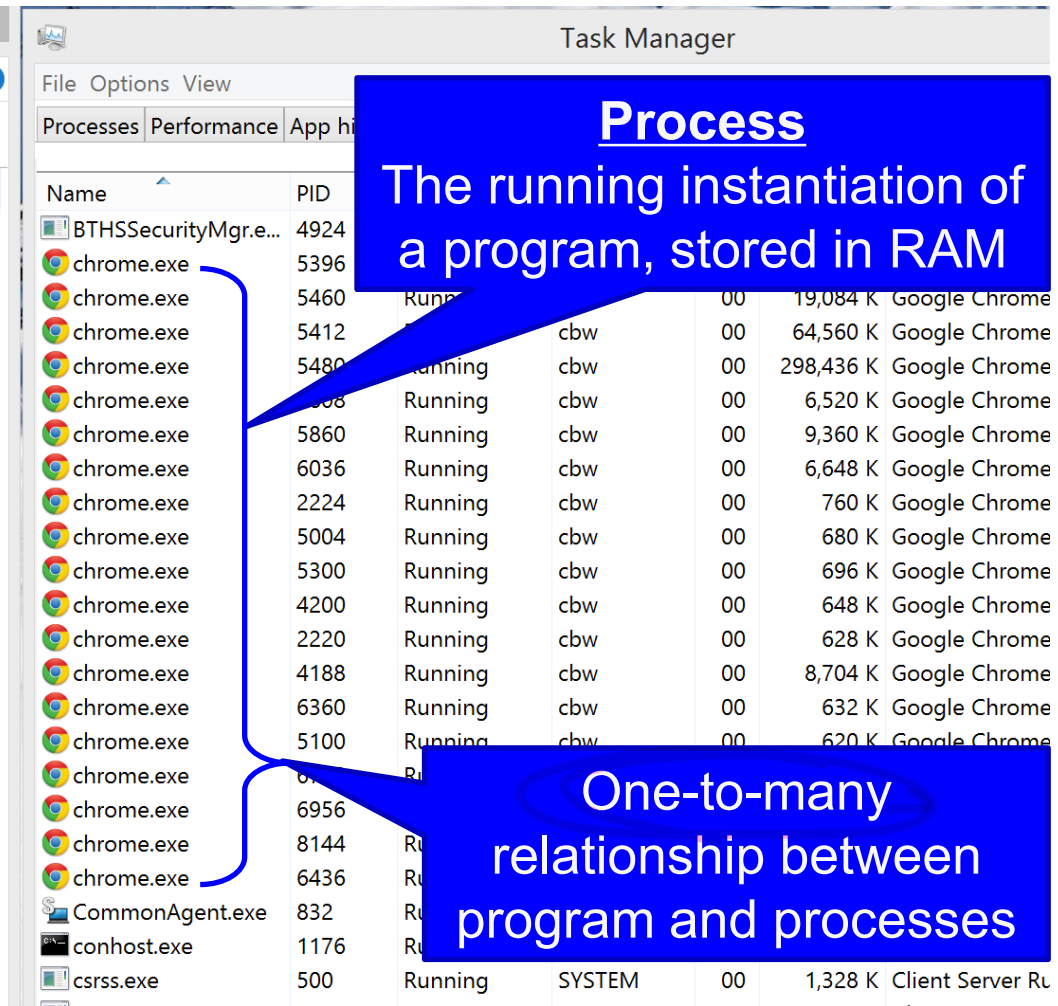
How to Provide the Illusion of Many CPUs?

- **The definition of a process, informally, is a running program**
 - The program itself is a lifeless thing; just sitting there on the disk, a bunch of instructions with some data
 - The OS gets them running, transforming it into something useful
 - We often want to run more than one program at once
- **The OS creates the illusion by virtualizing the CPU**
 - The OS can promote the illusion that many virtual CPUs exist when in fact there is only one physical CPU by **time sharing** of the CPU
 - Time sharing is implemented by running one process, then stopping it and running another, and so forth, and its potential **cost is performance**
- **To implement the virtualization, the OS will need both**
 - some low-level machinery: **mechanisms**; how to do something? (context switching)
 - some high-level intelligence: **policies**; what should be done? (scheduling)

Program vs Process



Program
An executable file in long-term storage



Process
The running instantiation of a program, stored in RAM

One-to-many relationship between program and processes


The Abstraction: A Process and API

- The abstraction provided by OS of a running program is **process**
- To understand what constitutes a **process**, we have to know its **machine state**
 - The machine state that comprises a process are:
 - 1) **Memory**: instructions and data lie in memory (address space)
 - 2) **Registers**: many instructions explicitly read or update registers, such as program counter, stack pointer, and frame pointer
 - 3) **Storage**: I/O information that the process is currently accessing (e.g. file)
- **Modern OS has an interface by providing APIs:**
 - **Create**: create a new process to run a program
 - **Destroy**: halt a runaway process (interface to destroy processes forcefully)
 - **Wait**: wait for a process to stop running
 - **Miscellaneous control**: other controls (e.g. suspend a process and resume it)
 - **Status**: get some status information about a process

Process Creation: A Little More Detail (1)

1) The first step is to load its code and any static data into memory, into the address space of the process


- Programs initially reside on disk in an **executable format** → 한 번에 실행, 프로그램은
- Early OSes load **eagerly** all at once before running the program ↓
현재는 LA2 체계
- Modern OSes perform the load process **lazily**, by loading pieces of code or data only as they are needed during program execution (swapping, paging)



2) Some memory is allocated for the program's run-time stack

- The stack for local variable, function parameters, and return addresses
- The OS initializes the stack with arguments (i.e. `argc` and `argv` of `main()`)

할당 해 주고 CPU가
실행할 수 있게끔
세팅 해놓음



Process Creation: A Little More Detail (2)

3) The OS allocates some memory for the program's heap.

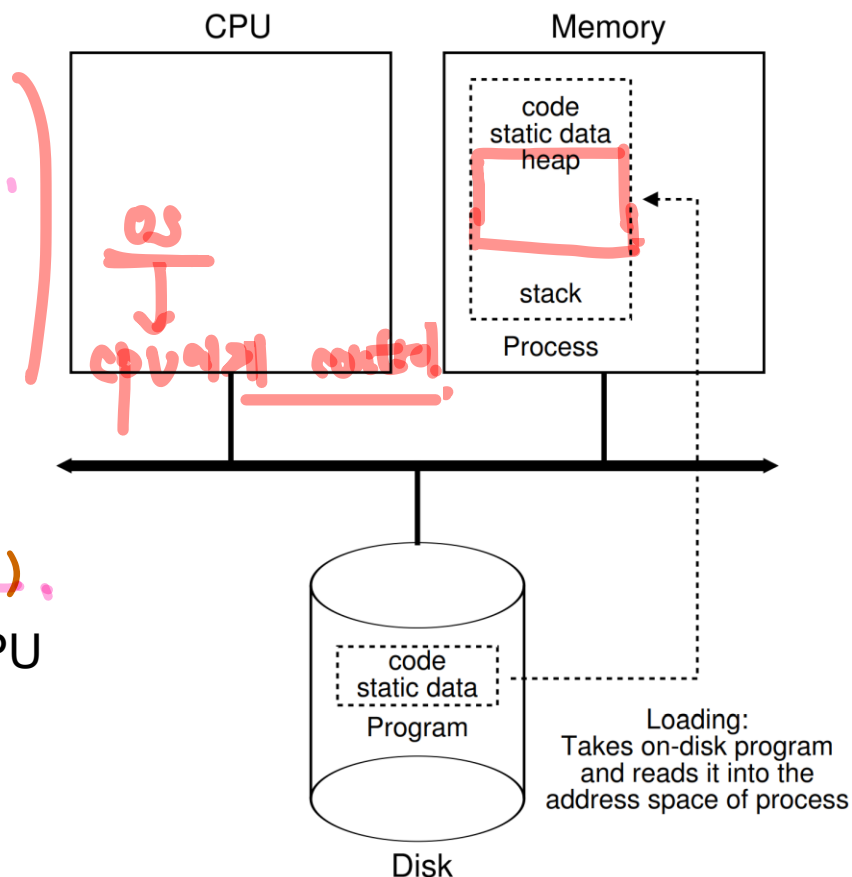
- The heap for explicitly requested dynamically-allocated data (e.g. `malloc()`)

4) The OS does some other initialization tasks, related to I/O.

- e.g.) each process in UNIX systems has three open file descriptors (standard input, output, error).

5) The OS starts the program running at the entry point, main().

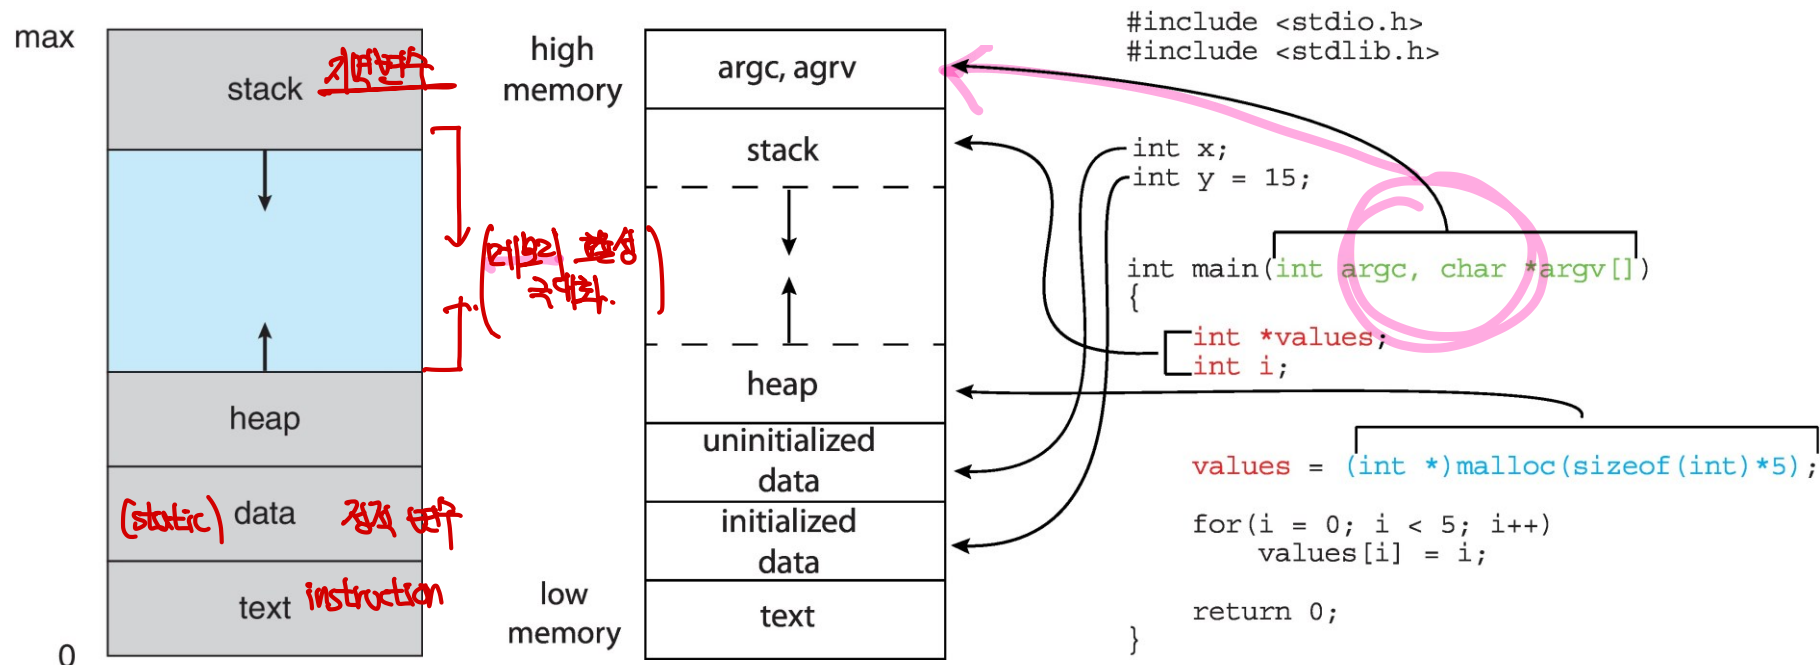
- Then, the OS transfers control of the CPU to the newly-created process
- Thus, the program begins its execution



Memory Layout of a C Program

실행 프로그램 → 인스트럭션

- The memory layout of a process is typically divided into:
 - 1) text (executable code), 2) data (global variable),
 - 3) heap (dynamically allocated memory), 4) stack (temporary data storage)
 - The data section is divided into (a) initialized data and (b) uninitialized data
 - A separate section is provided for the `argc` and `argv` parameters
- Remember that each process has its own private address space

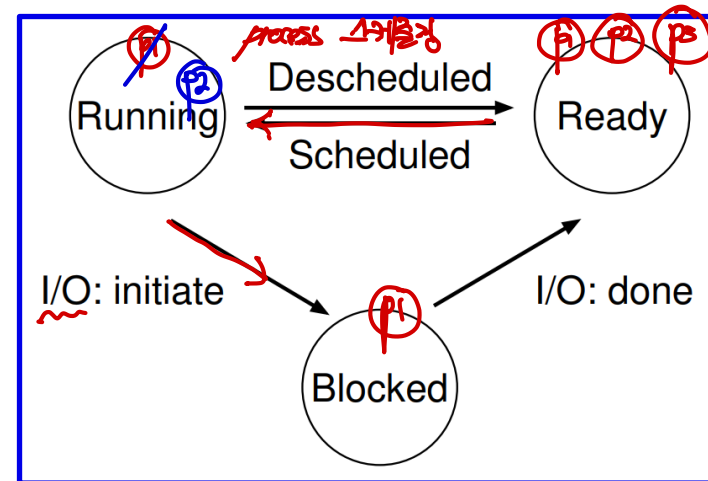


Process States

실행 상태 프로세스 수 = CPU의 개수

- In a simplified view, a process can be in one of three states:

- Running**: a process is running on a CPU (executing instructions)
- Ready**: a process is ready to run, but the OS chosen not to run at this moment
- Blocked**: a process has performed some kind of operation that makes it not ready to run until some other event takes place (e.g. when a process requests an I/O)



Examples

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ <u>now done</u>
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process ₁ now done

Tracing process states (CPU only)

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked, so <u>Process₁</u> runs
5	Blocked	Running	
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	<u>Process₁</u> now done
9	Running	–	
10	Running	–	Process ₀ now done

Tracing process states (CPU and I/O)

Data Structures

- The OS has some key data structures that track various relevant pieces of information:
 - Process list that includes all of the ready, blocked, and running processes
 - Register context to hold the register contents for a stopped process, which will be used for context switching (save them when stop → restore them when resume)
- Each process is represented by a process control block (PCB)
 - C-structure that contains all the information about a process

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};
```

```
// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
    RUNNABLE, RUNNING, ZOMBIE };
```

xv6's register context and process states

Machine state를 (최소 상태를 가) (Memory/register storage) ↑

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem; // Start of process memory
    uint sz; // Size of process memory
    char *kstack; // Bottom of kernel stack
                // for this process
    enum proc_state state; // Process state
    int pid; // Process ID
    struct proc *parent; // Parent process
    void *chan; // If !zero, sleeping on chan
    int killed; // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd; // Current directory
    struct context *context; // Switch here to run process
    struct trapframe *tf; // Trap frame for the
                          // current interrupt
};
```

운영체제
xv6's PCB (process descriptor)