

Interlude: Process API



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

The fork () System Call

- The **fork ()** system call is used to create a new process

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n",
               rc, (int) getpid());
    }
    return 0;
}
```

→ rc = child의 pid를 return

```
prompt> ./p1
hello world (pid:29146)
① hello, I am parent of 29147 (pid:29146)
② hello, I am child (pid:29147)
prompt>
```

```
prompt> ./p1
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

The output is **not deterministic**;
the **OS scheduler** determines which
one runs at a given moment in time

- The process starts with printing “hello world” with its process identifier (PID)
- The newly-created process (child) by **fork ()** is a copy of the calling process; but it does not start running at **main ()**, which is **not an exact copy**
- The child process now has its **own address space, registers, PC, and so forth**
- The value it returns to the call of **fork ()** is different (child: 0, parent: PID)

The wait () System Call

- The **wait ()** system call is used to wait for a child to finish

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {          // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {               // parent goes down this path (main)
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n",
               rc, rc_wait, (int) getpid());
    }
    return 0;
}
```

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (rc_wait:29267) (pid:29266)
prompt>
```

← child의 pid.

- The parent process calls **wait ()** to delay its execution the child finishes executing; when child is done, **wait ()** returns to the parents
- This **wait ()** call makes the output **deterministic**


The exec () System Call

- exec ()** is to run a program that is different from calling program


```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>
```

```
int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {          // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("p3.c"); // argument: file to count
        myargs[2] = NULL;          // marks end of array
        execvp(myargs[0], myargs); // runs word count
        printf("this shouldn't print out");
    } else {                // parent goes down this path (main)
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n",
               rc, rc_wait, (int) getpid());
    }
    return 0;
}
```

```
prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
      29      107      1030 p3.c
hello, I am parent of 29384 (rc_wait:29384) (pid:29383)
prompt>
```



- The **exec ()** loads code and static data from the executable and overwrites its current code segment and static data with it (**replacing** the current with a new)
- Also, the heap, stack and other parts of its memory space are **re-initialized**
- A successful call to **exec ()** never returns



- The **shell** is a user program that shows you a prompt and waits for you to type, then **fork()** and **exec()** with **wait()** when typed

- ```
prompt> wc p3.c > newfile.txt
```

```
prompt> ./p4
prompt> cat p4.output
 32 109 846 p4.c
prompt>
```

- The redirection can be implemented by closing STDOUT (screen) and opening a file to be redirected between **fork()** and **exec()**

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <fcntl.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
 int rc = fork();
 if (rc < 0) {
 // fork failed
 fprintf(stderr, "fork failed\n");
 exit(1);
 } else if (rc == 0) {
 // child: redirect standard output to a file
 close(STDOUT_FILENO);
 open("../p4.output", O_CREAT | O_WRONLY | O_TRUNC, S_IRWXU);
 // now exec "wc"...
 char *myargs[3];
 myargs[0] = strdup("wc"); // program: wc (word count)
 myargs[1] = strdup("p4.c"); // arg: file to count
 myargs[2] = NULL; // mark end of array
 execvp(myargs[0], myargs); // runs word count
 } else {
 // parent goes down this path (main)
 int rc_wait = wait(NULL);
 }
 return 0;
}
```