

Scheduling: Introduction




Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Workload Assumptions and Scheduling Metrics

- Now, we have to understand the high-level policies that an OS scheduler employs
 - Before diving into the scheduling, let's make some assumption about the processes running in the system, which is called the workload
- We make the following assumptions about the processes (jobs)
 - 1) Each job run for the same amount of time
 - 2) All jobs arrives at the same time
 - 3) Once started, each job runs to completion
 - 4) All jobs only use the CPU (no I/O)
 - 5) The run-time of each job is known
- We need one thing to compare different scheduling policies
 - Turnaround time is performance metric defined by $T_{turnaround} = T_{completion} - T_{arrival}$
 - In our assumptions, $T_{arrival} = 0$ and hence $T_{turnaround} = T_{completion}$
 - Another metric is fairness, which is often at odds with performance in scheduling



→ 항상 문제다

First In, First Out (FIFO) → (먼저 들어 프로세스 들어 먼저 나간다.)


- The most basic algorithm is **First In, First Out (FIFO)** scheduling

- This is sometimes called **First Come, First Served (FCFS)**
- It is simple and easy to implement, and works well in some cases


- **Let's do a quick example**

1) Three jobs (A, B, C) arrive at the same time ($T_{arrival} = 0$) and each job runs for 10; A arrived just a hair before B which did just before C

2) What if A runs for 100 while B and C run for 10 each? (**relaxing assumption #1**)



$$1) T_{turnaround} = (10+20+30)/3 = 20$$



$$2) T_{turnaround} = (100+110+120)/3 = 110$$

- The **FCFS** can cause a **convoy effect**.

- A number of relatively-short potential consumers of a resource get queued behind a heavyweight resource consumer → **long waiting time**

FIFO FCF

모든 프로세스에 선점권 있는 relaxing FIFO x

그럼 프로세스 단위 SJF

모든 프로세스에 선점권 있는

relaxing SJF x

모든 프로세스에 선점권 있는

STCF

다시 시작할 수 있는 (I/O x)

relaxing

→ preemptive SJF와 turnaround

시간

Response time 시간

모든 프로세스에 선점권 있는

선점권 있는

STCF

중단

turn-around time 다들

time slice context-switching

프로세스 단위로 나누기

PP 프로세스

relaxing

STCF x

선점권 있는

MLFQ

→ 프로세스 단위로 선점권 있는

선점권 있는

↓

PP

프로세스 단위로

Reset

MLFQ # 프로세스 단위로

프로세스 단위로 → PP

프로세스 단위로

프로세스 단위로

time slice

프로세스 단위로

프로세스

다들 ↓ batch

Shortest Job First (SJF)


- Shortest Job First (SJF) algorithm runs the shortest job first, then the next shortest, and so on

- This is a very simple approach to avoid the convoy effect


- Let's take the same example but with SJF

- 1) SJF reduces average turnaround time from 110 to 50

- 2) What if A arrives at $t=0$ and B and C arrive at $t=10$? (relaxing assumption #2)



$$1) T_{\text{turnaround}} = (10+20+120)/3 = 50$$



$$2) T_{\text{turnaround}} = (100 + (110-10) + (120-10))/3 = 103.33$$

- SJF is an **optimal** scheduling if all jobs arrive at the same time


- Otherwise, it can also suffer from the same convoy problem as shown in 2)

Shortest Time-to-Completion First (STCF)

- **Shortest Time-to-Completion First (STCF) algorithm adds the preemption on top of SJF**
 - This is also known as **Preemptive Shortest Job First (PSJF)** scheduler
 - Anytime a new job enters, the STCF scheduler determines which of remaining jobs (including the new job) has the least time left, and schedules that one
- **Preemptive vs Non-preemptive scheduling**
 - Preemptive: can stop a job to run another, require the context switch
 - Non-preemptive: run a job to completion (can't stop a job until finishing)

Let's take the same example but with STCF

- STCF reduces average turnaround time from 103.3 to 50



$$T_{\text{turnaround}} = ((120-0) + (20-10) + (30-10)) / 3 = 50$$

가장 먼저 turnaround 끝까지 줄었.

- **STCF is provably optimal** given our new assumptions

A New Metric: Response Time

- **The turnaround time is a good metric for early batch systems**
 - However, the introduction of time-shared machines changed all that
 - Now users would sit a terminal and demand interactive performance as well

- **A new metric was born: response time**

(batch-system) $\begin{pmatrix} A \\ B \\ C \end{pmatrix}$

→ 도착해서 작업을 시작하는 데에 걸리는 시간.

 - We define response time as time from when the job arrives to the first time it scheduled, more formally: $T_{response} = T_{firstrun} - T_{arrival}$

$T_{firstrun}$: Queue에 들어간 시간.

- **STCF and related disciplines are not particularly good for response time**
 - Imagine sitting at a terminal, typing, and having to wait 10 seconds to see a response from the system; not too pleasant
- **Then, how can we build a scheduler sensitive to response time?**

Round Robin (RR) → 표준, 다른 것들과의 경쟁에서 쓸.




context-switching이 너무 많이 발생.

■ To solve this, we introduce Round Robin (RR) scheduling


- Instead of running jobs to completion, RR runs a job for a time slice (scheduling quantum, time quantum) and then switches to the next job in the run queue
- It repeatedly does so until the jobs are finished; called time-slicing
- The length of a time slice must be a multiple of the timer-interrupt period

■ Let's look at an example

- Three jobs (A, B, C) arrive at the same time, each runs for 5, time slice is 1



SJF: $T_{\text{response}} = (0+5+10)/3 = 5$



turnaround time. 관점/문제. 존재.

RR: $T_{\text{response}} = (0+1+2)/3 = 1$

■ As can be seen, the length of the time slice is critical for RR


- The shorter it is, the better the response time; too short → context switch cost ↑
- Deciding on the length presents a trade-off to a system designer

Incorporating I/O *(printf)*


- Obviously, all programs perform I/O and let's relax assumption #4
 - A scheduler has a decision when a job has an I/O request and I/O complete

Let's consider an example

- Two jobs (A, B), which each needs 50 ms of CPU time. **A** runs for 10 ms and then requests I/O (takes 10 ms) and **B** uses the CPU for 50 ms with no I/O



Waiting for I/O completion:
Poor use of resources



Scheduling for another job:
Better use of resources

- For better utilizing the processor, the scheduler should make
 - the CPU-intensive jobs (job B) run while the interactive jobs (job A) are performing I/O