

# The Abstraction: Address Spaces



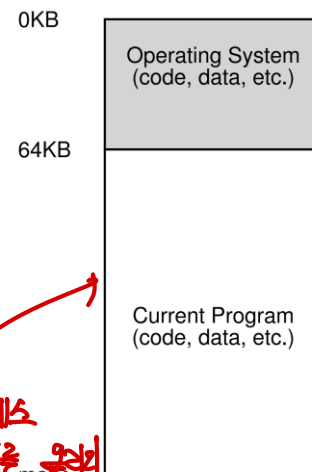
**Prof. Yongtae Kim**

Computer Science and Engineering  
Kyungpook National University

# Early Systems


## ■ In early days, building computer systems was easy

- From the memory perspective, early machines didn't provide much of an abstraction to users
- The physical memory of the machine contains the OS as a set of routines (library) and one running process
- Life was sure easy for OS developers in those days



## ■ Computer systems were getting more complicated

- In the **multiprogramming era**, multiple processes were ready to run at a given time, and the OS would switch between them
- Doing so increased the utilization of the CPU and efficiency
- For **time sharing era**, running one process with full access to all memory causes a big problem (memory swapping)
- So, processes were stayed in memory while switching between them, allowing OS to implement timing sharing efficiently
- This leads to **protection issue** (e.g. A can access B's memory)



# The Address Space

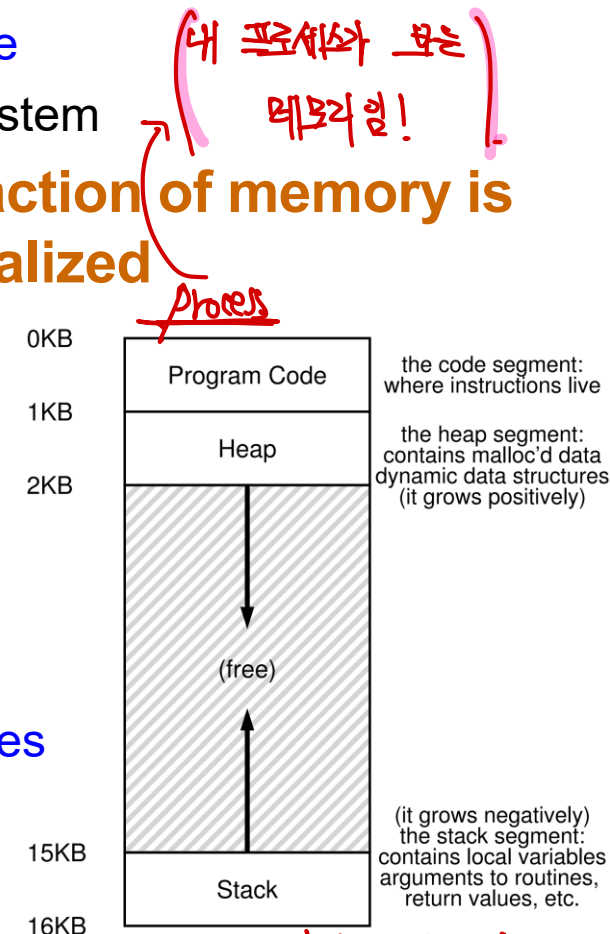
→ 메모리를 추상화!

## ▪ The OS creates an easy to use abstraction of physical memory

- This memory abstraction is called the **address space**
- It is the running program's view of memory in the system

## ▪ Understanding this fundamental OS abstraction of memory is key to understanding how memory is virtualized

- The address space of a process contains **all of the memory state** of running program
- **Code**: where instructions live
- **Stack**: local variables, parameters, return address
- **Heap**: dynamically allocated memory (`malloc()`)
- The program isn't in memory in physical addresses 0 ~ 16KB; rather at some **arbitrary physical addresses**
- How can the OS build this abstraction of a private, potentially large address space for multiple running processes on top of a single, physical memory?



(내 프로세스의 모든 메모리일!)

Process

→ 인접한 프로세스

## ▪ The OS is virtualizing memory with some hardware support

# Goals

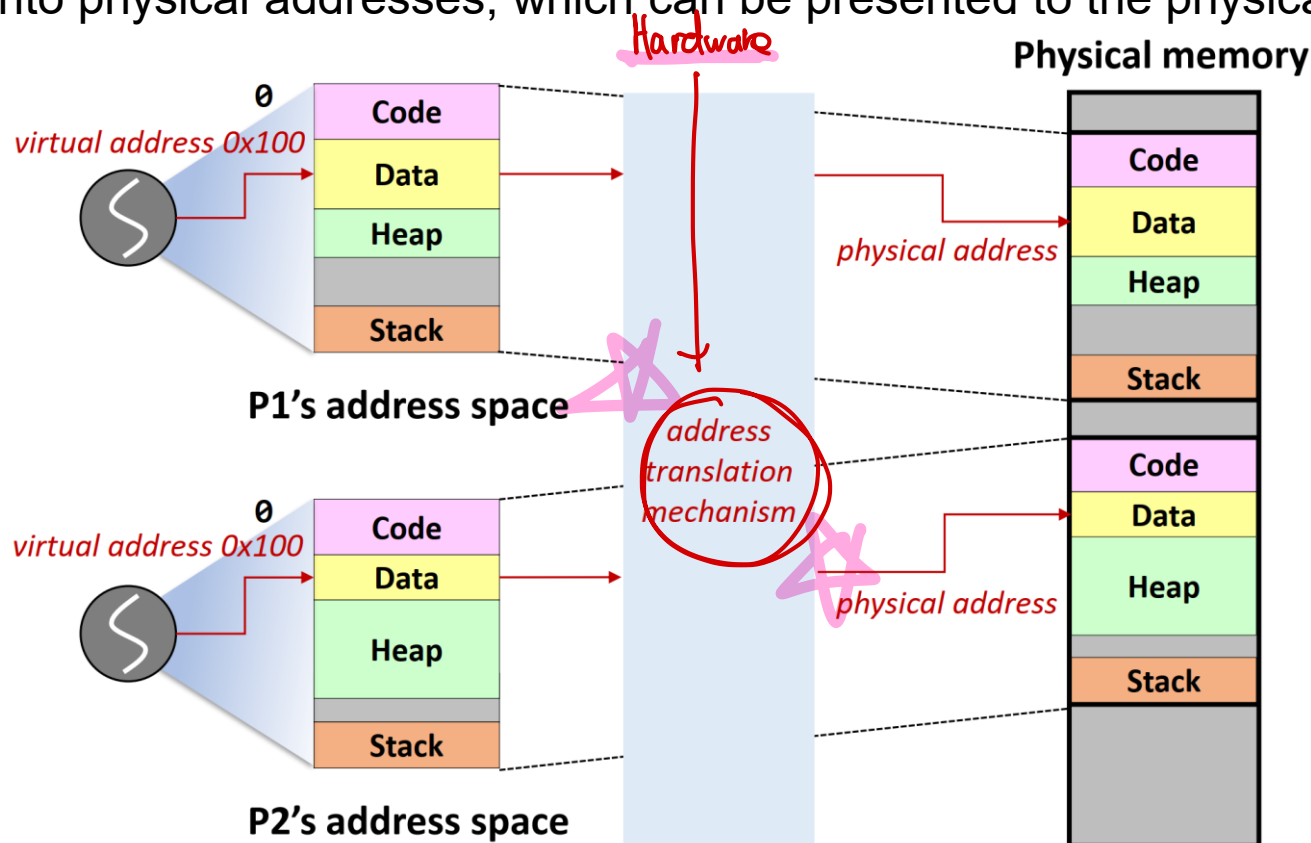
추가 4 항목을.

- **One major goal of a virtual memory (VM) system is transparency**
  - The program shouldn't be aware of the fact that memory is virtualized
  - The program behaves as if it has its own private physical memory
- **Another goal of VM is efficiency**
  - The OS should strive to make the virtualization as efficient as possible both in terms of time and space
  - The OS will have to rely on hardware support (e.g. TLBs)
- **Finally, a third VM goal is protection**
  - The OS should make sure to protect processes from one another as well as the OS itself from process
  - Protection thus enables us to deliver the property of isolation among processes
- **We will focus our exploration on the basic mechanisms needed to virtualize memory, including both OS and hardware support**

프로세스의 고립성 보장.

# Summary

- The virtual memory system is responsible for providing the **illusion of a large, sparse, private address space to programs**
  - The OS, with some hardware help, will take virtual memory references, and turn them into physical addresses, which can be presented to the physical memory



Courtesy of Prof. Jin-Soo Kim @ SNU