

Mechanism: Address Translation



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Virtualization of Memory

시스템 - 메모리 동시성

- **CPU virtualization is done by limited direct execution (LDE)**
 - The program run directly on the hardware through system calls (efficiency)
 - The OS ensures that it maintains control over the hardware (control)
- **Similarly, memory virtualization will attain both efficiency and control while providing the desired virtualization** → address translation.
 - **Efficiency** dictates that we make use of **hardware support** → software는 부족함.
 - **Control** implies that the OS ensures that no application is allowed to access any memory but its own; thus, to **protect** applications from one another
 - Finally, we will need a little more from the VM system; **flexibility** that programs to be able to use their address spaces in whatever they would like
- **How to efficiently and flexibly virtualize memory?**
 - **Hardware-based address translation** (or simply **address translation**) allows the hardware to change the virtual address to a physical address
 - The OS must get involved at key points to set up the hardware
- **The goal of the memory virtualization is to create a beautiful illusion that the program has its own private memory**

Assumption and Example

We will assume that:

- 1) The user's address space must be placed **contiguously** in physical memory
- 2) The size of address space is less than the size of physical memory
- 3) Each address space is exactly the same size

Let's look at a simple example

```
void func() {
    int x = 3000; // thanks, Perry.
    x = x + 3;    // line of code we are interested in
    ...
}
```

↓ Compiler

```
128: movl 0x0(%ebx), %eax    ;load 0+ebx into eax
132: addl $0x03, %eax       ;add 3 to eax register
135: movl %eax, 0x0(%ebx)    ;store eax back to mem
```

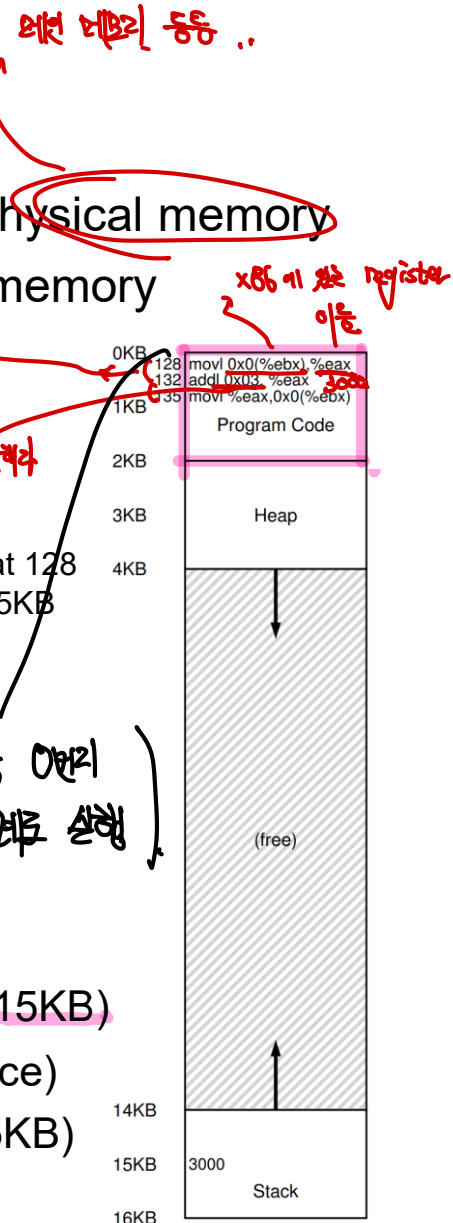
- Instructions are located at 128
- Variable x is located at 15KB

→ 항상 0번지
부터 차례로 실행

– To run the instructions, the following memory accesses:

- 1) Fetch instruction at address 128 & Execute (load from address 15KB)
- 2) Fetch instruction at address 132 & Execute (no memory reference)
- 3) Fetch instruction at address 135 & Execute (store to address 15KB)

3003



Handwritten notes in red ink on graph paper, including the word "Date" and various scribbles.

- **Static (Software-based) relocation**

- OS rewrites each program before loading



Dynamic (Hardware-based) Relocation (1)

- For efficient address translation, the **(base and bounds) scheme** (referred to as **dynamic relocation**) was introduced
 - We need two hardware registers within CPU: **base and bounds (limit) registers**
 - This base-and-bounds pair allows us to place **address space anywhere** in physical memory

- Each program is written and compiled as if it is loaded at **address zero**, and when it starts running


- The **OS** decides where in physical memory it should be loaded and thus **sets the base register** to this value

- Processor translates the address by

physical address = virtual address + base

- Every virtual address must not be greater than bound and negative

$$0 \leq \text{virtual address} < \text{bounds}$$



Dynamic (Hardware-based) Relocation (2)

Let's take a look at the detail

```
128: movl 0x0(%ebx), %eax
```

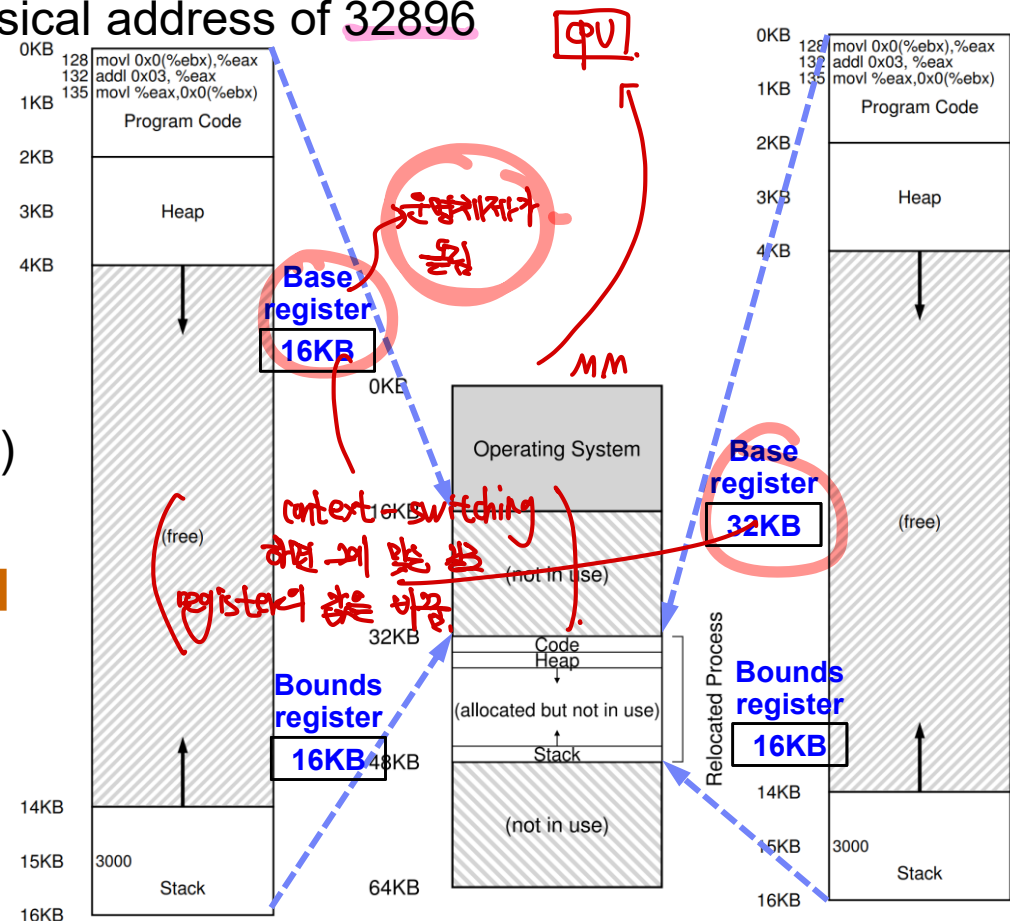


- The PC is set to 128; To fetch this, the processor adds the base register value of 32 KB (32768) to get a physical address of 32896

- The **base register** is set by OS when context switching
- In fact, the CPU first checks that the **memory reference is within bounds** to make sure it is legal; otherwise, it will raise an **exception** (protection)

Such a hardware is called **memory management unit (MMU)** → Me

- located inside the processor



Hardware Support: A Summary

Hardware requirements for dynamic relocation (MMU)

→ CPU 자체가 안됨, OS x

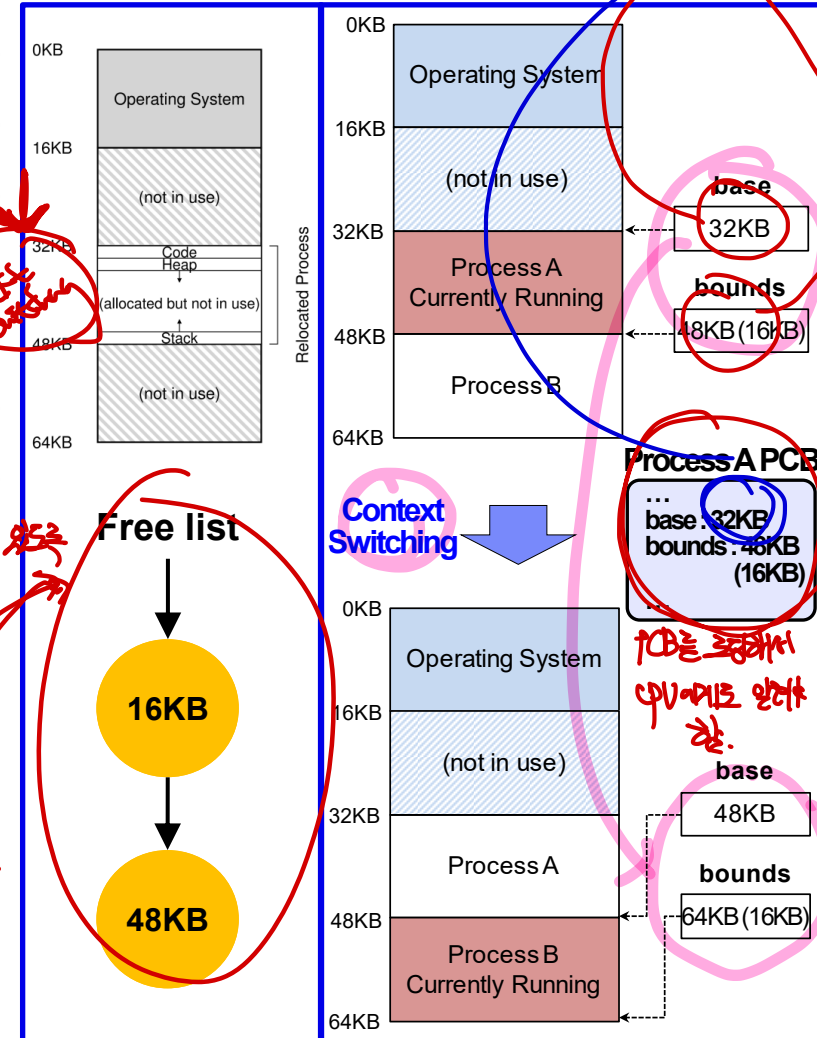
Hardware Requirements	Notes
Privileged mode	Needed to prevent user-mode processes from executing privileged operations
Base/bounds registers	Need pair of registers per CPU to support address translation and bounds checks
Ability to translate virtual addresses and check if within bounds	Circuitry to do translations and check limits; in this case, quite simple
Privileged instruction(s) to update base/bounds	OS must be able to set these values before letting a user program run
Privileged instruction(s) to register exception handlers	OS must be able to tell hardware what code to run if exception occurs
Ability to raise exceptions	When processes try to access privileged instructions or out-of-bounds memory

CPU가 메모리를 관리해줘야 함

OS가 메모리를 관리해줘야 함

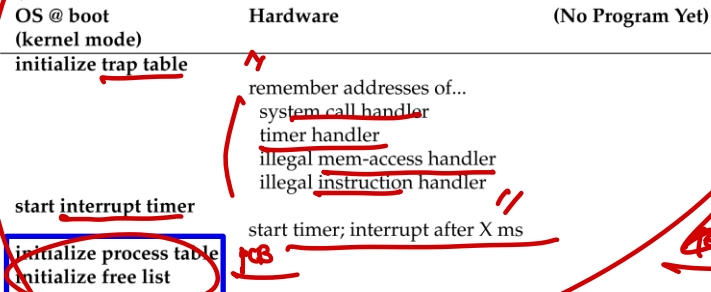
Operating System Issues

OS Requirements	Notes
Memory management	Need to allocate memory for new processes; Reclaim memory from terminated processes; Generally manage memory via free list.
Base/bounds management	Must set base/bounds properly upon context switch
Exception handling	Code to run when exceptions arise; likely action is to terminate offending process

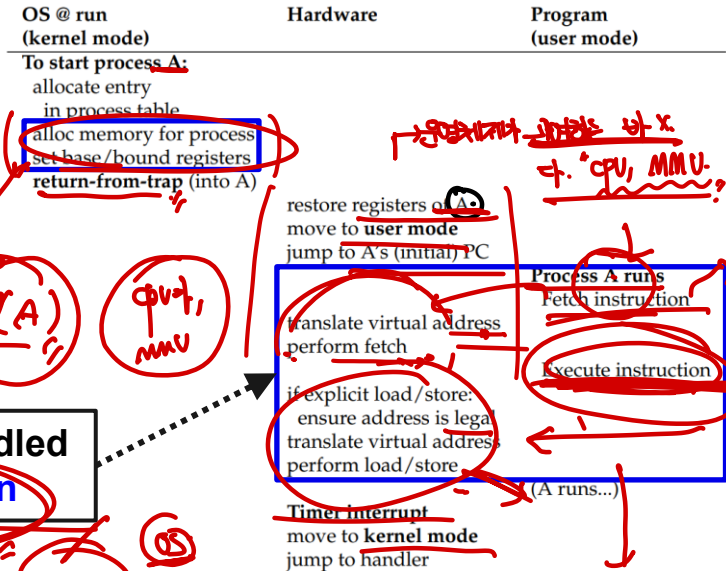


Limited Direct Execution Revisited

Initialization @ Booting

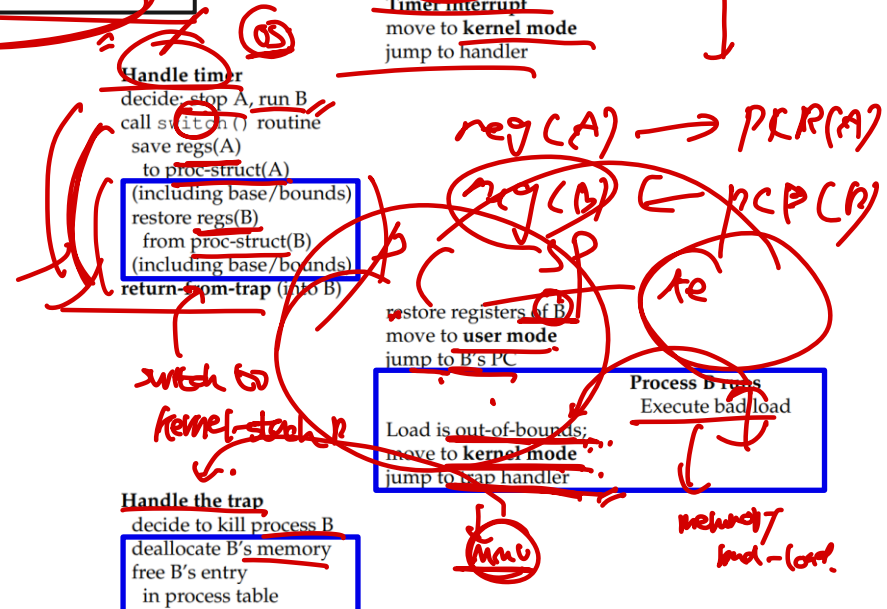
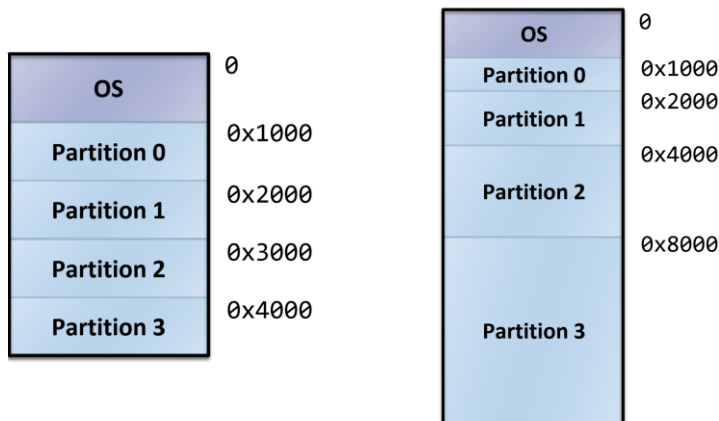


Note how its memory translations are handled by the hardware with **no OS intervention**



Fixed and Variable Partitions

- Partition size needs not be equal



os.

hw.

trap table

system-call handler.

timer handler

illegal mem-access handler

illegal instruct error.

interrupt-dq



start timer



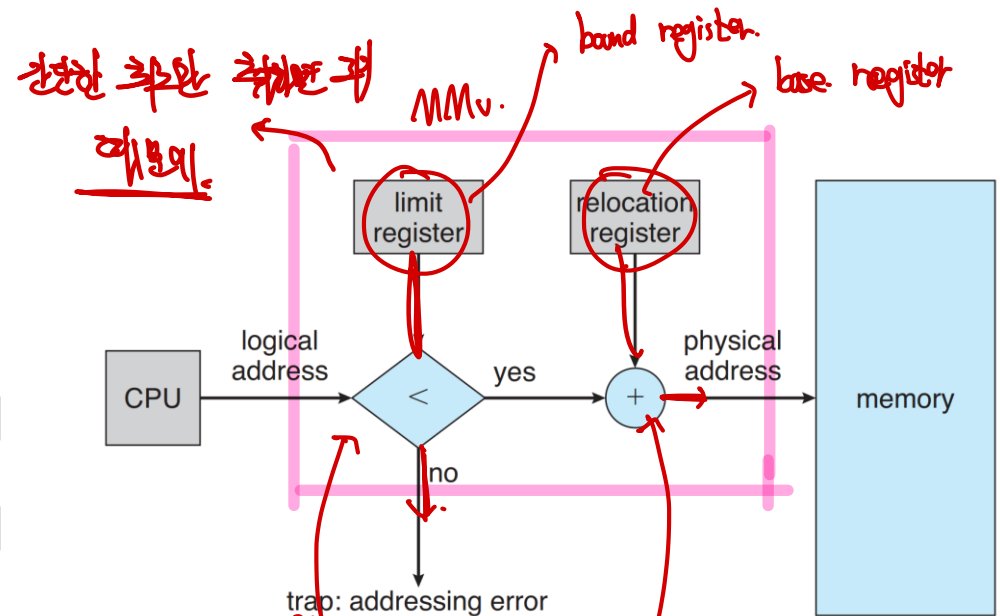
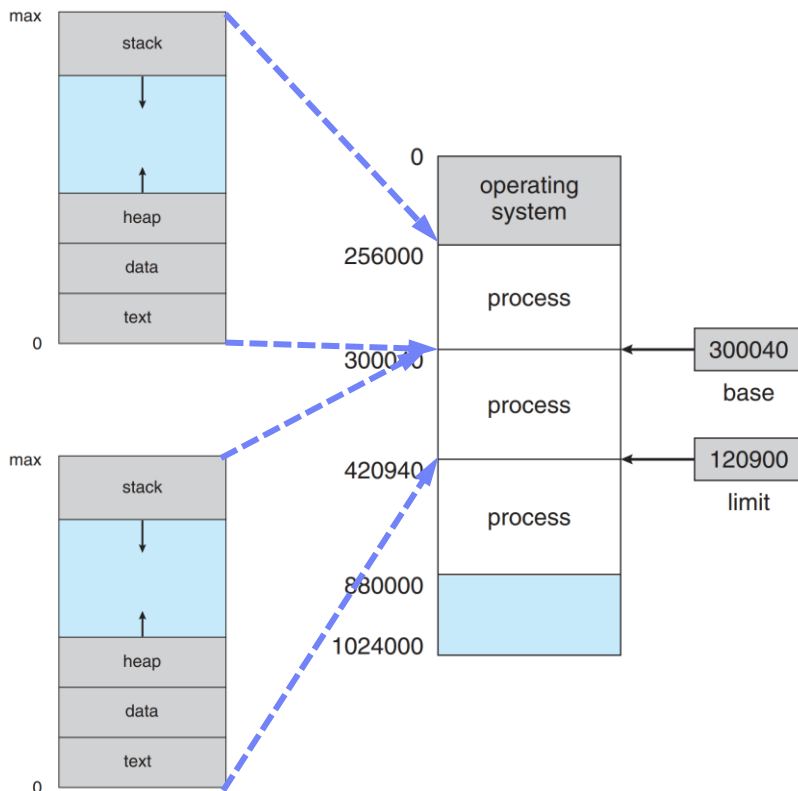
process table.

initializ. free-list.



Summary

- CPU generates a virtual (logical) address and MMU translates it into the physical address
 - MMU includes base and bound registers, which are used to load programs



$0 \leq \text{virtual address} < \text{bounds}$

$\text{physical address} = \text{virtual address} + \text{base}$