

Segmentation



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University


Limitation of Base-Bounds Scheme

- **The base-bounds approach to virtualize memory is wasteful**

- There is a big chunk of “free” space between heap and stack
- The space between stack and heap is not being used by process it still takes up physical memory when relocating
- It is quite hard to run a program when the entire address space does not fit into physical memory
- Thus, **base-bounds is not as flexible** as we would like

- **To solve this issue, segmentation was born**

- Instead of having one base and bounds pair in MMU, let's have **a base and bounds pair per logical segment**
- A segment is just a contiguous portion of the address space of a particular length
- We have three logically-different segments: **code, stack, heap**
- The segmentation allows the OS to place each segment in **different part of physical memory**




Segmentation

Let's look at an example

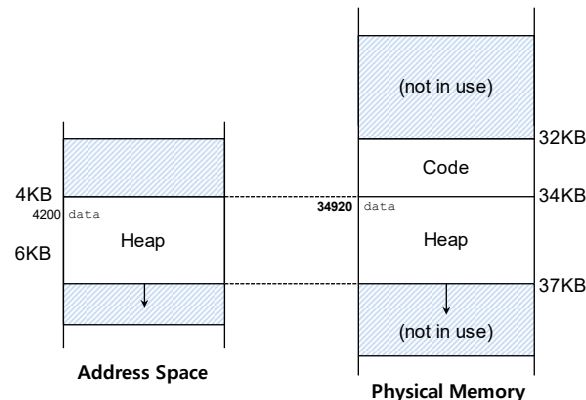
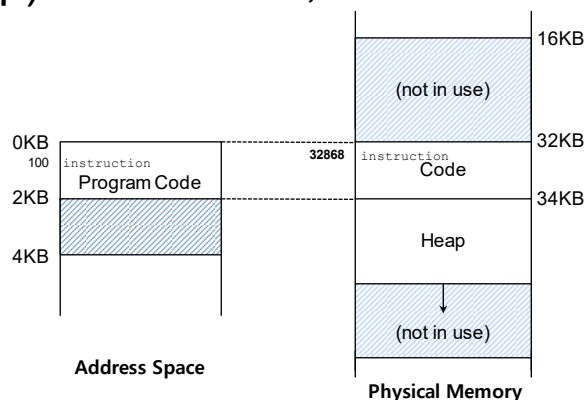
- 64KB physical memory with three segments in it
- Only used memory is allocated space in physical
- MMU has **three base and bounds register pairs**

Segment	Base	Size
Code	32K	2K
Heap	34K	3K
Stack	28K	2K





Address translation examples

- A reference is made to virtual address 100 (code segment):
 $32\text{KB}(32768) + 100 = 32868$ and the address is within bounds
- Virtual address 4200 (heap segment): heap starts at 4 KB (4096) and offset is $4200 - 4096 = 104$; thus $34\text{KB}(34816) + 104 = 34920$ (within bounds)
- OS occurs **segmentation fault** if an illegal address (e.g. >7KB; beyond the end of heap) is reference; the hardware detects that this address is out of bounds



Which Segment Are We Referring To?

- The hardware uses segment registers during translation
 - How does it know the offset into a segment, and to which segment an address refers?
- One common explicit approach is to chop up the address space into segments based on the top few bits of the virtual address



- 00: code segment, 01: heap segment, 11: stack segment, 10: unused
- Virtual address 4200 is encoded: **01_0000_0110_1000** (heap seg. + offset 104)

```
// get top 2 bits of 14-bit VA
Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
// now get offset
Offset = VirtualAddress & OFFSET_MASK
if (Offset >= Bounds[Segment])
    RaiseException(PROTECTION_FAULT)
else
    PhysAddr = Base[Segment] + Offset
    Register = AccessMemory(PhysAddr)
```

- SEG_MASK 0x3000
- SEG_SHIFT 12
- OFFSET_MASK 0xFFFF

- It limits use of virtual address space (2^{14} :16KB \rightarrow 2^{12} :4KB) \rightarrow implicit approach

What About The Stack?

- Stack has one critical difference: it grows backwards to lower

- It starts at 28KB and grows back to 26KB, virtual addresses 16KB to 14KB

- We need a little extra hardware support

- The hardware needs to know which way the segment grows and translates differently


- 1: positive direction, 0: negative direction

Segment	Base	Size (max 4K)	Grows Positive?
Code ₀₀	32K	2K	1
Heap ₀₁	34K	3K	1
Stack ₁₁	28K	2K	0


- Example of virtual address 15KB

- 11_1100_0000_0000 (offset 3KB) → 3KB-4KB=-1KB is actual offset for stack
→ 28KB - 1KB = 27KB is the physical address (within bound, $\text{abs}(-1\text{KB}) < 2\text{KB}$)


실제 3KB를 뺀 뒤 1KB가 남!



Offset in 12bit (4KB) space:
Forward 3KB == Backward 1KB



Positive direction



Negative direction

Address Translation for Stack

■ Example of virtual address 15KB (11_1100_0000_0000)

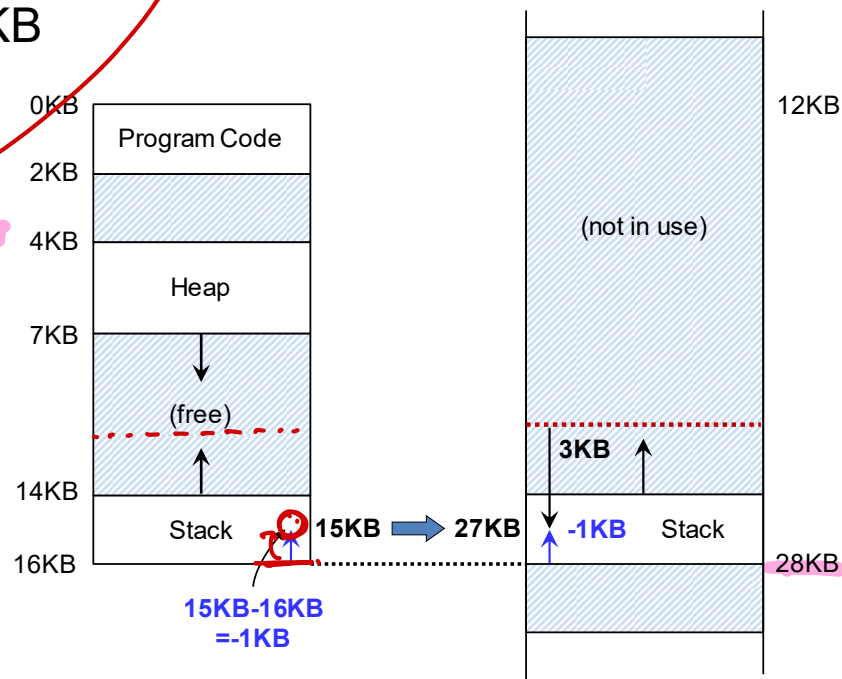
- To obtain an actual offset for backward growing stack, the virtual address needs to be **subtracted by virtual space size** (or **stack's virtual address**)
- Thus, offset = $15\text{KB} - 16\text{KB} = -1\text{KB}$

$$\text{offset (stack)} = \text{virtual address} - \text{virtual space size}$$

- Then, physical address = $28\text{KB} - 1\text{KB} = 27\text{KB}$
- Similarly, this can be obtained by
 $\text{offset (stack)} =$
 $\text{offset (from address)} - \text{max segment size}$
- Thus, offset = $3\text{KB} - 4\text{KB} = -1\text{KB}$

■ Still difficult?

- You take **2's complement** of the offset from address and attach **minus** sign
- e.g.) 1100_0000_0000 (3KB) \rightarrow
 0100_0000_0000 (1KB) $\rightarrow -1\text{KB}$



Support for Sharing

→ segment를 공유



It is useful to share certain segments between address spaces

– Code sharing is common and still in use in systems today

– We need a little extra hardware support:

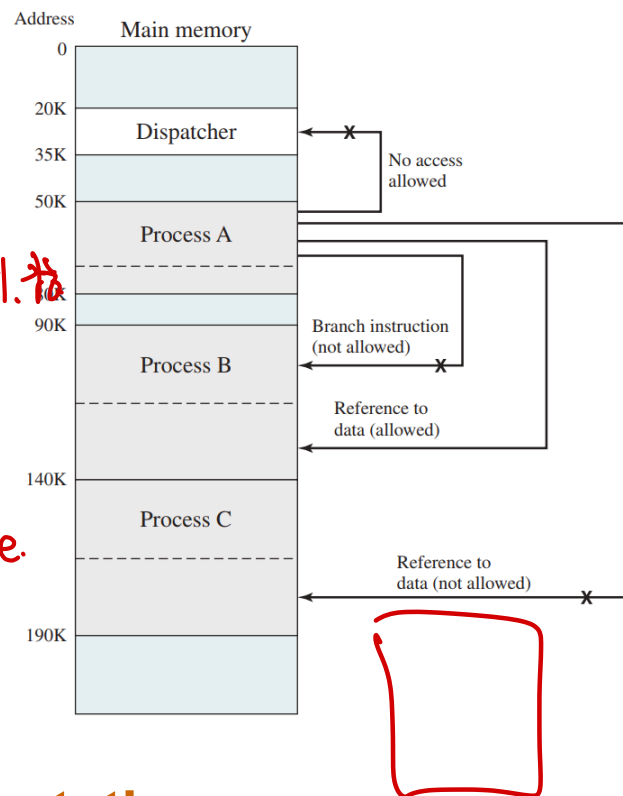
Protection bits per segment to indicate

permission of read, write, and execute. → 플래그로 표시

Segment	Base	Size (max 4K)	Grows Positive?	Protection
Code ₀₀	32K	2K	1	Read-Execute
Heap ₀₁	34K	3K	1	Read-Write
Stack ₁₁	28K	2K	0	Read-Write

→ segment table

– The hardware also has to **check whether a particular access is permissible or not**



Fine-grained and coarse-grained segmentation

– Coarse-grained segmentation: a few relatively large, coarse segments

– Fine-grained segmentation has a large number of smaller segments

– Supporting many segments requires further hardware support: **segment table**

OS Support

coalescing

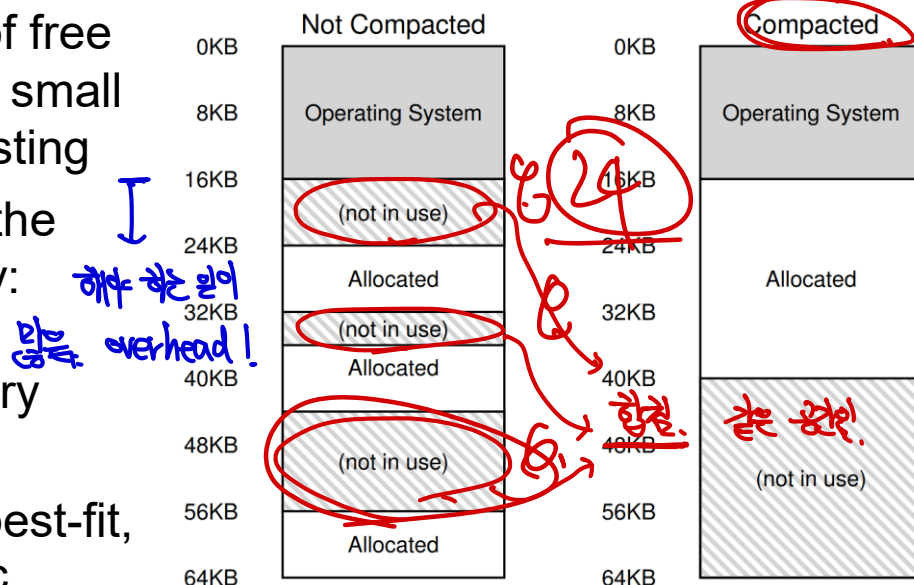
segment register.

Segmentation raises a number of new issues for OS

- What should the OS do on a context switch?
→ The segment registers must be saved and restored
- OS interaction when segments grows (or perhaps shrink)
e.g.) 1) the heap segment need to grow → 2) the memory-allocation library will perform a system call to grow the heap (e.g. brk()) → 3) OS provides more space, updating segment size register to the new bigger size

The critical issue is managing free space in physical memory

- External Fragmentation: little holes of free space in physical memory that is too small to allocate new segment or grow existing
- Solution is compaction: rearranging the exiting segments in physical memory:
 - stop running process
 - copy data to somewhere in memory
 - change segment register value
- New efficient allocation algorithms: best-fit, worst-fit, first-fit, buddy algorithm, etc



Summary

- **Segmentation utilizes a base and bounds register per segment**
 - MMU maintains multiple base and bound register pairs, which will be placed in the memory when the segment table is large

