

Free-Space Management



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

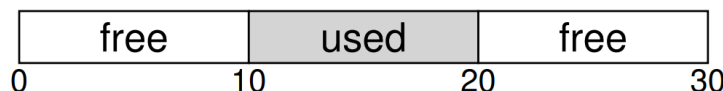
External Fragmentation

- Free-space management becomes more difficult when the free space consists of variable-sized units and this arises in:

- A user-level memory-allocation library (`malloc()` & `free()`)
- an OS managing physical memory when using segmentation for virtual memory

- In either case, the problem is external fragmentation

- The external fragmentation is that the free space get chopped into little pieces of different sizes and is thus fragmented
- Subsequent request may fail because there is no single contiguous space that can satisfy the request, although total free space exceeds the requested size




→ 다른 공간이 불완전 상황

- e.g.) total space in above figure is 20 bytes; unfortunately, it is fragmented into two chunks of size 10 bytes each
 → As a result, a request for 15 bytes will fail despite 20 bytes free


Assumptions

- We assume a basic interface of `malloc()` & `free()`
 - The space that the library manages is known historically as `heap`
 - The generic data structure to manage free space in heap is a kind of `free list`. *linked-list 구조*
- We assume that our primary concern is `external fragmentation`
 - Allocators could also have the problem of `internal fragmentation` *(차이)*
 - Any unasked space in an allocated chunk is considered internal fragmentation; the `waste occurs inside the allocated unit`; which is another space waste
- We assume that once memory is handed out to a client, it `cannot be relocated to another location memory`
 - No compaction of free space is possible
- We assume that allocator manages a `contiguous region of bytes`
 - The region is a single fixed size *→ 변하지 않음*




Low-Level Mechanisms: Splitting

- A free list contains a set of elements that describe the free space remaining in the heap; let's consider a 30-byte heap





- A request for ≤ 10 bytes will succeed; otherwise, it will fail (returning NULL)
- What happens if the request is for something < 10 bytes?
 - The allocator will perform **splitting**; it will find a free chunk and split it into "two"
 - The first chunk will return to the caller and the second will remain on the list
 - If a request for 1 byte were made, the call to `malloc()` would return the address 20, and the list will be:



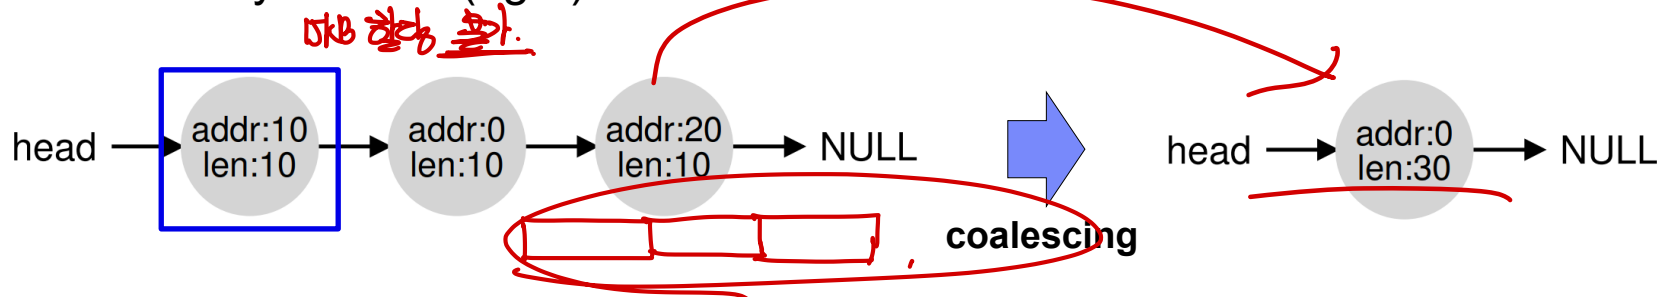
- As can be seen, the list basically stays **intact**; the only change is that the free region now starts at 21, instead of 20, and the length is now 9

Low-Level Mechanisms: Coalescing

- A corollary mechanism is known as **coalescing** of free space
 - Coalescing is to merge returning a free chunk with existing chunks into a larger single free chunk if addresses of them are nearby.
 - Let's consider the same example



- What happens if an application calls **free()** to free the chunk?
 - If we simply add this free space back into the list, you will see the list with three chunks (left)
 - To avoid this problem, the allocator **coalesces free space** when a chunk of memory is freed (right)



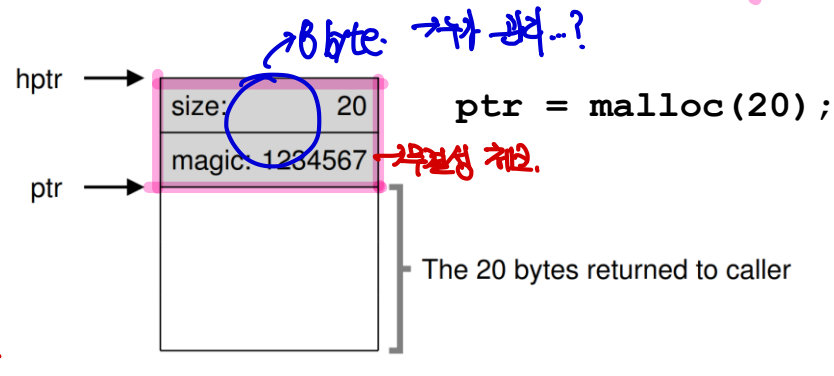
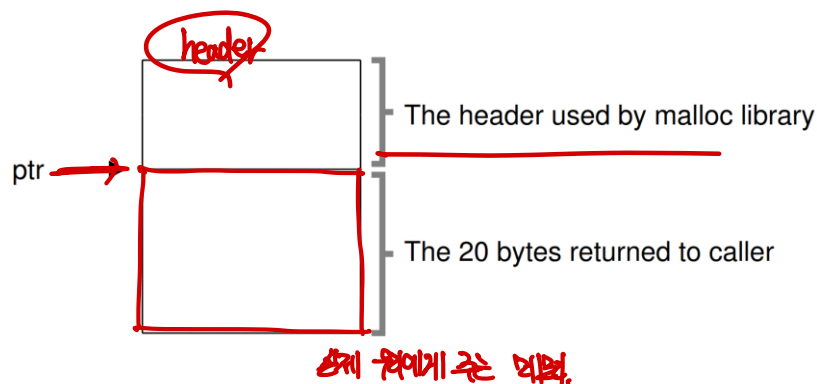
Tracking the Size of Allocated Regions

*int *p = malloc(10);
free(p) # 10 byte free*

누가 size 정보를 찾고 있겠어?

■ Inference to free (void *ptr) does not take size parameter

- The library can quickly determine the size of the region of memory being freed
- To accomplish this, allocators store a little bit of extra information in header



- The header contains the size of the allocated region and a magic number to offer additional integrity checking, and other information

```
typedef struct {
    int size;
    int magic;
} header_t;

void free(void *ptr) {
    header_t *hptr = (header_t *) ptr - 1;
    ...
    assert(hptr->magic == 1234567)
```

- Therefore, when a user request (N) bytes of memory, the library searches for a free chunk of size (N + header size)

size 정보는 header에도 추가 할당.

Embedding a Free List: Initialization

How do we build such a list inside the free space itself?


- Just `malloc()`? → No! you can't do this

Assume we have a 4096-byte chunk of memory to manage

- To manage this as a free list, we first have to initialize the list
- The list should have only one entry of size 4096 (minus the header size)
- We assume that the heap is built via a call to the system call `mmap()`

```
typedef struct __node_t {
    int size;
    struct __node_t *next;
} node_t;

// mmap() returns a pointer to a chunk of free space.
node_t *head = mmap(NULL, 4096, PROT_READ|PROT_WRITE,
                    MAP_ANON|MAP_PRIVATE, -1, 0);
head->size = 4096 - sizeof(node_t);
head->next = NULL;
```




- Once initialization, the actual heap size is 4088 bytes and the head pointer contains the beginning address of this range (e.g. virtual address 16KB)

Embedding a Free List: Allocation

head → 4096 → Null

- If a chunk of memory is request, the library will **(first find a chunk that is large enough)** to accommodate the request

- The chunk will be chosen as we have only one
- Then, it will be **split into two**: one chunk big enough to service the request and the remaining free chunk
- The library **allocates 108 bytes** (8 bytes header) and returns a pointer **ptr**
- The size of free space becomes 3980 bytes





After Two more 100 bytes chunks are allocated

head → 3980

108 - 116 = 3980

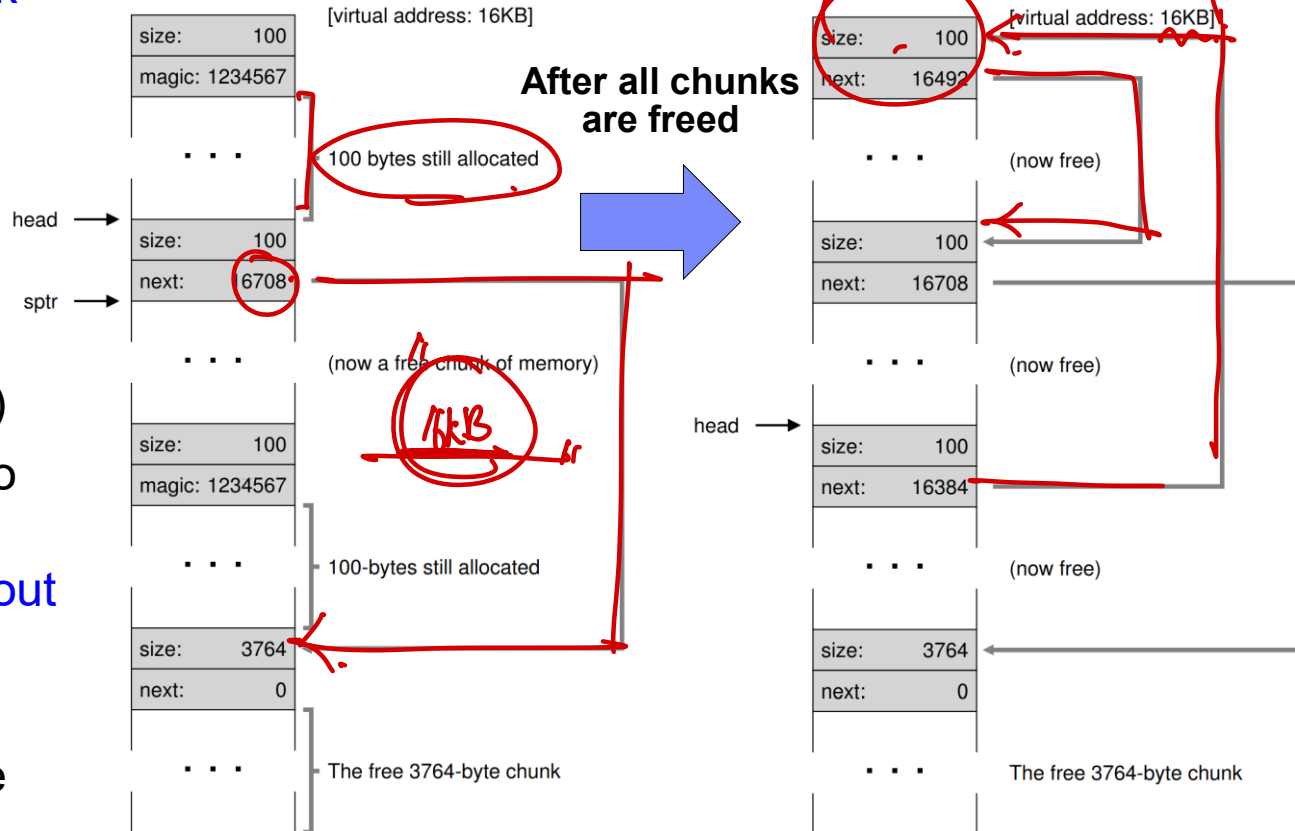
3980



Embedding a Free List: Free

What happens when the program returns memory via `free()`?

- The application returns the middle chunk of memory by calling `free(16500)`. 16500 is derived by $16384 + 108 + 8$ (`sptr`)
- The library immediately figures out the size of the free region, and then adds the free chunk back onto the free list
- Then, we have a list that starts with a small free chunk (100 bytes, `head`) and a large free chunk (3764 bytes)
- Freeing the last two chunks will lead to fragmentation without coalescing
- Merge neighboring chunks to coalesce



Basic Strategies for Managing Free Space

Strategy	Procedure
Best Fit	<ol style="list-style-type: none"> 1) search through the free list and find chunks of free memory that are <u>as big or bigger than the requested size</u> 2) return the one that is <u>the smallest in that group of candidates</u>; this is the so called best-fit chunk
Worst Fit	<ol style="list-style-type: none"> 1) find the <u>largest chunk</u> and return the requested amount 2) keep the remaining (large) chunk on the free list
First Fit	<ol style="list-style-type: none"> 1) find the <u>first block that is big enough</u> 2) return the requested amount to the user
Next Fit	<ol style="list-style-type: none"> 1) keep an <u>extra pointer</u> to the location within the list where one was looking last 2) return the <u>next big enough chunk of the pointer</u> on the free list



Other Approaches: Segregated Lists

■ The basic idea of the segregated list is that:

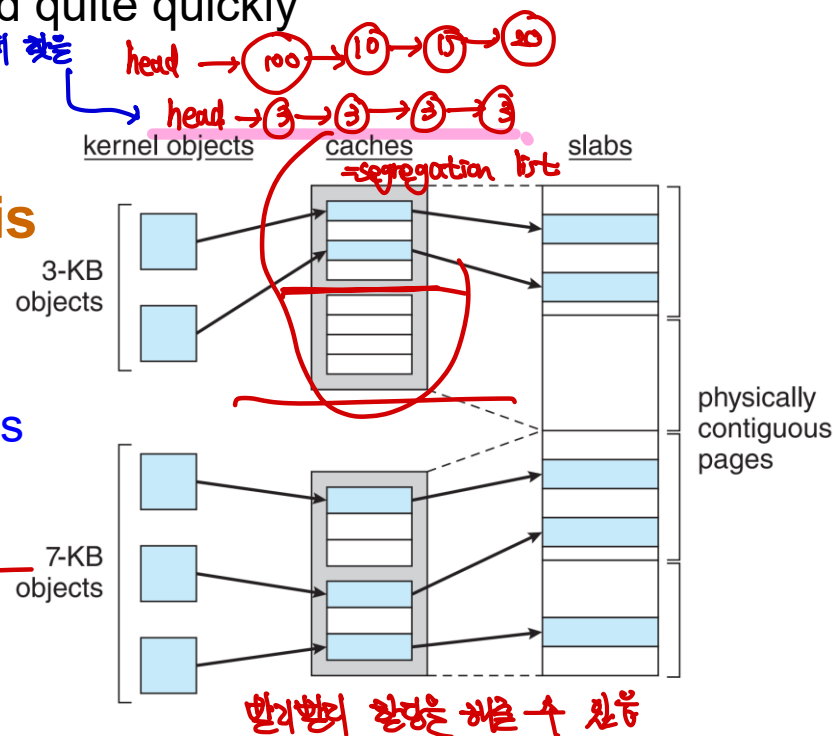
- If a particular application has one or a few popular (fixed) sized request, it keeps a separate list to manage objects of that size

■ The benefit of this approach are obvious

- Fragmentation is much less of a concern
- Allocation and free requests can be served quite quickly

■ Slab allocator is an advanced form of segregated list adopted in Solaris

- A slab is one or more physically contiguous pages in memory
- This scheme allocates some object caches for kernel object that are likely to be requested frequently (e.g., lock, i-node)
- Then, the object caches are each segregated free list



Other Approaches: Buddy Allocation

■ The binary buddy allocator makes coalescing simple

- In this system, free memory is first conceptually thought of as of one big space of size of 2^N 4096 = 2^{12}
- Each memory request, the search for free space recursively divides free space by two until a big enough block (best fit) to accommodate the request is found
- e.g.) 7KB request of 64KB memory: 64KB \rightarrow 32KB \rightarrow 16KB \rightarrow 8KB (allocate)
- This scheme suffers from internal fragmentation

■ The beauty of buddy allocation is found when the block is freed

- When 8KB block is freed, the allocator checks if “buddy” 8KB is free
- If so, it coalesces the two blocks into 16KB and checks again if buddy of 16KB is free and coalesces if it is free
- This recursive coalescing process continues up the tree

