

Paging: Introduction



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Paging: Concept

OS takes one of two approaches memory space-management:

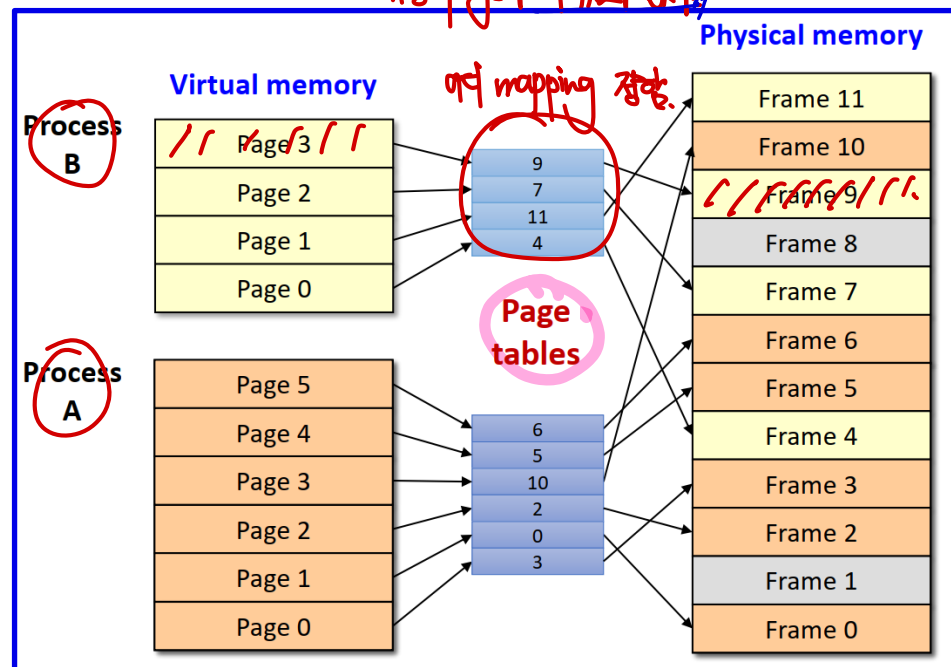
- 1) Chopping the space up into **variable-sized** pieces (Segmentation)
- 2) Chopping the space up into **fixed-sized** pieces (Paging)

Segmentation has inherent difficulties

- The space itself can become **fragmented** when dividing a space into different-size chunks
- Thus, the allocation becomes more challenging over time

Paging divides the space into fixed-sized units

- Each of the unit is called a **page**
- We view physical memory as an array of fixed-sized slots called **page frames**
- Each frame can contain a single virtual-memory page



Courtesy of Prof. Jin-Soo Kim @ SNU

external-frag 문제 x.


Simple Example

이러한 segmentation 방식 → (virtual, physical) = 3가지

(virtual)

Let's consider an example of a tiny address space:

- Only 64 bytes total in size, with four 16-byte pages (virtual pages 0, 1, 2, 3)
- Physical memory consists of eight fixed-sized slots, making 128-byte memory
- As can be seen, the pages of the virtual address space have been placed at different locations throughout physical memory



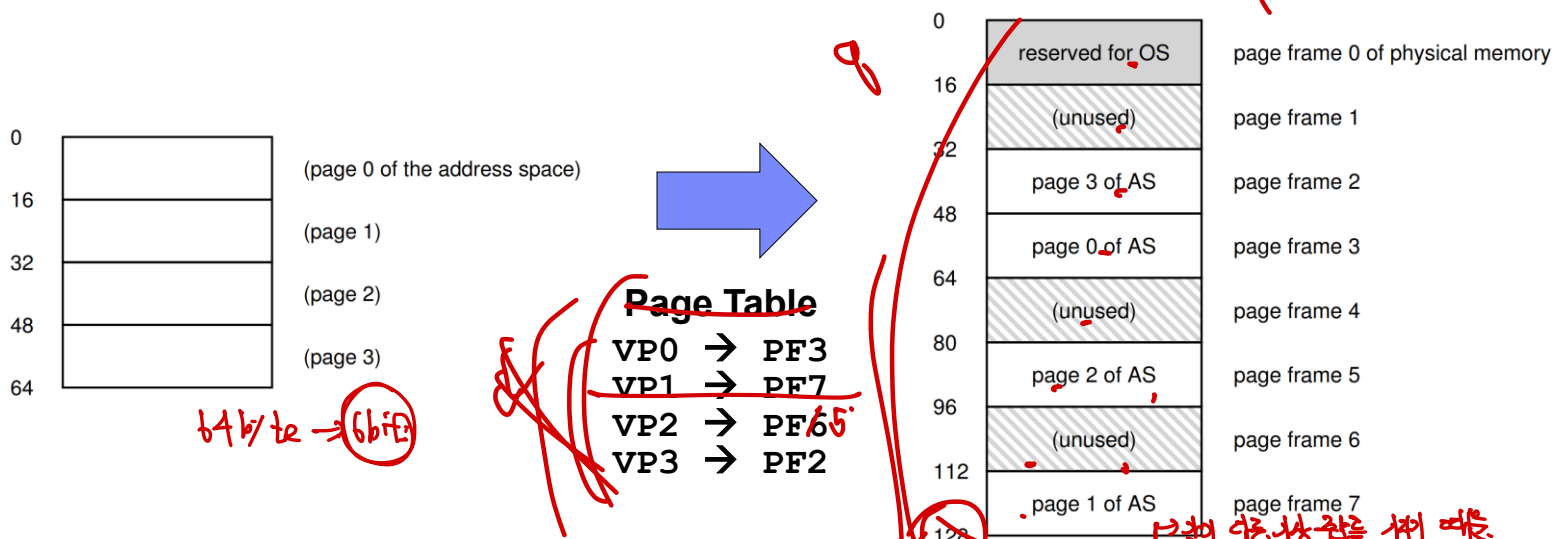
64-byte address space in 128-byte physical memory

Paging has a number of advantages over previous approaches

- Flexibility**: supporting (abstraction) of address space effectively; we don't need to know the growing directions for heap and stack
- Simplicity**: ease of free-space management; simple free list (same sized units)

Paging Overview

- To record where each virtual page is placed in physical memory, OS keeps a **per-process data structure known as a page table**.
 - The major role of the page table is to store address translations for each of the virtual pages of the address space (mapping from virtual to physical)
 - The page table for the previous example would have the following four entries: (Virtual Page 0 → Physical Frame 3), (VP 1 → PF 7), (VP 2 → PF 5), (VP 3 → PF 2)



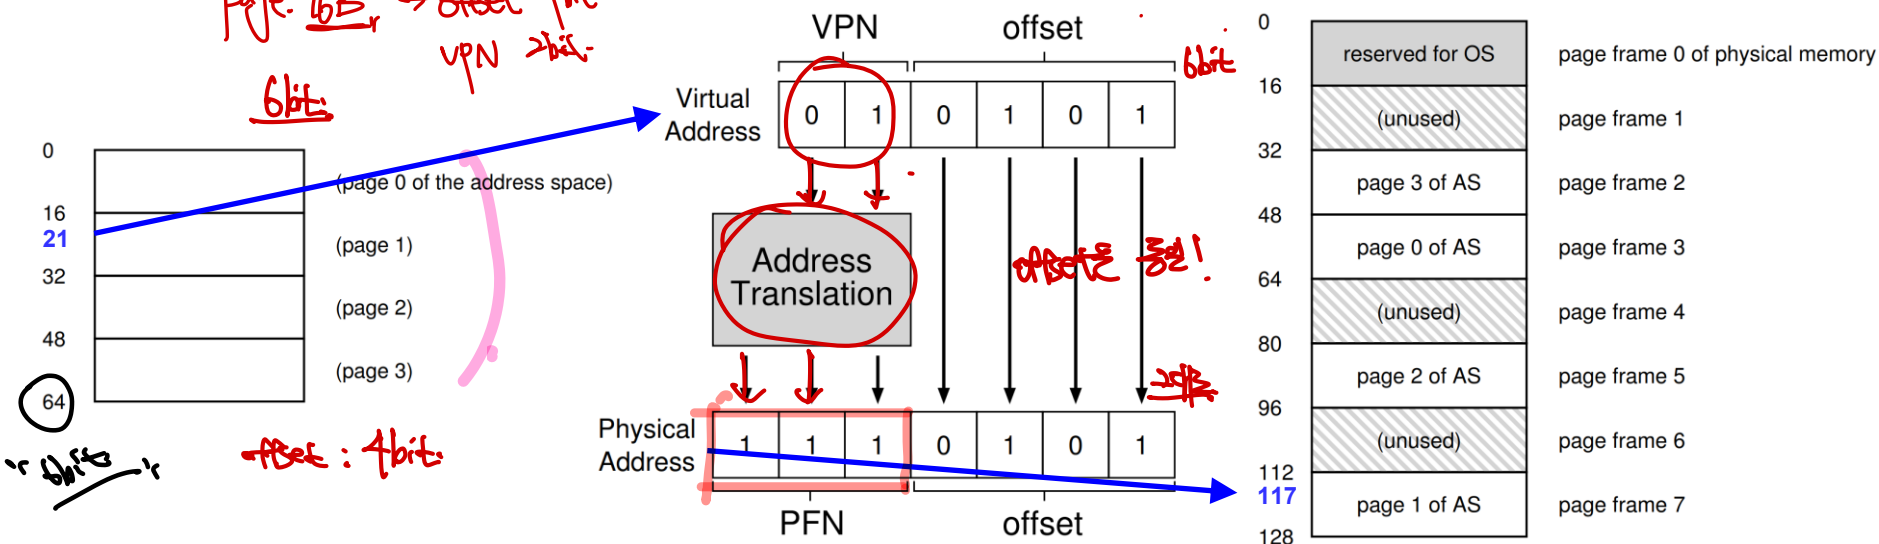
- Remember that this page table is a **per-process data structure**.
 - If another process were to run for the example, OS have to manage a different page table for it, as its virtual pages obviously map to different physical pages

Address Translation



- To translate virtual address, we first split it into two components: **virtual page number (VPN)** and **offset within the page**
 - In the example, we need 6-bit address ($2^6=64$); 4-bit offset ($2^4=16$ -byte page) and 2-bit VPN (4 pages: each page size is 16 bytes in a 64-byte address space)
 - To access the virtual address 21 (01_0101), VPN of 01 is translated into **physical page frame number (PFN)** of 7 (111); replacing VPN with PFN
 - PFN is referred to as **physical page number (PPN)**, Usually, # of VPN \geq # of PFN
 - Note the offset stays the same (i.e., it is not translated)

page: 16B → offset: 4bit
VPN: 2bit



Where Are Page Table Stored?

- Page tables can get terribly large, much bigger than the small segment table or base/bounds pair we discussed earlier

2^{20} 개의 translation $\approx 1\text{MB}$
= process 단.

- 32-bit address space with 4KB pages \rightarrow 20-bit VPN and 12-bit offset ($2^{12}=4\text{KB}$)
- A 20-bit VPN implies that there are 2^{20} translations that the OS would have to manage for each process
- If each page table entry (PTE) in the page table needs 4 bytes to hold address translation and useful stuffs, total 4MB ($4\text{B} \times 2^{20}$) is required for each page table

2^{20} 개 \rightarrow (4MB)

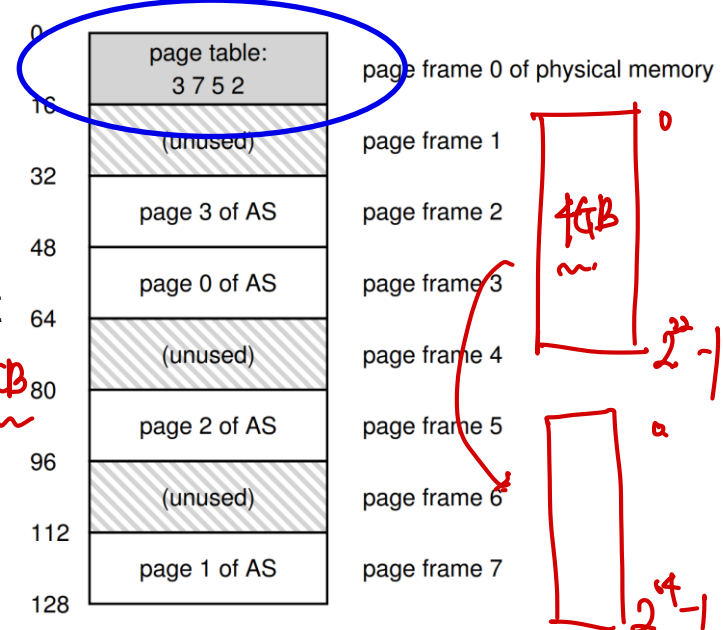
4B \times $2^{20} = 4\text{MB}$

- What if 100 processes is running? \rightarrow 400MB!
- What if 64-bit address space? \rightarrow !!

- Since page tables are so big, the MMU cannot hold them in it but store the tables in memory
- Usually stored in PCB in kernel space

산재한 메모리에
서장하는 것이 x.

OS가 관리한다.
다스리게 swap 되기도 한다.



- map virtual
- index 가 VPN 에기 때문에
하당 정호는 필드 x.
- at index to find
- indexing
- 하당 정호 2

입출 장치 휘발성.

→ brought into memory
(element)

"4byte" → several bits.

5	4	3	2	1	0
A	PCD	PWT	U/S	R/W	P

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
<div> <div>PFN</div> <div>실제 HW mem 어 역.</div> </div>																									G	PAT	D	A	PCD	PWT	U/S	R/W	P

7

Paging: Also Too Slow

- To access memory, the hardware must know where the page table is for the currently-running process.

- The page-table base register (PTBR) contains the physical address of the starting location of the page table
- Once the PTE's address is calculated, the CPU looks up the PFN, which is used to generate the corresponding physical address with offset

2번의 메모리 access

VPN → page table
index → PTE

실제 메모리
(여기 처리)

```
// Extract the VPN from the virtual address
VPN = (VirtualAddress & VPN_MASK) >> SHIFT
```

```
// Form the address of the page-table entry (PTE)
PTEAddr = PTBR + (VPN * sizeof PTE)
```

```
// Fetch the PTE
PTE = AccessMemory(PTEAddr)
```

```
// Check if process can access the page
if (PTE.Valid == False)
    RaiseException(SEGMENTATION_FAULT)
else if (CanAccess(PTE.ProtectBits) == False)
    RaiseException(PROTECTION_FAULT)
else
```

```
// Access is OK: form physical address and fetch it
offset = VirtualAddress & OFFSET_MASK
PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
Register = AccessMemory(PhysAddr)
```

Hw → "pagetable"

MMU → 관찰한다.

→ $PTBR + VPN \times \text{sizeof}(PTE)$ → MMU에 있는 address를 → Valid한 Bit를 → 다 통과하면
va를 처리하고 vpn을 가져옴



VPN



메모리 참조 2.

- For every memory reference, paging requires us to perform one extra memory reference to first fetch the translation from page table → a lot of extra work!

- ```
int array[1000];
..
for (i = 0; i < 1000; i++)
 array[i] = 0;
```


## Compile

```
1024 movl $0x0, (%edi, %eax, 4)
1028 incl %eax
1032 cmpl $0x03e8, %eax
1036 ine 0x1024
```

`%edi`: base address of array

```
%eax: i
```

-



1024

movl \$0x0, (%edi, %eax, 4)

∴ 인스트럭션을 PT[]를 참조해서 가져오

fetch를 해준다.

그리고, 바이트의 구성을 실제 주소로 매핑해

준다. 즉, 1 바이트 주소가 0

은 디바이스 fetch.

4바이트

+3x2바이트

= 10바이트

① PT[] VPN

② instruction fetch

③ PT[39] VPN 39

④ mem[%edi + 4x%eax]에 data fetch.

# Summary

## ■ Paging divides the space into fixed-sized units

- To record where each virtual page is placed in physical memory, OS keeps a per-process data structure known as a page table, which maps VPN to PFN
- OS indexes the array by the VPN, and looks up the PTE at that index to find the corresponding PFN

