


# Paging: Faster Translation (TLBs)



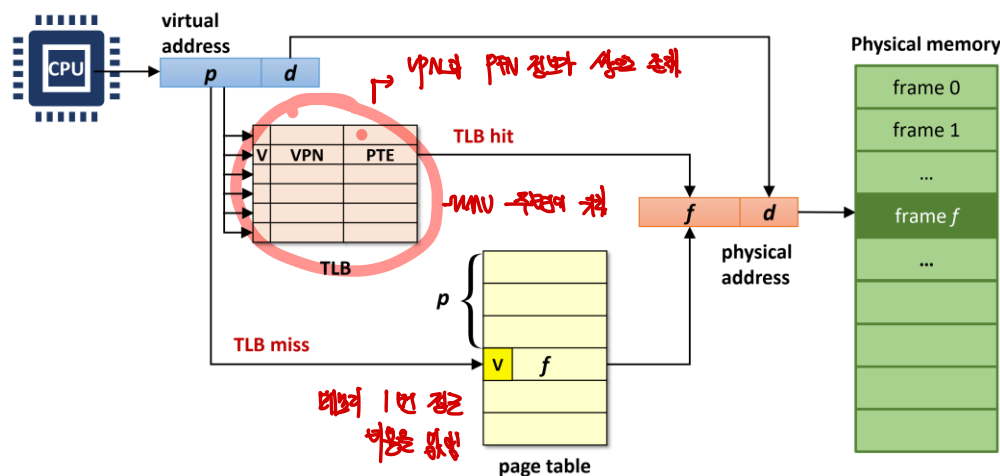

**Prof. Yongtae Kim**

Computer Science and Engineering  
Kyungpook National University



# Translation Lookaside Buffer (TLB): Concept

- Paging logically requires an extra memory lookup for each VA**
  - Going to memory for translation information before every instruction fetch or explicit load or store is prohibitively slow
  - To speed up the address translation, OS needs the hardware's help
- Translation lookaside buffer (TLB) expedites address translation**
  - A TLB, part of the chip's MMU, is simply a hardware cache of popular virtual-to-physical address translations; better name would be **address-translation cache**
  - TLB makes virtual memory possible as it affects the performance tremendously



Courtesy of Prof. Jin-Soo Kim @ SNU

# TLB Basic Algorithm

- Translation using hardware-managed TLB with linear page table

- TLB hit: 1) memory access
- TLB miss: 1) memory access for PTE → update TLB  
2) retry → TLB hit → memory access
- TLB, like all caches, is build on the premise that in common case, translations are found in cache

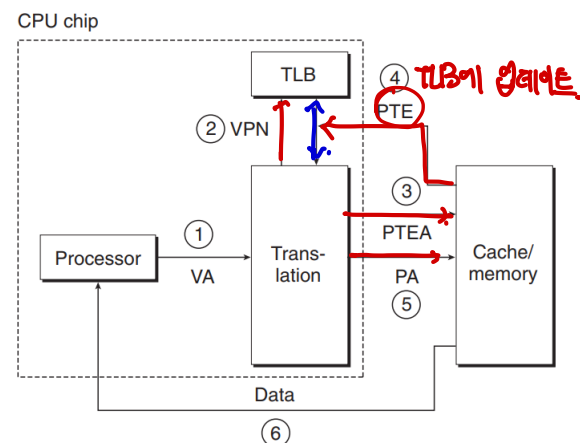
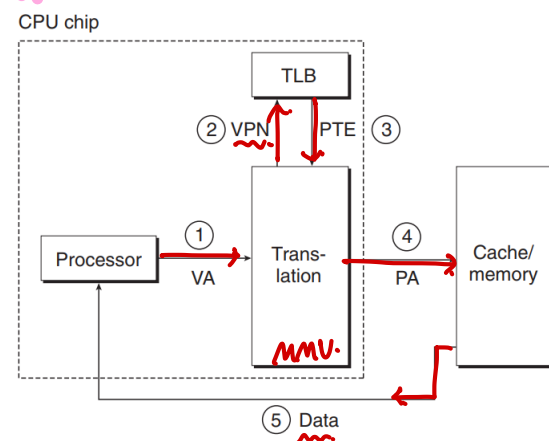
Typically, TLB hit rate  $\geq 99\%$

```

VPN = (VirtualAddress & VPN_MASK) >> SHIFT
(Success, TlbEntry) = TLB_Lookup(VPN)
if (Success == True) // TLB Hit
    if (CanAccess(TlbEntry.ProtectBits) == True)
        Offset = VirtualAddress & OFFSET_MASK
        PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
        Register = AccessMemory(PhysAddr)
    else
        RaiseException(PROTECTION_FAULT)
else // TLB Miss
    PTEAddr = PTBR + (VPN * sizeof(PTE))
    PTE = AccessMemory(PTEAddr)
    if (PTE.Valid == False)
        RaiseException(SEGMENTATION_FAULT)
    else if (CanAccess(PTE.ProtectBits) == False)
        RaiseException(PROTECTION_FAULT)
    else
        TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
        RetryInstruction()
  
```

TLB hit

TLB miss



# Example: Accessing an Array

int a[100]

## Let's consider an example with assumptions

- Array of 10 4-byte integers in memory, starting at VA 100
- Small 8-bit virtual address space with 16-byte pages (4-bit VPN & 4-bit offset)
- Only consider memory access for array

Memory access behavior:

1) access a[0] → TLB miss → TLB update (VPN 6)

2) access a[1], a[2] → TLB hit

3) access a[3] → TLB miss → TLB update (VPN 7)

4) access a[4], a[5], a[6] → TLB hit

5) access a[7] → TLB miss → TLB update (VPN 8)

6) access a[8], a[9], a[10] → TLB hit

TLB hit rate: 7 hit of total 10 access = 70%

10개의 access 16바이트 70%가 히트함.

일단 100% page의 크기가 4바이트이면 100% 히트함

**Spatial locality**: elements of the array are close to one to another in space (page size ↑ → TLB hit rate ↑)

**Temporal locality**: quick re-referencing of memory item in time

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

16/160  
a[0]

page의 크기가 4바이트

VPN

	Offset
	00 04 08 12 16
VPN = 00	
VPN = 01	
VPN = 02	
VPN = 03	
VPN = 04	
VPN = 05	
VPN = 06	a[0] a[1] a[2]
VPN = 07	a[3] a[4] a[5] a[6]
VPN = 08	a[7] a[8] a[9]
VPN = 09	
VPN = 10	
VPN = 11	
VPN = 12	
VPN = 13	
VPN = 14	
VPN = 15	

page size ↑ → TLB hit rate ↑

# Who Handles the TLB Miss?

flexible (x)

- Either **hardware (CPU)** or **software (OS)** can handle the **TLB miss**
- Earlier CISC (x86) has **hardware-managed TLB**
  - The hardware has to know the location of **page table** and its **format**
  - The hardware 1) walks the **page table**, 2) finds the **correct PTE**, 3) extracts the **desired translation**, 4) **update TLB** with the translation, 5) **retry the instruction**
- Newer RISC** has **software-managed TLB**
  - On a TLB miss, the hardware simply raises **an exception** and then: **trap**
  - The hardware 1) pauses execution, 2) switch to **kernel mode**, 3) jumps to **trap handler** → (OS to process TLB miss) → 4) **retries the instruction** (will TLB hit)

**HW TLB**

```

VPN = (VirtualAddress & VPN_MASK) >> SHIFT
(Success, TlbEntry) = TLB_Lookup(VPN)
if (Success == True)    // TLB Hit
    if (CanAccess(TlbEntry.ProtectBits) == True)
        Offset = VirtualAddress & OFFSET_MASK
        PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
        Register = AccessMemory(PhysAddr)
    else
        RaiseException(PROTECTION_FAULT)
else
    // TLB Miss
    PTEAddr = PTBR + (VPN * sizeof(PTE))
    PTE = AccessMemory(PTEAddr)
    if (PTE.Valid == False)
        RaiseException(SEGMENTATION_FAULT)
    else if (CanAccess(PTE.ProtectBits) == False)
        RaiseException(PROTECTION_FAULT)
    else
        TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
        RetryInstruction()
  
```

What hardware do  
on TLB miss

**SW TLB**


```

VPN = (VirtualAddress & VPN_MASK) >> SHIFT
(Success, TlbEntry) = TLB_Lookup(VPN)
if (Success == True)    // TLB Hit
    if (CanAccess(TlbEntry.ProtectBits) == True)
        Offset = VirtualAddress & OFFSET_MASK
        PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
        Register = AccessMemory(PhysAddr)
    else
        RaiseException(PROTECTION_FAULT)
else
    // TLB Miss
    RaiseException(TLB_MISS)
  
```

**TLB miss-handling trap**


- ※ This trap is different from the other trap (e.g. system call) in that the hardware must **resume execution at the instruction that caused the trap**, not the next instruction

# TLB Contents: What's in There?



- 스택이 되지 않는 운영체제 개념이 옳
- What if TLB miss when OS is running TLB miss-handling code?**
    - (Original) TLB miss → TLB miss handling → TLB miss → TLB miss handling → ...
    - Two simple solutions: → 항상 풀려 주신
      - Keeping TLB miss handler in physical memory (directly access to handler code)
      - Reserving some entries in TLB for permanently-valid translations and use some of those slots for handler code itself; these wired translations always hit the TLB
- The primary advantage of software approach is flexibility**
  - OS can use any data structure to implement page table without hardware change
  - Another advantage is simplicity; hardware doesn't do much on a miss (see codes)
- Typically, TLB is (fully associative) with 32, 64, or 128 entries**
  - The hardware search the entire TLB in parallel to find the desired translation
  - TLB entries would be like the following format:
 

VPN	PFN	other bits
  - Other bits usually contains:
    - valid bit to indicate whether the entry has a valid translation or not
    - protection bits to determine how a page can be accessed (**rwX**)
    - address-space identifier, dirty bit, etc





# TLB Issue: Context Switch & Replacement Policy

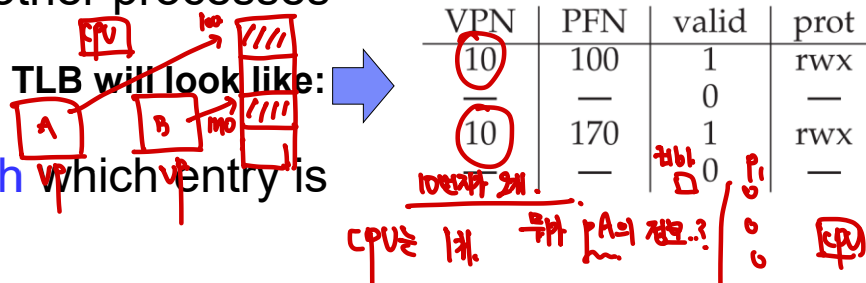
→ TLB에서 주어진 주소에 맞는 entry는 다시 찾아줘야 함.

## With TLBs, some new issues arise when context switching

- TLB contains virtual-to-physical translations that are only valid for the currently running process, not meaningful for other processes

- e.g.) Process A: VPN 10 → PFN 100,  
Process B: VPN 10 → PFN 170

Then, the hardware cannot distinguish which entry is meant for which process



## One solution is to simply flush the TLB on context switches.

- The flush operation simply sets all valid bits to 0; effectively clearing TLB contents
- There is cost: the flush incurs high TLB misses after context switches

## Some hardware can have address space identifier (ASID) in TLB

- ASID makes TLB be shared across context switches
- OS must, on a context switch, set a privileged register to the ASID of the current process

VPN	PFN	valid	prot	ASID
10	100	1	rwX	1
—	—	0	—	—
10	170	1	rwX	2
—	—	0	—	—

Handwritten notes: TLB를 새로 채우기 (Refill TLB), context switch할 때 프로세스 privileged register에 ASID를 (When context switching, set ASID in process's privileged register)

Handwritten notes: TLB에 실려 있는 entry가 (Entry stored in TLB), ASID가 바뀌면 (When ASID changes), TLB를 새로 채워야 함 (Need to refill TLB), ASID를 관리하는 register가 (Register managing ASID)

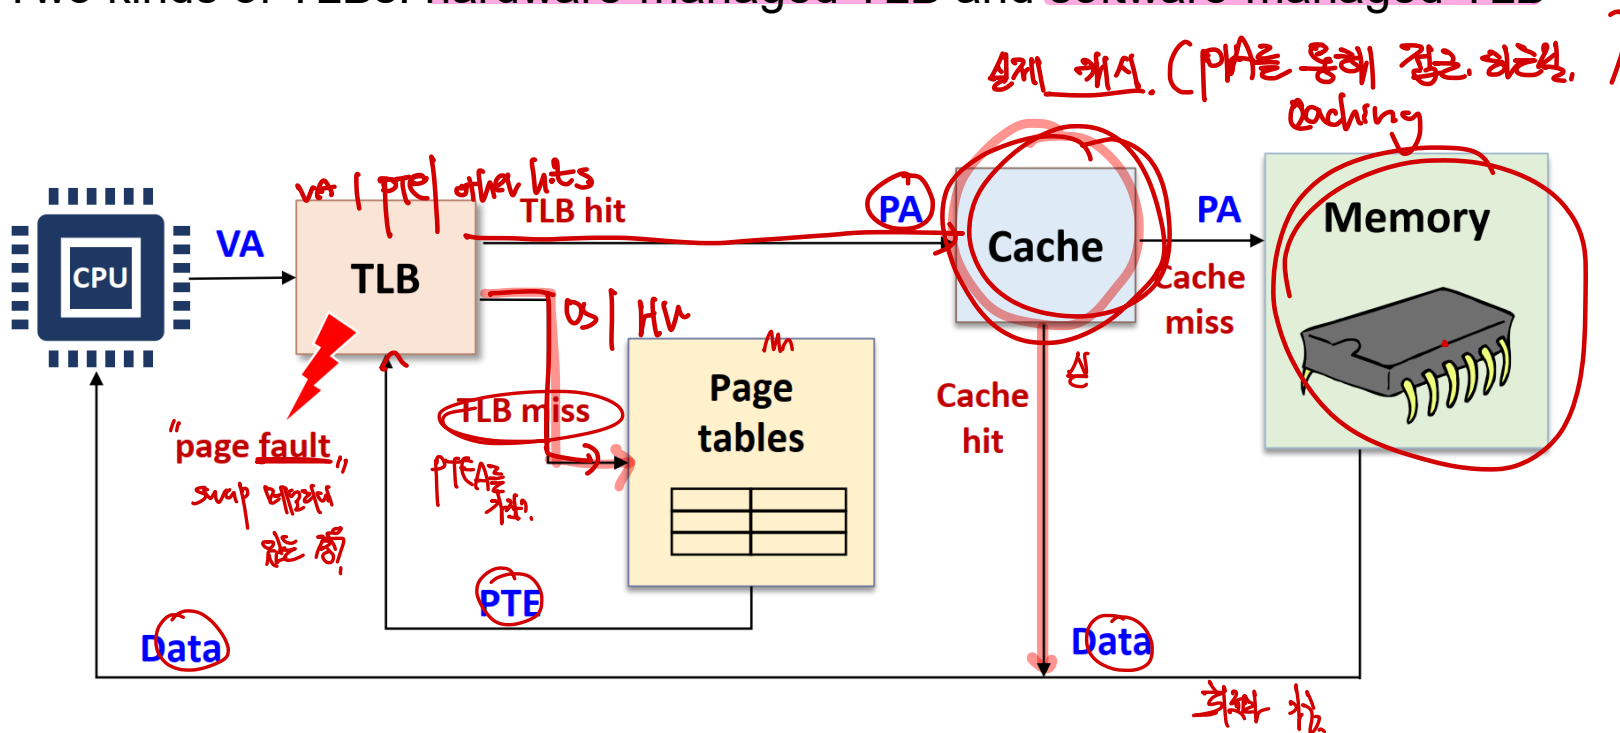
## Like cache, TLB replacement policy should also be considered

- Least-recently-used (LRU) policy (locality) or random policy (simplicity)

# Summary

- Translate lookaside buffer (TLB), part of MMU, is a hardware cache to expedite virtual-to-physical address translations

- A TLB entry includes both VPN and PFN + other bits (valid, protection, etc)
- Two kinds of TLBs: hardware-managed TLB and software-managed TLB



Courtesy of Prof. Jin-Soo Kim @ SNU