

생체 메모리를 아예 디스크로 저장하지.

↳ 메모리의 크기와 큰 프로세스를 실행하려면 how?

Beyond Physical Memory: Mechanisms

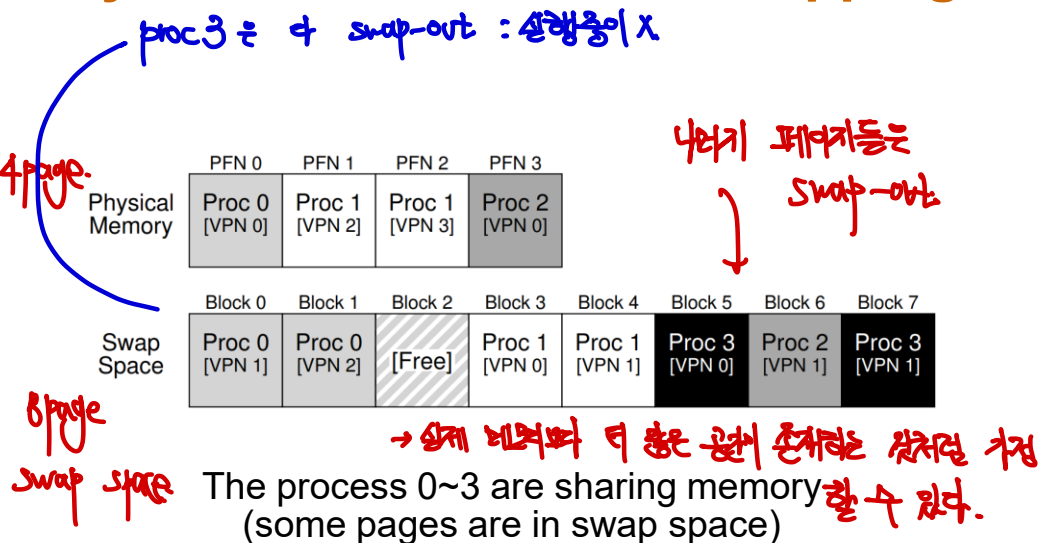
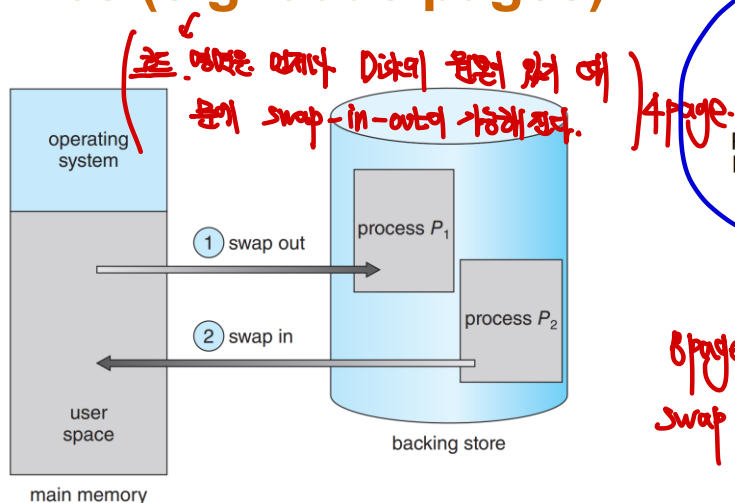


Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Swap Space

- Realistically, an address space doesn't fit into physical memory.
 - OS needs a place to stash away portions of address spaces that currently aren't in great demand → additional level in the memory hierarchy (e.g. disk)
- For moving pages back and forth, OS reserves some space, refer to as **swap space**, on the disk.
 - OS can read from and write to the space in page-sized units
 - OS needs to remember the disk address of a given page
- The swap space is not the only on-disk location for swapping traffics (e.g. code pages)



The Present Bit

→ 주안: 변환은 하려고 할 때 main 메모리에 있는 swap 공간에 있는지도 찾아 한다.

→ TLB swap 사용 방법에 차이 있다.

Overview of the memory reference

- The running process generates a virtual address
- The hardware extracts VPN, checks TLB, and obtains physical address if hit
- Otherwise, the hardware locates the page table in memory and looks up PTE
- If page is valid and present in memory, the hardware obtains physical address

→ present bit에 저장 → 메모리에 존재하지 않다면 (Page Fault) → page fault handler
→ swap 공간에 있는지 찾는

To allow pages to be swapped to disk, more machinery is needed

- When the hardware looks in the PTE, it may find that the page is not present in physical memory and this information is known as the present bit in the PTE

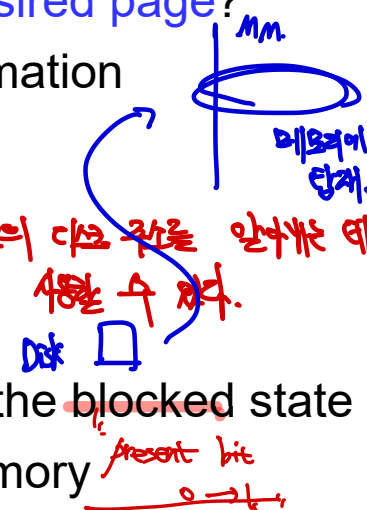
Present Bit	Meaning
0	The page is not in memory but rather on disk somewhere
1	The page is present in physical memory

The act of accessing a page that is not in physical memory is commonly referred to as a page fault

- Upon a page fault, OS is invoked to service the page fault (page fault handler)

Page Fault

- **The page faults are handled by software (OS) for the systems regardless of hardware- or software-managed TLBs.**
 - Note that TLB miss handling is done differently according to TLB type
- **If a page is not present and has been swapped to disk, OS will need to swap the page into memory to service the page fault**
 - Then, the question is how will OS know ^{→ 어디서 찾아야 할까.} where to find the desired page?
 - Usually, the page table is a natural place to store such information
- **When the OS receives a page fault for a page:**
 - 1) Looking in the PTE to find a disk address for the page
 - 2) Issuing the request to disk to fetch the page into memory
 - 3) During disk I/O to swap the page in → the process to be in the blocked state
 - 4) Updating the page table to mark the page as present in memory
 - 5) **Retrying** the instruction → This may result in a **TLB miss**
 - 6) Updating the TLB with the new entry
 - 7) **Restarting** the instruction and this will be a TLB hit → Finally, memory access!



Page Fault Control Flow

What hardware does during the address translation:

```

VPN = (VirtualAddress & VPN_MASK) >> SHIFT
(Success, TlbEntry) = TLB_Lookup(VPN)
if (Success == True)    // TLB Hit
    if (CanAccess(TlbEntry.ProtectBits) == True)
        Offset = VirtualAddress & OFFSET_MASK
        PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
        Register = AccessMemory(PhysAddr)
    else
        RaiseException(PROTECTION_FAULT)
else
    // TLB Miss
    PTEAddr = PTBR + (VPN * sizeof(PTE))
    PTE = AccessMemory(PTEAddr)
    if (PTE.Valid == False)
        RaiseException(SEGMENTATION_FAULT)
    else
        if (CanAccess(PTE.ProtectBits) == False)
            RaiseException(PROTECTION_FAULT)
        else if (PTE.Present == True)
            // assuming hardware-managed TLB
            TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
            RetryInstruction()
        else if (PTE.Present == False)
            RaiseException(PAGE_FAULT)

```

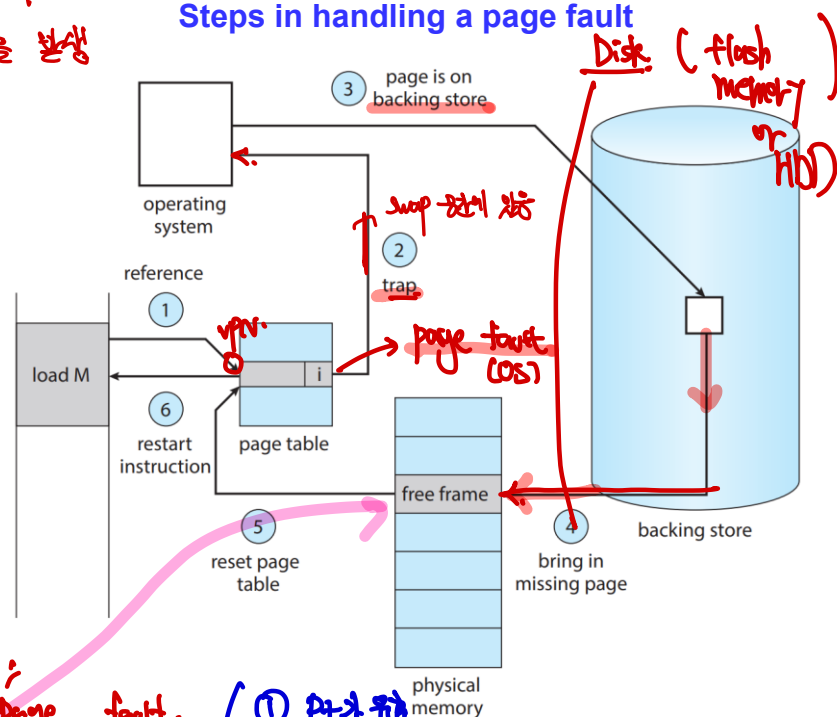
TLB miss (red arrow pointing to the TLB miss case)

valid & present (blue box around the PTE.Present == True case)

valid & !present (blue box around the PTE.Present == False case)

HW base 여는
exception을 발생

Steps in handling a page fault



What OS does upon a page fault:

```

PFN = FindFreePhysicalPage()
if (PFN == -1)
    PFN = EvictPage()
DiskRead(PTE.DiskAddr, PFN)
PTE.present = True
PTE.PFN = PFN
RetryInstruction()

```

OS to find a physical frame for the soon-be-faulted-in page to reside within and run replacement algorithm if failed (blue arrow pointing to the FindFreePhysicalPage() call)

PTE on disk (red text next to PTE.DiskAddr)

OS to find a physical frame for the soon-be-faulted-in page to reside within and run replacement algorithm if failed

When Page Replacement Occur?

- **When memory is full of pages, OS needs to swap old page(s) in memory out to disk to make room for new page(s) in disk**
 - The process of picking a page to kick out, or replace is known as the page-replacement policy.
- **When replacements really occur?**
 - To keep a small amount of memory free, OS usually has high watermark (HW) and low watermark (LW) to decide when to start evicting pages from memory
 - If ($\# \text{ of free pages} < \text{LW}$), a background thread that is responsible for freeing memory runs to evicts pages and goes to sleep if ($\# \text{ of free pages} > \text{HW}$)
 - This background thread is called swap daemon or page daemon
- **The performance optimizations become possible by performing a number of replacements at once**

→ page cluster → Disk I/O 비용 최소화.

 - Many systems will cluster or group a number of pages, and write them out at once to the swap partition, thus increasing the efficiency of the disk