

본 페이지 리스트에서 이미 있는 페이지를 찾아서 링크를
↑
없는 페이지의 링크를 삭제합니다.

Beyond Physical Memory: Policies

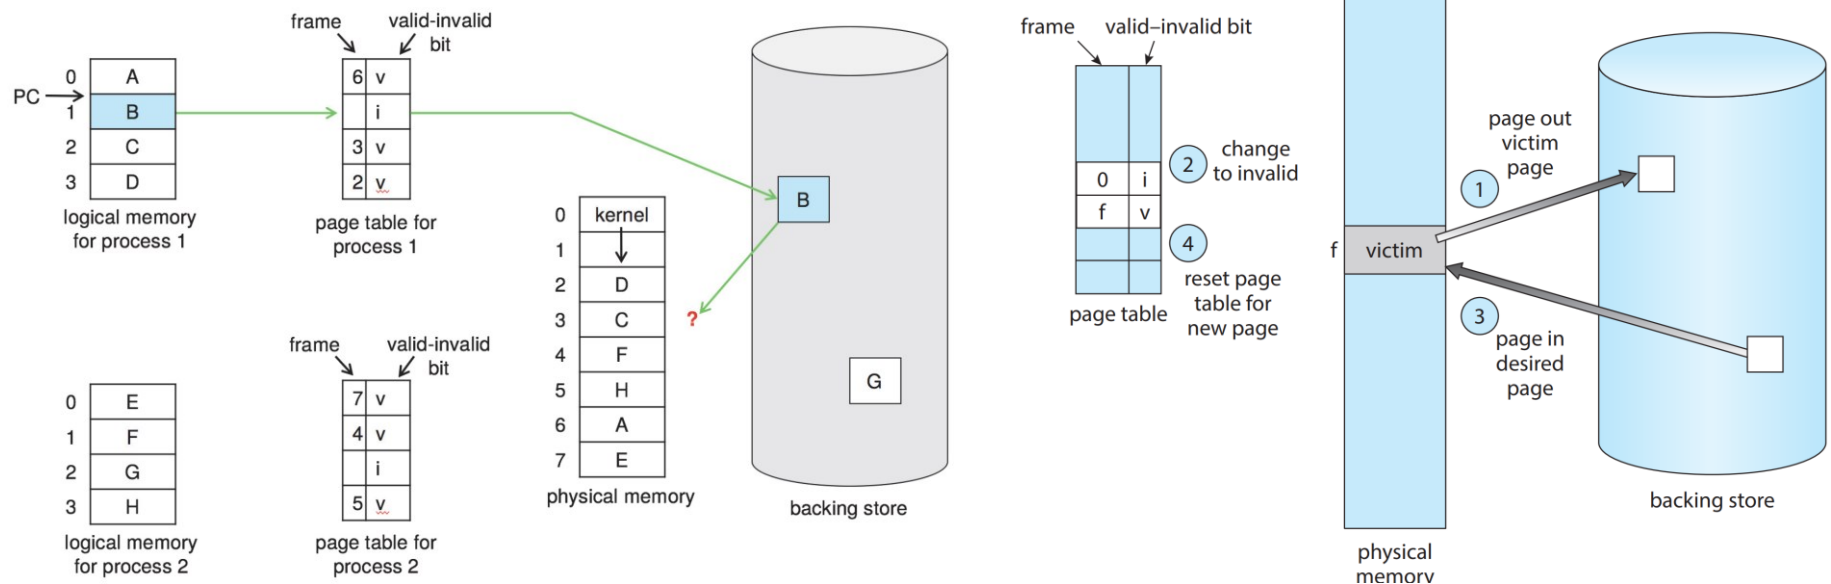


Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Page Replacement: Concept

- The memory pressure forces the OS to start paging out pages to make room for actively-used pages
 - Deciding which page(s) to evict is encapsulated within OS's replacement policy
- Since **main memory** holds a subset of all pages in system, it can be viewed as a **cache** for virtual memory pages in the system
 - Then, our primary goal in picking a replacement policy for this cache is to **minimize the number of cache misses** (i.e. maximize the cache hit rate)



Cache Management

→ Cache miss는 최악의 경우를 뜻함.

$$AMAT = T_M + (P_{Miss} \times T_D)$$

- Considering average memory access time (AMAT) for a program

$$AMAT = T_M + (P_{Miss} \cdot T_D)$$

↓
평균访问时间

Parameter	Meaning
T_M	The cost of <u>accessing memory</u>
T_D	The cost of <u>accessing disk</u>
P_{Miss}	The probability of <u>not</u> finding the data in the cache (0 ~ 1.0)

- AMAT examples** memory access disk access

- $T_M = 100 \text{ ns}$, $T_D = 10 \text{ ms} = 10,000,000 \text{ ns}$
- $P_{Miss} = 0.1$ (90% hit): $AMAT = 100 \text{ ns} + (0.1 \times 10 \text{ ms}) = 1.0001 \text{ ms} \approx 1 \text{ ms}$
- $P_{Miss} = 0.001$ (99.9% hit): $AMAT = 100 \text{ ns} + (0.001 \times 10 \text{ ms}) = 10.1 \mu\text{s}$
- Hit rate 90% → 99.9% (9.99%p difference) results in AMAT 100× faster

$$100\text{ns} + (0.1 \times 10\text{ms}) = 10 \times 10^6 \text{ ns}$$

$$= 10^4 \text{ ns} + 10^6 \text{ ns}$$

$$10^{-4} \text{ ms} + 1 \text{ ms} = 1.0001 \text{ ms}$$

$$\approx 1 \text{ ms}$$

- The cost of disk access is so high in modern systems**

- Therefore, even a tiny miss rate will quickly dominate the overall AMAT of running programs

The Optimal Replacement Policy (MIN)

- The optimal replacement policy leads to the fewest number of misses overall.

[^]Belady (optimal 정책)

가장 나중에 사용될 page를 교체하는 방법.

- The optimal policy is to replace the page that will be accessed furthest in the future, resulting in the fewest-possible cache misses

but 왜냐하면 어떤 page가 사용될 것인지 알 수 x

- Example with virtual page access of 0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1 in a row and a cache with capacity of 3 pages

그래서 다음 optimal policy/의 비교대상 요함.

- Not surprisingly, the first three accesses are misses due to empty cache and such a miss is sometimes called **cold-start miss** (or **compulsory miss**)
- When we reach another miss (page 3), the page 2 will be accessed furthest in the future → OS evicts page 2
- We can't build the optimal policy for general purpose OS and use it as **reference for comparison** as future is usually not known

$$\text{hit rate} = \frac{\text{hits}}{\text{hits} + \text{misses}} = \frac{6}{6 + 5} = 54.5\%$$

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2
3	Miss	2	0, 1, 3
0	Hit		0, 1, 3
3	Hit		0, 1, 3
1	Hit		0, 1, 3
2	Miss	3	0, 1, 2
1	Hit		0, 1, 2

한시 ←
발생하는
miss

할 수 없음
↑ (알다 가정)

아래의 가장 나중에 사용될
page를 교체.

A Simple Policy: FIFO

→ 가장 먼저 시스템 page를 miss 발생했을 때 교체할 것임
여기

■ FIFO (first-in, first-out) replacement was used early systems

- Pages were simply placed in a queue when they enter the system
- When a replacement occurs, the page on the front of the queue is evicted
- FIFO has one great strength: quite simple to implement
- Cons: hit rate may not increase when the cache gets larger (Belady's Anomaly)

■ Example with the same access sequence and cache

cache가 커질수록 hit-rate가 높아져 였다.

- Three compulsory misses to pages 0, 1, and 2, and then hit on both 0 and 1
- When page 3 is referenced, causing a miss, OS evicts page 0 that was the first one in order

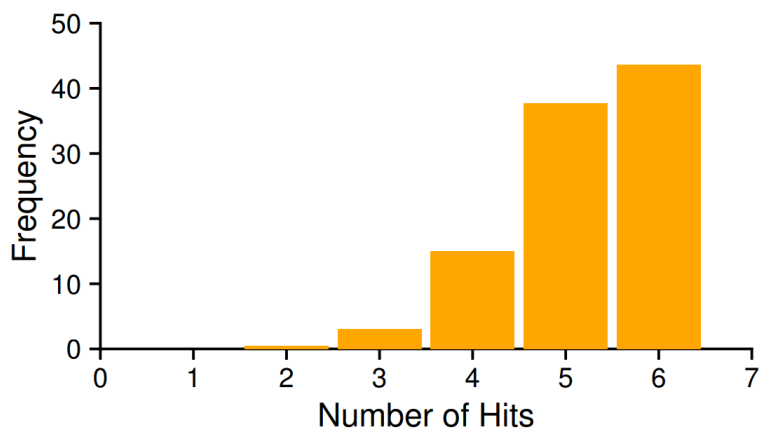
$$\text{hit rate} = \frac{\text{hits}}{\text{hits} + \text{misses}} = \frac{4}{4 + 7} = 36.4\%$$



Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		First-in→ 0
1	Miss		First-in→ 0, 1
2	Miss		First-in→ 0, 1, 2
0	Hit		First-in→ 0, 1, 2
1	Hit		First-in→ 0, 1, 2
3	Miss	0	First-in→ 1, 2, 3
0	Miss	1	First-in→ 2, 3, 0
3	Hit		First-in→ 2, 3, 0
1	Miss	2	First-in→ 3, 0, 1
2	Miss	3	First-in→ 0, 1, 2
1	Hit		First-in→ 0, 1, 2

Another Simple Policy: Random

- **Random replacement policy** simply picks a random page to **replace under memory pressure**
 - Random has properties similar to FIFO; simple to implement
 - Cons: unpredictable → performance varies
- **Example with the same access sequence and cache**
 - Three compulsory misses and OS randomly evicts the page upon a miss



Random performance over 10,000 trials:
 sometimes it is as good as optimal (6 hits)
 sometimes it does much worse (2 hits or fewer)

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2
3	Miss	0	1, 2, 3
0	Miss	1	2, 3, 0
3	Hit		2, 3, 0
1	Miss	3	2, 0, 1
2	Hit		2, 0, 1
1	Hit		2, 0, 1

$$\text{hit rate} = \frac{\text{hits}}{\text{hits} + \text{misses}} = \frac{5}{5 + 6} = 45.5\%$$

Using History: LRU

- **FIFO or random policy might kick out an important page, one that is about to be referenced again**
 - We lean on the past and use history to improve our guess at the future (locality)
- **A family of simple historically-based algorithms are born**
 - Least-Frequently-Used (LFU) policy replaces the least-frequently-used pages
 - Least-Recently-Used (LRU) policy replaces the least-recently used ones
- **Example with the same access sequence and cache (LRU)**
 - Three compulsory misses and OS picks least-recently-used page for eviction

$$\text{hit rate} = \frac{\text{hits}}{\text{hits} + \text{misses}} = \frac{6}{6 + 5} = 54.5\%$$



- Opposites of these algorithms:
 - 1) Most-Frequently-Used (MFU)
 - 2) Most-Recently-Used (MRU)

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		LRU → 0
1	Miss		LRU → 0, 1
2	Miss		LRU → 0, 1, 2
0	Hit		LRU → 1, 2, 0
1	Hit		LRU → 2, 0, 1
3	Miss	2	LRU → 0, 1, 3
0	Hit		LRU → 1, 3, 0
3	Hit		LRU → 1, 0, 3
1	Hit		LRU → 0, 3, 1
2	Miss	0	LRU → 3, 1, 2
1	Hit		LRU → 3, 2, 1

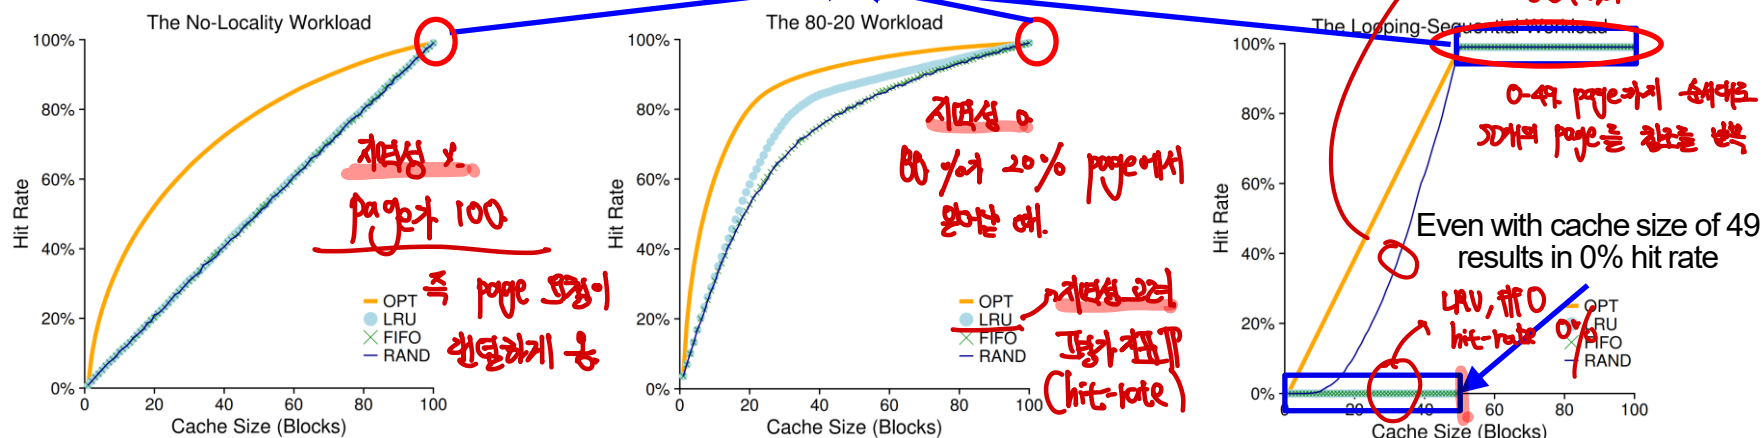
Workload Examples

Examining more complex workloads to better understand

- 50 or 100 unique page accesses over time, choosing next page randomly → overall 10,000 accesses; cache size from 1 to 100 pages (holding 1 ~ all pages)

- No-locality workload:** LRU, FIFO, and random all perform the same, with the hit rate exactly determined by the size of the cache, while optimal is the best
- 80-20 workload (80% references by 20% pages):** while both random and FIFO do reasonably well, LRU does better (20% hot pages referenced again)
- Looping-sequence workload (0~49 page in a row and repeat):** while both LRU and FIFO show a worst (kicking out older pages), random is better

When the cache is large enough to fit the entire workload, it doesn't matter which policy you use; 100% hit rate

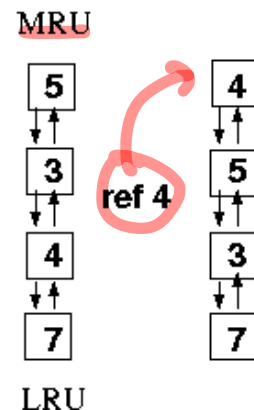


Implementing Historical Algorithms

→ policy 자체는 어렵지 않음. 구현이 중요함.

■ To implement LRU perfectly, a lot of work is required

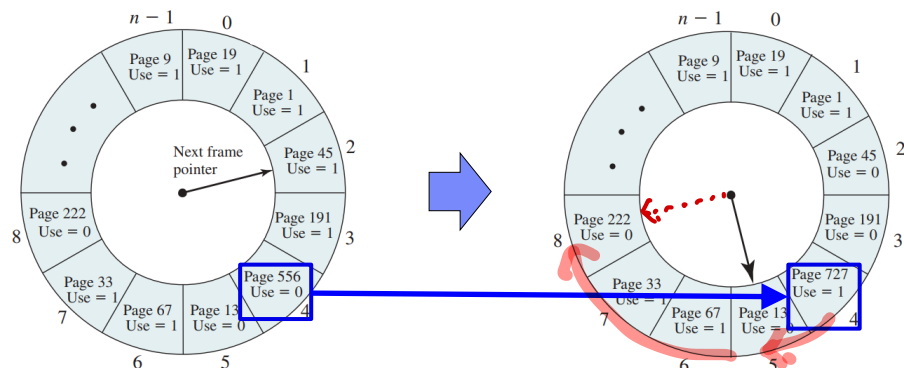
- Upon each page access, some data structure needs to be updated to **move this page to the front of this list** (MRU side)
- To keep track which pages were least- and most-recently used, accounting work on every memory reference is required
- To speed-up, **hardware's help** is possible (e.g. time stamp)



■ Approximating LRU is more feasible for computation efficiency

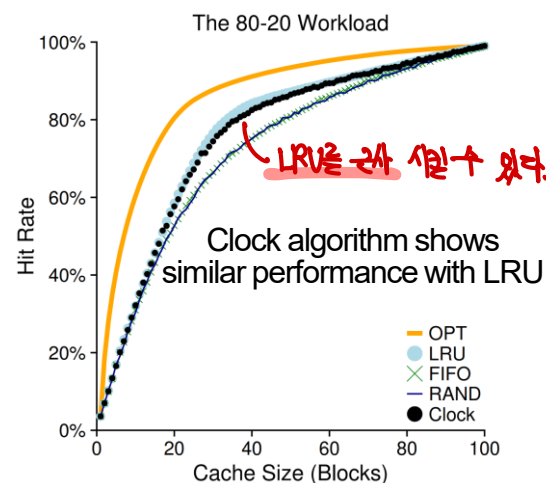
→ 페이지 사용 횟수를 기록하는 방법 → 레퍼런스

- This requires some hardware support, a use (reference) bit; whenever a page is referenced, the bit is set by hardware to 1, then OS sets it to 0 by own algorithm
- **Clock algorithm**: moving clock until finding a page with use bit = 0; then replacing it to new with use bit = 1



page 121 swap in

use = 1
→ 바뀌어 있음



Considering Dirty Pages & Other Policies



- If a page has been **modified** and is thus **dirty**, it must be **written back to disk to evict it, which is expensive** CPU가 재한 것들

- If it has not been modified (and is thus clean), the eviction is free; the physical frame can be reused without additional I/O
- Thus, some systems prefer to evict clean pages over dirty pages
- To support this behavior, the hardware should include a (modified (dirty) bit) //

→ 그래서 수정되지 않은 page 먼저 교체하도록 함.

■ Other policies for virtual memory system

- OS has to decide **when to bring** a page into memory; **page selection policy**:
 - 1) For most pages, OS brings the page into memory when it is accessed, “on demand”; **demand paging**
 - 2) OS could **guess** that a page is about to be used, and thus bring it in ahead of time; **prefetching** (e.g. code pages P & P+1 together)
- Another policy determines **how the OS writes** pages out to disk:

Many system collect a number of pending writes together in memory and write them to disk in **one** (more efficient) write; (**clustering** or **grouping**) of writes

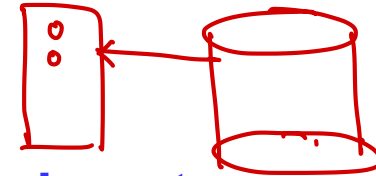
(예측)

먼저 가져올 수도 있음 ↔ on-demand.
T = prefetch

→ delay write (거면 쓰지)를 위해 page에 관한 작은 한개. 무한. 쓰기 하.
정해진 어느 정도와 한 번에 disk에 저장한다

Thrashing

= 실행 중인 프로세스들이 요하는 메모리 크기가 실제 메모리 크기보다 큰 경우.



What should OS do when physical memory is not enough to hold all the working sets of processes?

- The **working set** indicates a set of pages that a process is using actively
- The system will **constantly be paging**; most of the time is spent by OS paging data back and forth over the disk, called **thrashing**, (degrading CPU utilization)

Possible solutions are

- ① **Admission control**: reducing set of working set to fit in memory by not running a subset of processes
- ② **Out-of-memory killer**: killing some memory-intensive processes when memory is oversubscribed; causes problems
- ③ Buy more memory!

mm ← disk → paging mm이 부족하면
swap한 데만 실행

시스템의 사용 가능한 리소스를 바탕으로
새로운 프로세스를 수락할지 여부를 결정

