

Locks



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Locks: The Basic Idea & Pthread Locks

lock의 사용 매커니즘

- To use a lock, some code around the critical section is required

```
lock_t mutex; // some globally-allocated lock 'mutex'
...
lock(&mutex);
balance = balance + 1;
unlock(&mutex);
```

← critical section

- A **lock** is a variable that holds state of lock at any instant in time

- The state is either 1) available (or unlocked or free) that no thread holds the lock or 2) acquired (or locked or held) that exactly one thread holds the lock

- Calling the routine `lock()` tries to acquire the lock

- If no other thread hold the lock, the thread will acquire it and become the **owner**
- If another thread calls `locks()` on the same lock variable, it will not return and thus other threads are prevented from entering the critical section
- Once the owner calls `unlock()`, the lock is available again

스레드 자원 공유
→ /연산을 위해 원자적으로
수행 → 데이터 무결성 보장

- The name that the **POSIX** library uses for a lock is a **mutex**, as it is used to provide mutual exclusion between threads

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

```
Pthread_mutex_lock(&lock); // wrapper; exits on failure
```

```
balance = balance + 1;
```

```
Pthread_mutex_unlock(&lock);
```

→ 락에 취해 가림.
spin으로 처리.

Building and Evaluating a Lock

- **To build a working lock, (hardware and OS) need collaboration**
 - Over the years, a number of different hardware primitives have been added to the instruction sets of various computer architectures
 - OS gets involved to complete building a sophisticated locking library

명령어 셋
↓
OS의 도움을 받아야 함
- **To evaluate whether a lock works (and works well), some basic criteria need to be established**

상호배제. (호출을 고려해야 함)

 - 1) **Mutual exclusion**: does the lock work, preventing multiple threads from entering a critical section?

공정성 → 가아 상태가 있어야만 안된다.
 - 2) **Fairness**: does each thread contending for the lock get a fair shot at acquiring it once it is free? (lock starvation)

상호 → 실패이라는 자원의 overhead.
 - 3) **Performance**: the time overheads added by using the lock; considering three cases of no contention, multiple threads on a single CPU or multiple CPUs

→ 고려하여 lock/unlock 과정이 불필요

 - ① thread가 1개일 때 lock.
 - ② cpu가 1개일 때 여러 thread가. 경쟁하는 상태.
 - ③ cpu가 여러개일 때. 여러 thread가. lock은 최위에서 경쟁하는 상태

Controlling Interrupts

- One of the earliest solution to provide **mutual exclusion** was to **disable interrupts** for critical sections


- The code inside the critical section (will not be interrupted), thus will execute as if it's atomic; its **simplicity** is the main benefit but quite **slow** due to turning it on/off.

```
void lock() {
    DisableInterrupts();
}
void unlock() {
    EnableInterrupts();
}
```

- Since disabling interrupts requires a **privileged operation**, we need **trust** that this facility is not abused (e.g. infinite loop can monopolize the processor)
- This **does not work on multiprocessors**; each of multiple threads on different CPUs tries to enter the same critical section, disabling on a CPU doesn't matter
- Turning off interrupts for long time can lead to interrupt loss

- For these reason, turning off interrupts is only used in **limited context as a mutual-exclusion primitive**

- e.g.) OS itself will use interrupt masking to guarantee atomicity when accessing its own data structure, or at least to prevent certain messy interrupt handling



↑ 이 부분은 critical section에서 interrupt를
기면 됨.

효율성 (느임)

느리게 처리 함.

이것이 중요

→ timer interrupt x.

↓
다른 프로세스 환경 문제 x.
있을 수 있음.

cpu disk I/O interrupt
무시 x.

→ OS 자체 작동에 lock, unlock은 시스템 레이어나
시스템 수준

A Failed Attempt: Just Using Loads/Stores

→ 하드웨어 지원 필요.


Let's try to build a simple lock by using a single flag variable

```
typedef struct __lock_t { int flag; } lock_t;
```

```
void init(lock_t *mutex) {
    // 0 -> lock is available, 1 -> held
    mutex->flag = 0;
}
```

```
void lock(lock_t *mutex) {
    while (mutex->flag == 1) // TEST the flag
        ; // spin-wait (do nothing)
    mutex->flag = 1; // now SET it!
}
```

```
void unlock(lock_t *mutex) {
    mutex->flag = 0;
}
```



Assuming flag = 0 to begin

Thread 1

call lock ()
while (flag == 1)
interrupt: switch to Thread 2

Thread 2

call lock ()
while (flag == 1)
flag = 1;
interrupt: switch to Thread 1

flag = 1; // set flag to 1 (too!)

context-switching 때 진행 상황 기록

원자성 보장
되지 않음.

동시 access 가능함

- Using a simple variable (flag) to indicate if any thread has the lock
- lock () to test whether flag = 0/1 and set the flag to 1 to acquire lock if flag = 0
- unlock () to clear the flag that indicates that the lock is no longer held

This software-only approach has two problems

- Correctness: multiple threads are able to enter the critical section (right figure)
- Performance: checking the flag continuously while flag=1; spin-waiting.

test-and-set | atomic exchange.

cpu 처리 속도 유 유 한 (수요)

To resolve issues, we need hardware support to implement lock

Building Working Spin Locks with Test-And-Set


■ The simplest hardware support is test-and-set instruction

- It returns the old value pointed by `old_ptr`, and simultaneously updates the value `new`; the key is that this sequence is performed **ATOMICALLY!**

It is sometimes called atomic exchange memory with register

```
int TestAndSet(int *old_ptr, int new) {
    int old = *old_ptr; // fetch old value at old_ptr
    *old_ptr = new;      // store 'new' into old_ptr
    return old;          // return the old value
}
```

x86: `xchg`
RISC-V: `amoswap`
SPARC: `ldsturb`

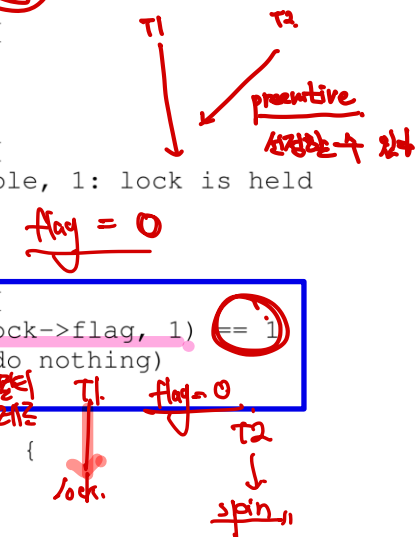


```
typedef struct __lock_t {
    int flag;
} lock_t;

void init(lock_t *lock) {
    // 0: lock is available, 1: lock is held
    lock->flag = 0;
}
```

```
void lock(lock_t *lock) {
    while (TestAndSet(&lock->flag, 1) == 1)
        ; // spin-wait (do nothing)
}
```

```
void unlock(lock_t *lock) {
    lock->flag = 0;
}
```



- flag = 0: `TestAndSet()` tests the flag and sets it to 1; return 0 and acquire lock
- flag = 1: `TestAndSet()` tests the flag and sets it to 1 again; return 1 and spinning until the lock is released → when flag=0, then can acquire lock

■ To work correctly on a single processor, this approach requires preemptive scheduler.

- Without preemption, a thread spinning on a CPU will never relinquish it

Evaluating Spin Locks

- The most important aspect of a lock is **correctness**:

- Does the spin lock provide mutual exclusion? **Yes** we have a correct lock

정확성

↓
- lock 보지 TestAndSet

- The next aspect is **fairness**:

- How fair is a spin lock to a waiting thread?
- In other words, can it guarantee that a waiting thread will ever enter the critical section? Unfortunately, **not fair**; this may lead to starvation

spin-lock은 공평성을 보장하지 않음 x.

- The final aspect is **performance**:

- In **single CPU**: the performance overheads can be quite **painful**; one of a lot of threads is holding the lock and the others keep spinning, a **waste of CPU cycle**
e.g.) T0~T9 try to acquire a lock, T0 has the lock for 30ms, time slice: 20ms
Single CPU: T0 runs for 20ms (lock) → T1~T9 **spin one by one** for 20ms
→ T1 releases the lock after 180ms → one of T1~T9 can obtain the lock
- In **multiple CPUs**: spin locks **work reasonably well** (# of threads \approx # of CPUs);
e.g.) T0~T9 on CPU0~9 try to get a lock and T0 acquires the lock for 30 ms
T0 runs and T1~T9 **spins simultaneously** for 30ms → one of T1~T9 will get the lock after 30ms (right after T1 releases the lock)

T0 ~ T9 RR 20ms
lock. C T1 ~ T9

spin-wait 20ms를 해야함

→ 비로 끝나지 않을 수 있다.

Compare-And-Swap

- Another hardware primitive is **compare-and-swap** or **compare-and-exchange (x86) instruction**

- It tests whether the value at the address specified by **ptr** is equal to **expected**
- If so, **update the memory pointed by ptr with new**, otherwise, do nothing
- In either case, it returns the original value at that memory location

```
int CompareAndSwap(int *ptr, int expected, int new) {
    int original = *ptr;
    if (original == expected)
        *ptr = new;
    return original;
}
```

x86: **cmpxchg**

```
void lock(lock_t *lock) {
    while (CompareAndSwap(&lock->flag, 0, 1) == 1)
        ; // spin
}
```

- CompareAndSwap()** simply checks if flag = 0 and if so, it atomically swaps in a 1 thus acquiring the lock
- Actually, **compare-and-swap** is more powerful instruction than test-and-set in lock-free synchronization, but their behavior is identical in terms of spin-lock

- In x86, **cmpxchg** is the instruction for **compare-and-swap**

- If you want to see how to make a C-callable x86-version of compare-and-swap:
<https://github.com/remzi-arpacidusseau/ostep-code/tree/master/threads-locks>

Load-Linked and Store-Conditional

- The **load-linked and store-conditional instruction pair** can be used to build locks and other concurrent structures
 - The load-linked instruction simply **fetches a value from memory to a register**
 - With no intervening store to the address, the store-conditional instruction **returns 1 with updating the value at ptr to value**, otherwise, it returns 0 with no update

RISC-V, MIPS, Alpha, PowerPC, ARM support these instructions

```
int LoadLinked(int *ptr) {
    return *ptr;
}
```

→ load 명령어.

```
int StoreConditional(int *ptr, int value) {
    if (no update to *ptr since LoadLinked to this address) {
        *ptr = value;
        return 1; // success!
    } else {
        return 0; // failed to update
    }
}
```

lock T1 T2 T3

→ lock을 가한 것이 변형되지 않았을 때만.

```
void lock(lock_t *lock) {
    while (1) {
        while (LoadLinked(&lock->flag) == 1)
            ; // spin until it's zero
        if (StoreConditional(&lock->flag, 1) == 1)
            return; // if set-it-to-1 was a success: all done
        // otherwise: try it all over again
    }
}
```

연에 끼어 spin

반복 실행을 통해 lock을 얻기 위해 시도하시 동시에 성공하러

```
void lock(lock_t *lock) {
    while (LoadLinked(&lock->flag) ||
           !StoreConditional(&lock->flag, 1))
        ; // spin
}
```

concise form

- **LoadLinked()** spins waiting for the flag to be set to 0 (lock is not held)
- **StoreConditional()** tries to **acquire lock** and atomically changes the flag to 1
- If two threads call **LoadLinked()** and both get 0 (lock is free), **one of them will get the lock** with **StoreConditional()** and the **other will fail** since the winning thread updated the flag between its **LoadLinked()** and **StoreConditional()**

Fetch-And-Add

turn을 도출해서. / 모든 thread가 인젠 lock을 얻을 수 있게 되는데.

One final hardware primitive is fetch-and-add instruction

- It atomically increments a value while returning the old value, at a particular address and can be used to build a ticket lock.

```
int FetchAndAdd(int *ptr) {
    int old = *ptr;
    *ptr = old + 1;
    return old;
}
```

x86: xadd

```
typedef struct __lock_t {
    int ticket;
    int turn;
} lock_t;

void lock_init(lock_t *lock) {
    lock->ticket = 0;
    lock->turn = 0;
}

void lock(lock_t *lock) {
    int myturn = FetchAndAdd(&lock->ticket);
    while (lock->turn != myturn)
        ; // spin
}

void unlock(lock_t *lock) {
    lock->turn = lock->turn + 1;
}
```

$T_0 = 0$ $T_1 = 1$ $T_2 = 2$
 \vdots
 $T_9 = 9$

turn 0부터 증가

fairness가 보장된다.

- FetchAndAdd()** increases the ticket value that is considered the thread's turn
- The globally shared **lock->turn** is used to determined the next turn
- Unlock is done by incrementing the turn so next thread can have the lock if any
- This approach ensures progress for all threads.** fairness
 - Once a thread has its ticket value, it will be scheduled at some point in future

Too Much Spinning

→ lock을 가지고 있지 않은 thread는
스피닝으로 리소스 작업을 낭비함! } → OS 개입해!

■ The hardware-based locks are simple and works well.

- However, in some cases, these solutions can be quite inefficient
- Threads not holding the lock keep spinning, wasting an entire time slice
- To solve this issue, OS support is necessary

spin 하는 동안에는 다른 스레드에게

■ A simple approach to avoid spin is to yield CPU를 할당하는 것.

- OS primitive yield() is required to give up the CPU and let another thread run




기전

```
void init() {
    flag = 0;
}
```

```
void lock() {
    while (TestAndSet(&flag, 1) == 1)
        yield(); // give up the CPU
}
```

```
void unlock() {
    flag = 0;
}
```

sleep mode로 전환
다른 프로세스가 실행할 수 있음.



- Two threads: a thread has a lock and the other yields the CPU, then thread with the lock will run and finish faster its critical section; this works well
- Many threads: a thread holds a lock, each of the others execute this run-and-yield before getting the lock; this still has cost (context switch ↑)
- Worse, this approach does not resolve the starvation problem at all

(차 문제 해결 X)

문제 해결 X

overhead ↑

Using Queues: Sleeping instead of Spinning

→ 만약 lock을 가진게라 unlock후에 다른 많은 스레드를 직접 깨우면

- To provide fairness, we need some **explicit control** for threads

- To do this, we will need **queue** to keep track of which threads are waiting for lock
- Two calls: **park()** to sleep, **unpark(threadID)** to wake a thread with threadID

```
typedef struct lock_t {
    int flag;
    int guard;
    queue_t *q;
} lock_t;
```

```
void lock_init(lock_t *m) {
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void lock_acquire(lock_t *m) {
    while (TestAndSet(&m->guard, 1) == 1)
        ; //acquire guard lock by spinning
    if (m->flag == 0) {
        m->flag = 1; // lock is acquired
        m->guard = 0;
    } else {
        queue_add(m->q, gettid());
        m->guard = 0;
        park();
    }
}
```

```
void lock_release(lock_t *m) {
    void unlock(lock_t *m) {
        while (TestAndSet(&m->guard, 1) == 1)
            ; //acquire guard lock by spinning
        if (queue_empty(m->q))
            m->flag = 0; // let go of lock; no one wants it
        else
            unpark(queue_remove(m->q)); // hold lock
            // (for next thread!)
        m->guard = 0;
    }
}
```

setpark() is added to resolve wakeup/waiting race

```
queue_add(m->q, gettid());
setpark(); // new code
m->guard = 0;
```

- The **guard** is used as **spin lock** around flag and queue; Acquiring the guard lock is first step and this lock is released within a few instructions, → **spin overhead?**
- When a thread can't acquire the lock, it is added to a queue, sleeps with guard=0
- What if T1 released the lock just before **park()** in T2? → T2 parks and sleeps forever (**wakeup/waiting race**) as no thread will call **unlock()** → **setpark()** indicates it is about to park → **park()** returns immediately instead of sleeping

Different OS, Different Support

(운영체제마다 lock을 제공하는 scheme 이 다름.)

OSes provide their own lock scheme in the thread library

- Linux provides a **futex** which provides some in-kernel functionality,
- Each futex has associated with it a specific physical memory location and a per-futex in-kernel queue
- Futex offers two calls: **futex_wait()** and **futex_wake()**
- Please see the Linux-based futex lock code at pp. 349

Two-phase lock realizes that spinning can be useful, particularly if the lock is about to be released

- 1st phase: spinning for a while, hoping that it can acquire the lock
- 2nd phase: if the lock is not acquired in 1st spin phase, the caller sleeps when 2nd phase is entered and only woken up when the lock becomes free later.

↓
second phase의 sleep
sleeping

‘자빠리 spin이 더 효율적일 수 있다.’

→ flag