



# Lock-based Concurrent Data Structures


→ 자료구조는 어떻게 thread-safety 하게!



**Prof. Yongtae Kim**

Computer Science and Engineering  
Kyungpook National University

# Concurrent Counters



## ■ A concurrent counter requires locks to make it work correctly

- Before accessing the counter variable, each thread must acquire the lock

```
typedef struct __counter_t {
    int      value;
    pthread_mutex_t lock;
} counter_t;

void init(counter_t *c) {
    c->value = 0;
    Pthread_mutex_init(&c->lock, NULL);
}

void increment(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    c->value++;
    Pthread_mutex_unlock(&c->lock);
}
```

*lock 추가.*

```
void decrement(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    c->value--;
    Pthread_mutex_unlock(&c->lock);
}

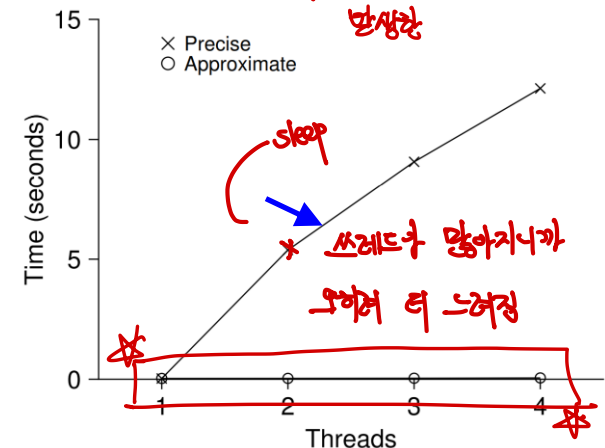
int get(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    int rc = c->value;
    Pthread_mutex_unlock(&c->lock);
    return rc;
}
```

*lock, unlock은 사용하는 데미지 행위 들*

*→ 스케줄러는 어떤 사용자를 대하  
게를 사용할 때부터  
이 빠른 반응이  
발생함*

## ■ The problem of this implementation is performance

- Each thread update the counter one million times
- Varying the number of threads (1~4)
- Running on iMac with four Intel 2.7GHz i5 CPUs
- Unfortunately, the performance scales poorly
- Adding thread leads to massive slowdown



# Scalable Counting

(이제 다른 스레드도 가능 X)

→ 스레드와 어레이도 스레드 증가 불능 X

- **Perfect scaling:** though more work is done, it is done in parallel

- The approximate counter works by representing a single logical counter via many (local physical counters), one per CPU core, and a single global counter
- Each local counter and the global counter have their own lock

- The basic idea of approximate counting is

thread 마다 로컬 (local)  
→ 4개의 다. 카운트 (global)

- 1) Each thread running on a given core increments its local counter using lock
- 2) The local values are periodically transferred to the global counter using lock
- 3) The local counter resets to zero after the transfer


- Tracing the approximate counters:

(이제 4)

Time	$L_1$	$L_2$	$L_3$	$L_4$	$G$
0	0	0	0	0	0
1	0	0	1	1	0
2	1	0	2	1	0
3	2	0	3	1	0
4	3	0	3	2	0
5	4	1	3	3	0
6	5 → 0	1	3	4	5 (from $L_1$ )
7	0	2	4	5 → 0	10 (from $L_4$ )

5를 넘으면 local counter의  
값을 global counter에 반영

threshold  
↓  
global counter를 업데이트  
↑  
해결  
후 리셋



S = 5

# Scalability Factor

이 값이 낮을수록 성능이 좋을까?

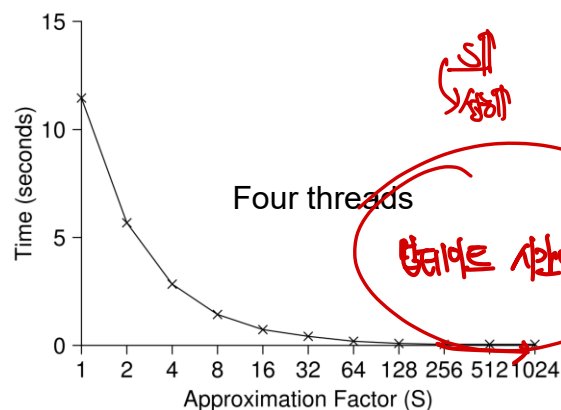

- How often this local-to-global transfer occurs is determined by a threshold  $S$

속도, 정확도 trade-off.

- The smaller  $S$  is, the more the counter behaves like the non-scalable counter
- The bigger  $S$  is, the more scalable the counter, but the further off the global value might be from the actual count

- The performance of approximate counter is excellent

- The time does not increase significantly as the number of thread increases
- If  $S$  is high, the performance is excellent but the global counter lags
- If  $S$  is low, the performance is poor (accurate)



```
typedef struct __counter_t {
    int global; // global count
    pthread_mutex_t glock; // global lock
    int local[NUMCPUS]; // per-CPU count
    pthread_mutex_t llock[NUMCPUS]; // ... and locks
    int threshold; // update frequency
} counter_t;

void update(counter_t *c, int threadID, int amt) {
    int cpu = threadID % NUMCPUS;
    pthread_mutex_lock(&c->llock[cpu]);
    c->local[cpu] += amt;
    if (c->local[cpu] >= c->threshold) {
        // transfer to global (assumes amt>0)
        pthread_mutex_lock(&c->glock);
        c->global += c->local[cpu];
        pthread_mutex_unlock(&c->glock);
        c->local[cpu] = 0;
    }
    pthread_mutex_unlock(&c->llock[cpu]);
}
```

cpu 개수를 늘리면 counter 성능

S가 클수록

가장 빠른 상태

# Concurrent Linked Lists

## ■ We next examine a more complicated structure, the linked list

```
// basic node structure
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;
```

```
int List_Insert(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        pthread_mutex_unlock(&L->lock);
        return -1; // fail
    }
    new->key = key;
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
    return 0; // success
}
```

lock 세션티

unlock을 위한 코드가 왜?

↓  
변경

상호작용에 mutex를  
참고 하기

```
// basic list structure (one used per list)
typedef struct __list_t {
    node_t *head;
    pthread_mutex_t lock;
} list_t;
```

```
int List_Lookup(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *curr = L->head;
    while (curr) {
        if (curr->key == key) {
            pthread_mutex_unlock(&L->lock);
            return 0; // success
        }
        curr = curr->next;
    }
    pthread_mutex_unlock(&L->lock);
    return -1; // failure
}
```

이런 때는 lock을  
풀어야 함

- The code acquires a lock in the insert routine upon entry and releases upon exit
- If `malloc()` fails, the code must also release the lock before failing the insert

## ■ This kind of exceptional control flow is to be quite error prone

- Can we rewrite the insert and lookup codes to remain correct under concurrent insert but avoid the case where the failure path requires the call to unlock?
- The answer is YES


# Revised Concurrent Linked List

- We can do the lock and release only surround the actual critical section and a common exit path is used in the lookup code

- The former works because part of the insert actually need not be locked
- Only when updating the shared list does a lock need to be held

```
void List_Insert(list_t *L, int key) {
    // synchronization not needed
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        return;
    }
    new->key = key;

    // just lock critical section
    pthread_mutex_lock(&L->lock);
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
}
```



→ 삽입할 때만 mutex-lock을 해줘야 함.

```
int List_Lookup(list_t *L, int key) {
    int rv = -1;
    pthread_mutex_lock(&L->lock);
    node_t *curr = L->head;
    while (curr) {
        if (curr->key == key) {
            rv = 0;
            break;
        }
        curr = curr->next;
    }
    pthread_mutex_unlock(&L->lock);
    return rv; // now both success and failure
}
```

↓  
찾는 경우는 0


- Scalable linked list can be achieved by adding a lock per node

- This enables a high degree of concurrency in list operations
- However, it is hard to make it faster than the single lock approach

노드마다 lock

# Concurrent Queues

→ head와 tail 각각에 lock을 준다면  
동시에 pop, push를 할 수 있게.



- There is always a standard method to make a concurrent data structure: add a big lock

- For concurrent queues, we can achieve scalability by adding two locks, each for head and tail of the queue, enabling concurrency of enqueue and dequeue

```
typedef struct __node_t {
    int value;
    struct __node_t *next;
} node_t;

void Queue_Init(queue_t *q) {
    node_t *tmp = malloc(sizeof(node_t));
    tmp->next = NULL;
    q->head = q->tail = tmp;
    pthread_mutex_init(&q->head_lock, NULL);
    pthread_mutex_init(&q->tail_lock, NULL);
}

void Queue_Enqueue(queue_t *q, int value) {
    node_t *tmp = malloc(sizeof(node_t));
    assert(tmp != NULL);
    tmp->value = value;
    tmp->next = NULL;

    pthread_mutex_lock(&q->tail_lock);
    q->tail->next = tmp;
    q->tail = tmp;
    pthread_mutex_unlock(&q->tail_lock);
}
```

```
typedef struct __queue_t {
    node_t *head;
    node_t *tail;
    pthread_mutex_t head_lock, tail_lock;
} queue_t;
```

→ 앞쪽 데이터 읽을 때 lock 2개

```
int Queue_Dequeue(queue_t *q, int *value) {
    pthread_mutex_lock(&q->head_lock);
    node_t *tmp = q->head;
    node_t *new_head = tmp->next;
    if (new_head == NULL) {
        pthread_mutex_unlock(&q->head_lock);
        return -1; // queue was empty
    }
    *value = new_head->value;
    q->head = new_head;
    pthread_mutex_unlock(&q->head_lock);
    free(tmp);
    return 0;
}
```

원래 비어 있었지만 데이터가 들어오는 상황.

head가 데이터를 읽을 수 있도록 하는 "bounded queue"

# Concurrent Hash Table

- We finally consider a simple and widely applicable concurrent data structure, the **hash table** "제본 중독한"
- Let's focus on a simple hash table that does not resize
- This concurrent hash table is straightforward, is built using the concurrent lists we developed earlier
- The performance of this hash table is good "bucket이 다른 경우만!"
- Instead of a single lock for the entire structure, it uses a **lock per hash bucket** "키"

```
#define BUCKETS (101)

typedef struct __hash_t {
    list_t lists[BUCKETS];
} hash_t;

void Hash_Init(hash_t *H) {
    int i;
    for (i = 0; i < BUCKETS; i++)
        List_Init(&H->lists[i]);
}

int Hash_Insert(hash_t *H, int key) {
    return List_Insert(&H->lists[key % BUCKETS], key);
}

int Hash_Lookup(hash_t *H, int key) {
    return List_Lookup(&H->lists[key % BUCKETS], key);
}
```

bucket이 다른 동시에 가능.

/linked-list의 장점

그림 사용

