

Condition Variables

순서 지정과 평가 체계 등.



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Condition Variable: Concept

동기화 (synchronization)은 lock만으로 구현하기는 어렵다

- There are many cases where a thread wishes to check whether a condition is true before continuing its execution

- e.g.) a parent thread might wish to check whether a child thread has completed before continuing

```
void *child(void *arg) {
    printf("child\n");
    // XXX how to indicate we are done?
    return NULL;
}

int main(int argc, char *argv[]) {
    printf("parent: begin\n");
    pthread_t c;
    Pthread_create(&c, NULL, child, NULL); // create child
    // XXX how to wait for child?
    printf("parent: end\n");
    return 0;
}
```

What we want to see:

```
parent: begin
child
parent: end
```

최소화하여 바라는 뜻 (컴파일러도)

```
volatile int done = 0;

void *child(void *arg) {
    printf("child\n");
    done = 1;
    return NULL;
}

int main(int argc, char *argv[]) {
    printf("parent: begin\n");
    pthread_t c;
    Pthread_create(&c, NULL, child, NULL); // create child
    while (done == 0)
        ; // spin
    printf("parent: end\n");
    return 0;
}
```

그때가 왔네!
그때가 왔네!
done = 1
done = 0
done = 1
done = 1
HW 실행

spin

- We could try using a shared variable

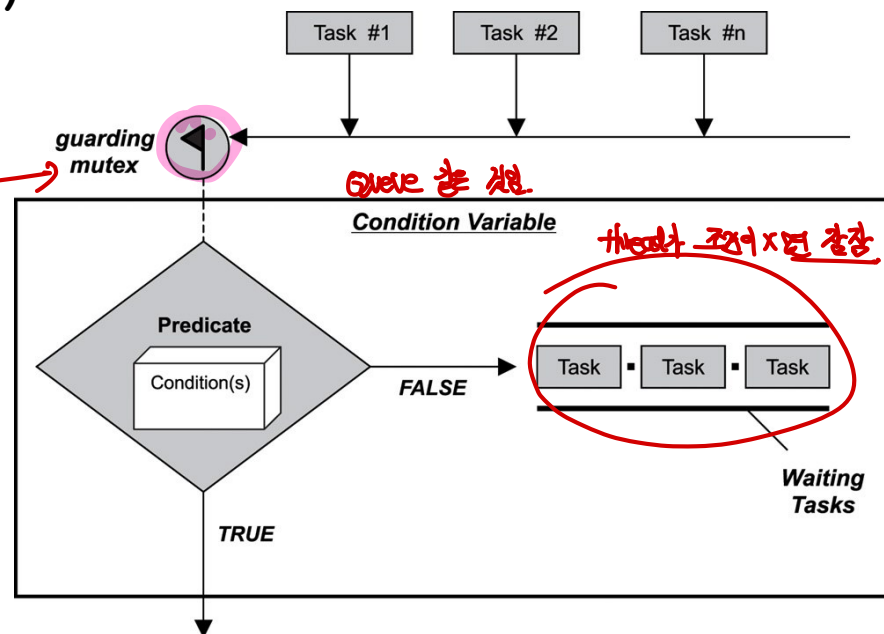
- This solution will generally work, but it is hugely inefficient as a parent spins, wasting CPU time

완전 상태
→ 모든 스레드를 sleep으로 깨워.

Definition and Routines

쓰레드 특정 조건이 만족되지 않았을 때
 큐 안에서 대기하게 만들어 준다
 → 다른 스레드에서 큐 안에서 대기 상태
 로 바뀐 스레드를 깨워. 다시
 Ready 상태가 됨

- To wait until condition is met, thread can use condition variable.
 - It is an explicit queue that threads can put themselves on when some state of execution (i.e. some condition) is not as desired by waiting on the condition
 - Some other thread, when it changes the state, can then wake one (or more) of the waiting threads and thus allow them to continue by signaling on the condition
- To use condition variable, you simply declare and initialize first
 - Two operations: `wait()` and `signal()`
 - `wait()` to put itself to sleep when a condition is not met
 - `signal()` to wake a sleeping thread waiting on this condition
 - `wait()` needs a mutex as parameter
 - 1) When calling, mutex must be locked
 - 2) Then, it releases the lock and put the thread to sleep atomically.
 - 3) When waking up, it must re-acquire the lock before returning



Parent Waiting for Child: Use a Condition Variable

Let's look at a simple example: two cases to consider

```
int done = 0;
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t c = PTHREAD_COND_INITIALIZER;
```

```
void thr_exit() {
  (4) Pthread_mutex_lock(&m);
  done = 1;
  Pthread_cond_signal(&c);
  Pthread_mutex_unlock(&m);
}
```

```
void *child(void *arg) {
  (3) printf("child\n");
  thr_exit();
  return NULL;
}
```

```
void thr_join() {
  (2) Pthread_mutex_lock(&m);
  while (done == 0)
    Pthread_cond_wait(&c, &m);
  Pthread_mutex_unlock(&m);
}
```

```
int main(int argc, char *argv[]) {
  printf("parent: begin\n");
  pthread_t p;
  Pthread_create(&p, NULL, child, NULL);
  (1) thr_join();
  (5) printf("parent: end\n");
  return 0;
}
```

1) The parent creates the child thread but continues running itself (parent → child):

- (1) The parent immediately calls `thr_join()` to wait for the child to complete
- (2) `thr_join()` acquires the lock and puts itself to sleep by calling `wait()`
- (3) The child runs, print "child", and call `thr_exit()` to wake the parent → child's context-switch.
- (4) `thr_exit()` acquires the lock, `done=1`, and signals to parent
- (5) the parent run, print "parent: end" and return

2) The child runs immediately upon creation (child → parent): → 이 순서 안 X

- (1) The child performs (3) and (4) above but nothing sleeping, and it just returns
- (2) The parent runs, calls `thr_joins()`, sees `done=1`, and doesn't wait and returns

Parent Waiting: No State Variable & No Lock

Why do we need the state variable done?

- Consider the case where the child runs immediately and calls `thr_exit()`
- The child signals but there is no thread asleep on the condition → When the parent runs, it will just call wait() and be stuck; no thread will ever wake it

No State Variable

```

void thr_exit() {
    Pthread_mutex_lock(&m);
    Pthread_cond_signal(&c);
    Pthread_mutex_unlock(&m);
}

void thr_join() {
    Pthread_mutex_lock(&m);
    Pthread_cond_wait(&c, &m);
    Pthread_mutex_unlock(&m);
}
  
```

child parent:
done이 없잖아?
부모 쓰러져 있는 거야

No Lock

```

void thr_exit() {
    done = 1;
    Pthread_cond_signal(&c);
}

void thr_join() {
    if (done == 0)
        Pthread_cond_wait(&c);
}
  
```

child 수행
parent가 서있
race condition 발생
atomic 하지 않아서 문제가 발생
시작에서 interrupt

What if threads does not hold a lock to signal and wait?

- The issue here is a subtle race condition
- Assume that the parent runs upon the child thread creation:
 - It calls `thr_join()`, sees `done=0`, is interrupted before sleep and the child runs
 - The child makes `done=1` and signals but no thread is waiting and nothing wakes
 - When the parent runs again, it sleeps forever

From these, we saw basic requirements for condition variable


~~The~~ Producer/Consumer Problem

→ bounded problem.

→ buffer 한정 환경

- The producer/consumer problem, is sometimes called bounded buffer problem, is a famous synchronization problem.

- Imagine one or more producer thread and one or more consumer threads
- Producers generate data items and place them in a buffer
- Consumers grab said items from the buffer and consume them in some way
- e.g.) multi-threaded web server (producer: http requests, consumer: processing requests), pipe (`grep foo file.txt` | `wc -l`) → connected through in-kernel bounded buffer



- The bounded buffer is a shared resource = 공유 자원 = 여러 스레드에서 접근 가능
= race-condition이 발생할 수 있음

- We must require synchronized access to it, otherwise, a race condition arise

Producer/Consumer Threads: Initial

- The first thing we need is a **shared buffer**, into which a producer puts data, and out of which a consumer takes data
 - Let's use a single integer for simplicity and we will generalize it later

→ int 1개 buffer.

```

int buffer;
int count = 0; // initially, empty

void put(int value) {
    assert(count == 0);
    count = 1;
    buffer = value;
}

int get() {
    assert(count == 1);
    count = 0;
    return buffer;
}

```

buffer
count=0/1

→ 0이 되면 (공백) 상태
producer가 data를 넣는 것
다한 것
다한 것

```

void *producer(void *arg) {
    int i;
    int loops = (int) arg;
    for (i = 0; i < loops; i++) {
        put(i);
    }
}

```

```

void *consumer(void *arg) {
    while (1) {
        int tmp = get();
        printf("%d\n", tmp);
    }
}

```

→ 상수 배열이 구현 되어 있지 않음.

← 계속 반복
get() 호출

- put () checks if the buffer is empty and puts a value into buffer and count=1
- get () does the opposite, setting the buffer to empty and returning the value
- We need two types of threads: **producer and consumer threads**
 - producer () puts an integer into the shared buffer loops times
 - consumer () gets the data out of the buffer (forever) and each time printing
 - However, this code does not work properly and an assertion will fire

2개의 thread.

→ sink가 많지 않음.

A Broken Solution



→ 여러개의 consumer들은
데이터 생각해 봐야 함.

Consider mutual exclusion (lock) and order (condition variable)

- Consider two consumers (T_{c1} , T_{c2}) and a producer (T_p): **1**) T_{c1} sleeps due to no data, **2**) T_p runs, produces a value, and sleeps due to buffer full, **3**) T_{c2} runs and consumes (then no data in buffer), **4**) T_{c1} runs and tries to consume but nothing
- Mesa semantics**: no guarantee that when the woken thread runs, the state will still be as desired (virtually every system ever built employs this semantics)
- Hoare semantics**: strong guarantee that the woken thread will run immediately

This code with a single producer and a single consumer works fine

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex); // p1
        if (count == 1) // p2
            Pthread_cond_wait(&cond, &mutex); // p3
        put(i); // p4
        Pthread_cond_signal(&cond); // p5
        Pthread_mutex_unlock(&mutex); // p6
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex); // c1
        if (count == 0) // c2
            Pthread_cond_wait(&cond, &mutex); // c3
        int tmp = get(); // c4
        Pthread_cond_signal(&cond); // c5
        Pthread_mutex_unlock(&mutex); // c6
        printf("%d\n", tmp);
    }
}
```

정리: T_{c1} 실행과 같이 어떤의 상태 변경이 있음. → 어떤 다른 실행도 어떤 용지점의

	T_{c1}	State	T_{c2}	State	T_p	State	Count	Comment
	c1	Run		Ready		Ready	0	
	c2	Run		Ready		Ready	0	
	c3	Sleep		Ready		Ready	0	Nothing to get
		Sleep		Ready	p1	Run	0	
		Sleep		Ready	p2	Run	0	
		Sleep		Ready	p4	Run	1	Buffer now full
		Ready		Ready	p5	Run	1	T_{c1} awoken
		Ready		Ready	p6	Run	1	
		Ready		Ready	p1	Run	1	
		Ready		Ready	p2	Run	1	
		Ready		Ready	p3	Sleep	1	Buffer full; sleep
		Ready	c1	Run		Sleep	1	T_{c2} sneaks in ...
		Ready	c2	Run		Sleep	1	
		Ready	c4	Run		Sleep	0	... and grabs data
		Ready	c5	Run		Ready	0	T_p awoken
		Ready	c6	Run		Ready	0	
	c4	Run		Ready		Ready	0	(?) Oh oh! No data

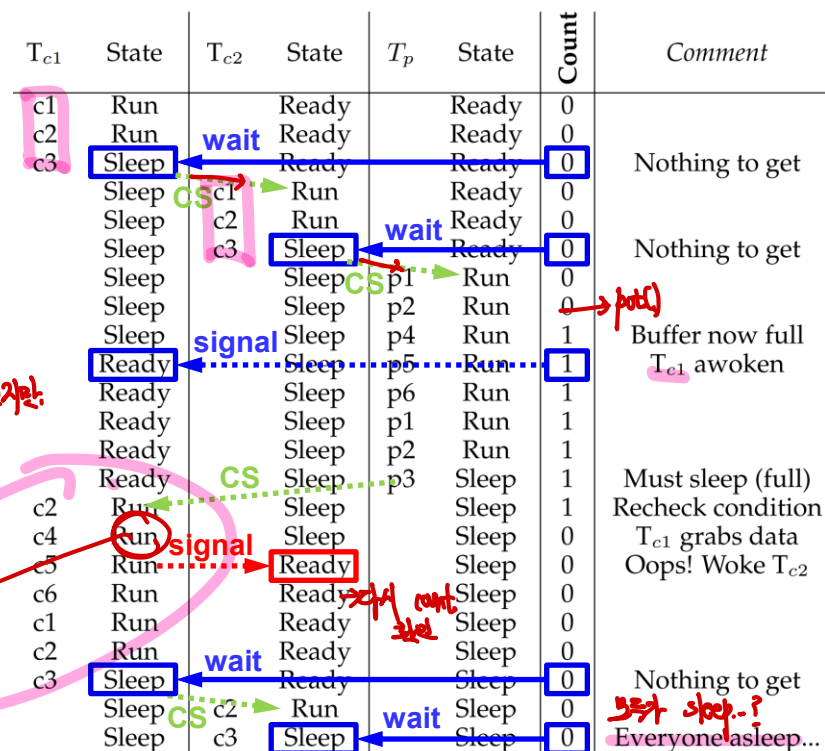
Better, But Still Broken: While, Not If

- The earlier issue is fixed by changing if to while for re-check**
 - However, a situation where **1)** T_{c1} , T_{c2} run, sleep (no data), **2)** T_p runs, produces a value, wakes T_{c1} , **3)** T_p sleeps (full), **4)** T_{c1} runs, re-checks, consumes, signals to T_{c2} (not T_p), sleeps (no data), **5)** T_{c2} runs, sleeps \rightarrow everyone asleep...
 - Signaling** is clearly needed, but must be **more directed**; a consumer should not wake other consumers, only producers, and vice-versa

Thanks to Mesa semantics, a rule to remember for condition variable is to **always use while loops**

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);          // p1
        while (count == 1)                  // p2
            Pthread_cond_wait(&cond, &mutex); // p3
        put(i);                             // p4
        Pthread_cond_signal(&cond);          // p5
        Pthread_mutex_unlock(&mutex);        // p6
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);          // c1
        while (count == 0)                  // c2
            Pthread_cond_wait(&cond, &mutex); // c3
        int tmp = get();                   // c4
        Pthread_cond_signal(&cond);          // c5
        Pthread_mutex_unlock(&mutex);        // c6
        printf("%d\n", tmp);
    }
}
```



The Single Buffer Producer/Consumer Solution

- The solution is a small one: use **two condition variables**, not one

- Then, threads can properly signal (which type of thread should be wake up) when the state of the system changes

```
cond_t empty, fill;
mutex_t mutex;
```

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 1)
            Pthread_cond_wait(&empty, &mutex);
        put(i);
        Pthread_cond_signal(&fill);
        Pthread_mutex_unlock(&mutex);
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 0)
            Pthread_cond_wait(&fill, &mutex);
        int tmp = get();
        Pthread_cond_signal(&empty);
        Pthread_mutex_unlock(&mutex);
        printf("%d\n", tmp);
    }
}
```

생산자는 소비자를 깨우고

소비자는 생산자만 깨우게 해야함

consumer가 consumer 버퍼를 깨워 줌

- The producer threads wait on the condition empty, and signals fill
- Conversely, the consumer threads wait on fill and signal empty
- By doing so, a consumer can never accidentally wake a consumer, and a producer can never accidentally wake a producer.

- Now, we know that the use of **two condition variables** and **while** is necessary to properly handle the produce/consumer problem


The Correct Producer/Consumer Solution

- The last change is to enable more concurrency and efficiency
 - We add more buffer slots, so that multiple values can be produced before sleeping, and similarly multiple values can be consumed before sleeping
 - The first change for this correct solution is within the buffer structure itself and the corresponding `put()` and `get()` → circular buffer (queue)

```
int buffer[MAX];
int fill_ptr = 0;
int use_ptr = 0;
int count = 0;

void put(int value) {
    buffer[fill_ptr] = value;
    fill_ptr = (fill_ptr + 1) % MAX;
    count++;
}
```

```
int get() {
    int tmp = buffer[use_ptr];
    use_ptr = (use_ptr + 1) % MAX;
    count--;
    return tmp;
}
```



- We also slightly change the conditions that producers check to determine whether to sleep or not


```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex); // p1
        while (count == MAX) // p2
            Pthread_cond_wait(&empty, &mutex); // p3
        put(i); // p4
        Pthread_cond_signal(&fill); // p5
        Pthread_mutex_unlock(&mutex); // p6
    }
}
```

buffer가 꽉차면 못 넣음

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex); // c1
        while (count == 0) // c2
            Pthread_cond_wait(&fill, &mutex); // c3
        int tmp = get(); // c4
        Pthread_cond_signal(&empty); // c5
        Pthread_mutex_unlock(&mutex); // c6
        printf("%d\n", tmp);
    }
}
```

꽉차 있으면 producer 멈춤

Covering Conditions



~

Consider a simple multi-threaded memory allocation library

- When a thread calls into the memory allocation code, it might have to wait for more memory to become free
- Then, which of waiting threads should be woken up when a thread frees memory?
- A scenario that when there is zero byte free, 1) T_a calls `allocate(100)`, T_b calls `allocate(10)`, and both wait on the condition and sleep, 2) T_c calls `free(50)` and what if it signals to T_a instead of T_b ?
→ T_a sleeps again

```
// how many bytes of the heap are free?
int bytesLeft = MAX_HEAP_SIZE;
```

```
// need lock and condition too
cond_t c;
mutex_t m;
```

```
void *
allocate(int size) {
    Pthread_mutex_lock(&m);
    while (bytesLeft < size)
        Pthread_cond_wait(&c, &m);
    void *ptr = ...; // get mem from heap
    bytesLeft -= size;
    Pthread_mutex_unlock(&m);
    return ptr;
```

```
void free(void *ptr, int size) {
    Pthread_mutex_lock(&m);
    bytesLeft += size;
    Pthread_cond_signal(&c); // whom to signal??
    Pthread_mutex_unlock(&m);
```

- The solution is simple: replace `signal()` to `broadcast()`, waking up all
- This scheme guarantees that any threads that should be woken are
- However, the threads that should not be awake will also wake up, re-check the condition, and then go immediately back to sleep → performance ↓ context-switch., condition check.
- Covering condition: covers all the cases where a thread needs to wake up