

Semaphores



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

Semaphores: A Definition

→ OS concurrency 의 핵심.

상호 배제 + 동기화 (한번에!)

- A semaphore invented by Dijkstra is an object with an integer value that we can manipulate with two routines

- In the POSIX standard, these routines are `sem_wait()` and `sem_post()`
- The semaphore must be initialized since its initial value determines the behavior

```
#include <semaphore.h>
sem_t s;
sem_init(&s, 0, 1);
```

semaphore를 초기

→ 초기화 값이 1이다. → lock도도 쓰면 안 되고
condition variable은 쓰면 안 됨

This code declares a semaphore s and initialize it to the value 1
0 indicates a shared semaphore between threads in the same process

- After a semaphore is initialized, we can call one of two functions to interact with it, `sem_wait()` or `sem_post()`

```
int sem_wait(sem_t *s) {
    decrement the value of semaphore s by one
    wait if value of semaphore s is negative
}

int sem_post(sem_t *s) {
    increment the value of semaphore s by one
    if there are one or more threads waiting, wake one
}
```

- `sem_wait()` either returns right away (`s--`) or causes the caller to suspend execution waiting (if `s < 0`) for a subsequent post (`s++`)
- `sem_post()` simply increases the value of the semaphore (`s++`) and then, wakes one of them up if there is any waiting thread
- The semaphore value, when negative, is equal to the number of waiting threads

Binary Semaphores (Locks)

Consider a code that tries to use a semaphore as a lock

- What should the initial value of the semaphore X be? By definition, $X=1$

Consider scenario with two threads

- T_0 calls `sem_wait()` ($m=0$), enters critical section, calls `sem_post()` ($m=1$)
- 1) T_0 calls `sem_wait()` ($m=0$), begins critical section, 2) T_1 runs, tries critical section, calls `sem_wait()` ($m=-1$), sleeps, 3) T_0 runs, ends critical section, calls `sem_post()` ($m=0$), wakes, 4) T_1 runs, critical section, calls `sem_post()` ($m=1$)

```
sem_t m;
sem_init(&m, 0, X); // initialize to X; what should X be?
```

```
sem_wait(&m);
// critical section here
sem_post(&m);
```

이것은 lock 기능
= binary semaphore

Value of Semaphore	Thread 0	Thread 1
1		
1	call sem_wait()	
0	sem_wait() returns	
0	(crit sect)	
0	call sem_post()	
1	sem_post() returns	

sleeping 상태
기다림

Val	Thread 0	State	Thread 1	State
1		Run		Ready
1	call sem_wait()	Run		Ready
0	sem_wait() returns	Run		Ready
0	(crit sect begin)	Run		Ready
0	Interrupt; Switch $\rightarrow T1$	Ready		Run
0		Ready	call sem_wait()	Run
-1		Ready	decr sem	Run
-1		Ready	(sem < 0) \rightarrow sleep	Sleep
-1		Run	Switch $\rightarrow T0$	Sleep
-1	(crit sect end)	Run		Sleep
-1	call sem_post()	Run		Sleep
0	incr sem	Run		Sleep
0	wake (T1)	Run		Ready
0	sem_post() returns	Run		Ready
0	Interrupt; Switch $\rightarrow T1$	Ready		Run
0		Ready	sem_wait() returns	Run
0		Ready	(crit sect)	Run
0		Ready	call sem_post()	Run
1		Ready	sem_post() returns	Run

sleep

이제 switching 가능함.

Semaphores For Ordering → initial value 0 (24 18)

- Semaphores are useful to **order events** in a concurrent program
 - What should the **initial value** of this semaphore be? It should be set to **0**
- There are two cases to consider for the parent-child program
 - Parent→Child: **1)** parent calls `sem_wait()` ($s=-1$), sleeps, **2)** child runs, calls `sem_post()` ($s=0$), wakes (P), **3)** parent runs, returns from `sem_wait()`
 - Child→Parent: **1)** child runs, calls `sem_post()` ($s=1$), wakes (nothing), **2)** parent runs, calls `sem_wait()` ($s=0$) and `sem_wait()` returns immediately due to $s \geq 0$

```
sem_t s;

void *child(void *arg) {
    printf("child\n");
    sem_post(&s); // signal here: child is done
    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    sem_init(&s, 0, X); // what should X be?
    printf("parent: begin\n");
    pthread_t c;
    pthread_create(&c, NULL, child, NULL);
    sem_wait(&s); // wait here for child
    printf("parent: end\n");
    return 0;
}
```

개별 스레드 실행.

Val	Parent	State	Child	State
0	create (Child)	Run	(Child exists, can run)	Ready
0	call sem_wait()	Run		Ready
-1	decr sem	Run		Ready
-1	(sem<0)→sleep	Sleep		Ready
-1	Switch→Child	Sleep	child runs	Run
-1		Sleep	call sem_post()	Run
0		Sleep	inc sem	Run
0		Ready	wake (Parent)	Run
0		Ready	sem_post() returns	Run
0		Ready	Interrupt→Parent	Ready
0	sem_wait() returns	Run		Ready

Val	Parent	State	Child	State
0	create (Child)	Run	(Child exists; can run)	Ready
0	Interrupt→Child	Ready	child runs	Run
0		Ready	call sem_post()	Run
1		Ready	inc sem	Run
1		Ready	wake (nobody)	Run
1		Ready	sem_post() returns	Run
1	parent runs	Run	Interrupt→Parent	Ready
1	call sem_wait()	Run		Ready
0	decrement sem	Run		Ready
0	(sem≥0)→awake	Run		Ready
0	sem_wait() returns	Run		Ready

The Producer/Consumer Problem: First Attempt

■ The first attempt introduces two semaphores, empty and full

→ 이거 세마포어로 잘 되나?

- Imagine there is two threads, a producer and a consumer, with $MAX=1$
- 1) consumer runs, calls `sem_wait(&full)` ($f=-1$), sleeps, 2) producer runs, calls `sem_wait(&empty)` ($e=0$), `put()`, `sem_post(&full)` ($f=0$), wakes: Then, 3) if the producer continues, calls `sem_wait(&empty)` ($e=-1$), sleeps or 3) if the consumer runs, returns from `sem_wait(&full)`, consumes

– It works with two or more threads, but what if multiple threads with $MAX > 1$?

- Imagine two producers (P_a, P_b) call `put()`: 1) P_a puts a value to `buffer[0]` and interrupted before `fill++`, 2) P_b puts a value to `buffer[0]`?...NO!

semaphore 2개

```
int buffer[MAX];
int fill = 0;
int use = 0;

void put(int value) {
    buffer[fill] = value; // Line F1
    fill = (fill + 1) % MAX; // Line F2
}

int get() {
    int tmp = buffer[use]; // Line G1
    use = (use + 1) % MAX; // Line G2
    return tmp;
}
```

또 produce 하도.

sleep 하지 않을 수 있다.

→ Mutual exclusion 이
필요하다.

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&empty); // Line P1
        put(i); // Line P2
        sem_post(&full); // Line P3
    }
}
```

buffer의 크기가 1인 상황.

Initialization

```
sem_init(&empty, 0, MAX);
sem_init(&full, 0, 0);
```

```
void *consumer(void *arg) {
    int tmp = 0;
    while (tmp != -1) {
        sem_wait(&full); // Line C1
        tmp = get(); // Line C2
        sem_post(&empty); // Line C3
        printf("%d\n", tmp);
    }
}
```

정신차
이름

만약
공유하면

A Solution: Adding Mutual Exclusion

dead-lock ...
↑
binary semaphore

■ Filling a buffer and increasing the index is a critical section

- We can use the binary semaphore to add some locks (left code)
- Unfortunately, it does not work due to deadlock; imagine a producer, a consumer:
 1) consumer calls `sem_wait(&mutex)` ($m=0$), `sem_wait(&full)` ($f=-1$), sleeps; holds the lock, 2) producer runs, calls `sem_wait(&mutex)` ($m=-1$), sleeps; stuck
 → consumer holds lock and waiting for signal full and producer is waiting for lock
- The solution is to reduce the scope of the lock (see right code)

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex); // Line P0 (NEW LINE)
        sem_wait(&empty); // Line P1
        put(i);           // Line P2
        sem_post(&full);  // Line P3
        sem_post(&mutex); // Line P4 (NEW LINE)
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex); // Line C0 (NEW LINE)
        sem_wait(&full);  // Line C1
        int tmp = get();  // Line C2
        sem_post(&empty); // Line C3
        sem_post(&mutex); // Line C4 (NEW LINE)
        printf("%d\n", tmp);
    }
}
```

mutex 1 → 0 → producer sleep
full 0 → -1 → consumer sleep
empty 1

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&empty); // Line P1
        sem_wait(&mutex); // Line P1.5 (MUTEX HERE)
        put(i);           // Line P2
        sem_post(&mutex); // Line P2.5 (AND HERE)
        sem_post(&full);  // Line P3
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&full);  // Line C1
        sem_wait(&mutex); // Line C1.5 (MUTEX HERE)
        int tmp = get();  // Line C2
        sem_post(&mutex); // Line C2.5 (AND HERE)
        sem_post(&empty); // Line C3
        printf("%d\n", tmp);
    }
}
```

dead-lock x.
lock의 범위를 줄여주
mutex 1 → 0 → 1
full 0 → 1 → sleep
empty 1 → 0

Reader-Writer Locks



Read-writer lock allows many concurrent reads and single write

- A single writer can acquire the lock (**writelock**) to update the data structure
- When acquiring a **read lock**, the reader first acquires **lock** and **increase the reader variable** to track how many reads are currently inside the data structure
- When the **first reader** acquires the lock, the reader **also acquires the write lock (**writelock**)**, and then will release the lock later after the last reader is done
→ 읽는 동안에 쓰면 안되니까.
- Once a reader has **lock**, more readers can acquire, but any thread to write has to wait until finishing all readers → relatively **easy for readers to starve writers**

```
typedef struct _rwlock_t {
    sem_t lock;        // binary semaphore (basic lock)
    sem_t writelock;   // allow ONE writer/MANY readers
    int  readers;      // #readers in critical section
} rwlock_t;
```

```
void rwlock_init(rwlock_t *rw) {
    rw->readers = 0;
    sem_init(&rw->lock, 0, 1);
    sem_init(&rw->writelock, 0, 1);
}
```

```
void rwlock_acquire_readlock(rwlock_t *rw) {
    sem_wait(&rw->lock); → read 할때 lock 받.
    rw->readers++;
    if (rw->readers == 1) // first reader gets writelock
        sem_wait(&rw->writelock);
    sem_post(&rw->lock);
}
```

→ 첫번째 reader가 lock 받음

```
void rwlock_release_readlock(rwlock_t *rw) {
    sem_wait(&rw->lock);
    rw->readers--;
    if (rw->readers == 0) // last reader lets it go
        sem_post(&rw->writelock);
    sem_post(&rw->lock);
}
```

```
void rwlock_acquire_writelock(rwlock_t *rw) {
    sem_wait(&rw->writelock);
}
```

```
void rwlock_release_writelock(rwlock_t *rw) {
    sem_post(&rw->writelock);
}
```

starvation.

writer는 계속 기다릴 수 밖에 없음 (reader 우선)

The Dining Philosophers: Introduction


손 안 잡지마!

■ Dining philosopher's problem is a famous concurrency problem

- It is fun and intellectually interesting; however, its practical utility is low
- Assume there are five “philosophers” on a table and between each pair of philosophers is a single fork (and thus, five forks total)
- The philosophers each have times where they **think**, and don't need any forks, and where they **eat**, and a philosopher needs two forks to eat (left and right)
- Consider the basic loop of each philosopher; the key is to write `get_forks()` and `put_forks()` to attain **no deadlock**, **no starvation**, and **high concurrency**.

2개의 포크 필요

→ 어떻게 해볼까?



```
while (1) {
    think();
    get_forks(p);
    eat();
    put_forks(p);
}
```

Each philosopher has a unique thread identifier `p` from 0 to 4 (inclusive)

The Dining Philosophers: Solutions

■ We need some semaphores to solve the problem

- Let's assume we have **five semaphores**, one for each fork 5개 semaphore
- Two helper functions to get left and right forks when philosopher p wants to eat

```
int left(int p) { return p; }
int right(int p) { return (p + 1) % 5; }
```

왼쪽, 오른쪽 포크 번호를 반환하는?

■ Broken solution: deadlock 철학자 번호

- To pick the forks, we simply grab a lock on each; one on the left and then one on the right and release them when done eating → deadlock

평 무 포크를 얻지 못하고 기다릴 때

Deadlock

Each will be stuck holding one left fork and waiting for another, forever

dead lock

무한 대기 발생함

```
void get_forks(int p) {
    sem_wait(&forks[left(p)]);
    sem_wait(&forks[right(p)]);
}

void put_forks(int p) {
    sem_post(&forks[left(p)]);
    sem_post(&forks[right(p)]);
}
```

왼쪽, 오른쪽 포크를 얻음

포크를 놓음

```
void get_forks(int p) {
    if (p == 4) {
        sem_wait(&forks[right(p)]);
        sem_wait(&forks[left(p)]);
    } else {
        sem_wait(&forks[left(p)]);
        sem_wait(&forks[right(p)]);
    }
}
```

한명은 오른쪽 포크를 먼저 잡는다.

마지막은 오른쪽

No deadlock

■ A solution: breaking the dependency

- The simplest way to attack this problem is to **change how forks are acquired**
- Assume that philosopher 4 gets the forks in a different order than the others
- There is no situation where each philosopher grabs one fork and is stuck waiting for another → the **cycle of waiting** is broken

How to Implement Semaphores

Building our own semaphores using lock and condition variable

- One lock and one condition variable, plus a state variable to track the value of the semaphore are used to implement a semaphore
- What the subtle difference between ours and Dijkstra's is that the negative value of the semaphore doesn't indicate the number of waiting threads
- The value will never be lower than zero → current Linux implementation

lock + condition variable + state variable

```
typedef struct __Zem_t {
    int value;           → state variable.
    pthread_cond_t cond;
    pthread_mutex_t lock;
} Zem_t;
```

```
// only one thread can call this
void Zem_init(Zem_t *s, int value) {
    s->value = value;
    Cond_init(&s->cond);
    Mutex_init(&s->lock);
}
```

```
void Zem_wait(Zem_t *s) {
    Mutex_lock(&s->lock);
    while (s->value <= 0) → negative value
        Cond_wait(&s->cond, &s->lock);
    s->value--;
    Mutex_unlock(&s->lock);
}
```

```
void Zem_post(Zem_t *s) {
    Mutex_lock(&s->lock);
    s->value++;
    Cond_signal(&s->cond);
    Mutex_unlock(&s->lock);
}
```

Too many threads are running and bogging the system down?

- Then, decide upon a threshold for "too many", and then use a semaphore to limit the number of threads concurrently executing the piece of code
- This approach is called throttling, is a form of admission control

Summary

- A semaphore is an object with an integer value that is manipulated by two routines, `sem_wait()` and `sem_post()`
 - The semaphore must be initialized since its initial value determines the behavior
- A semaphore is like a bouncer at the front of the lineup
 - The bouncer only allows threads to proceed when instructed to do so

