

I/O Devices

→ 커널, 루틴을 포함한 모든 임베디드 시스템

↓

이름 OS에서 사용할 수 있도록 하는

Device Driver

Linux 90%



Prof. Yongtae Kim

Computer Science and Engineering
Kyungpook National University

System Architecture HW 72.

■ We introduce the concept of an input/output (I/O) device

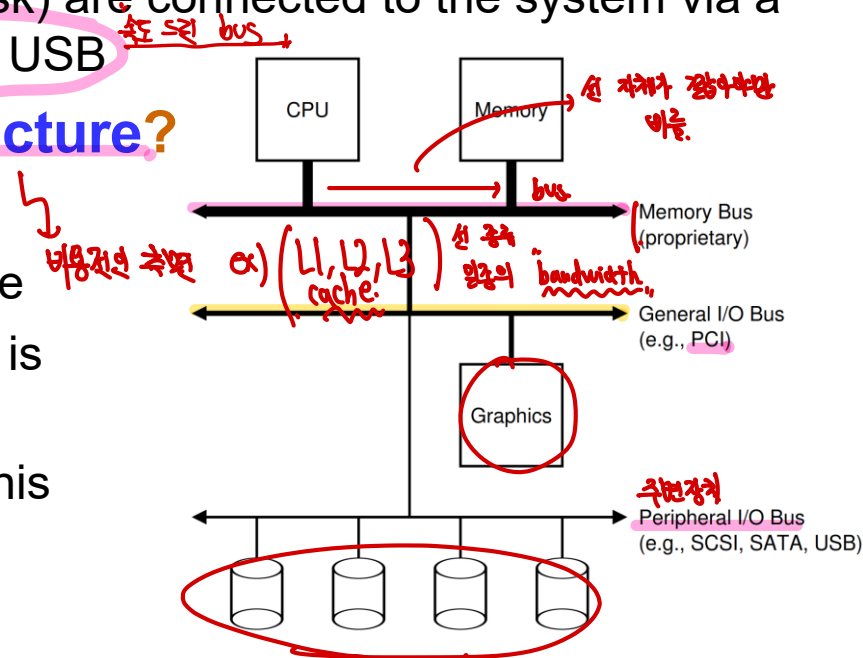
- I/O is quite critical to computer systems; if there is no keyboard and monitor?

■ Let's look at a classical diagram of a typical system


- A single CPU attached to the main memory via memory bus or interconnect
- Some devices are connected to the system via a general I/O bus (e.g. PCI); graphics and some other higher-performance I/O devices might be found here
- Slow devices (e.g. keyboard, mouse, disk) are connected to the system via a peripheral bus, such as SCSI, SATA, or USB

■ Why do we need hierarchical structure?

- The answer is simple: physics and cost
- The faster a bus is, the shorter it must be
- Engineering a bus for high performance is quite costly
- Thus, system designers have adopted this hierarchical approach

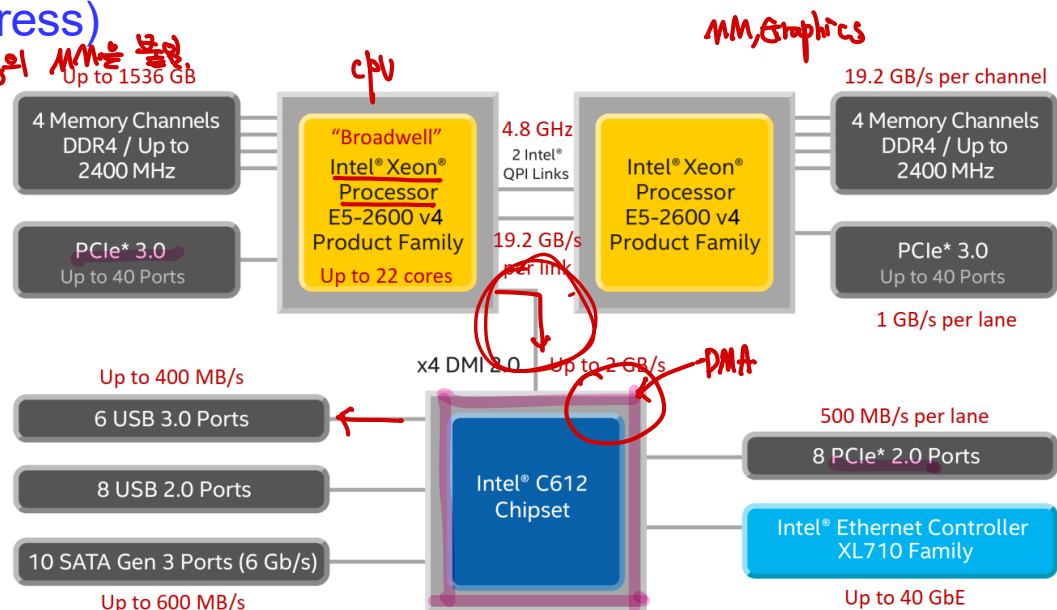
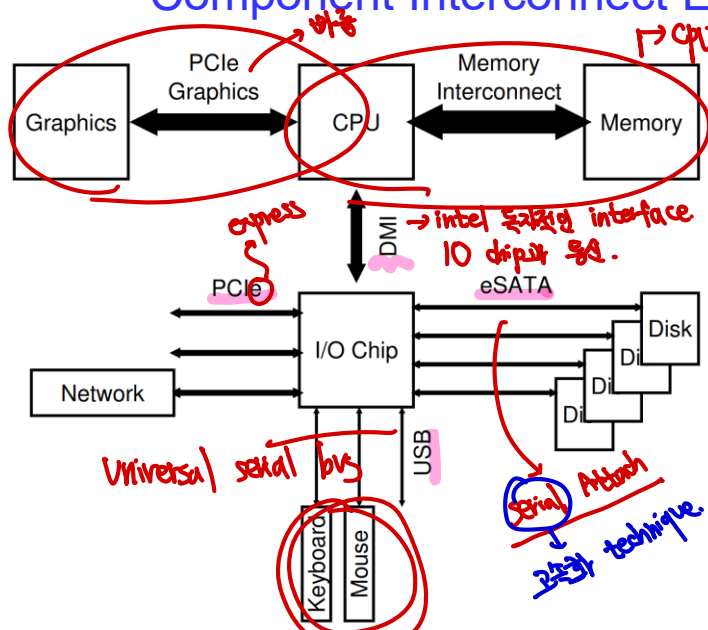


Modern System Architecture



- Modern systems increasingly use specialized chipsets and faster point-to-point interconnects to improve performance

- A CPU connects to an I/O chip via Intel's proprietary DMI (Direct Media Interface) and the rest of the devices connect to this chip via various different interconnects
- One or more hard drives connect to the system via the eSATA interface
- The I/O chip includes a number of USB connections for low performance devices
- Other higher performance devices can be connected via PCIe (Peripheral Component Interconnect Express)



Platform Controller Hub (PCH)

Courtesy of Prof. Jin-Soo Kim @ SNU

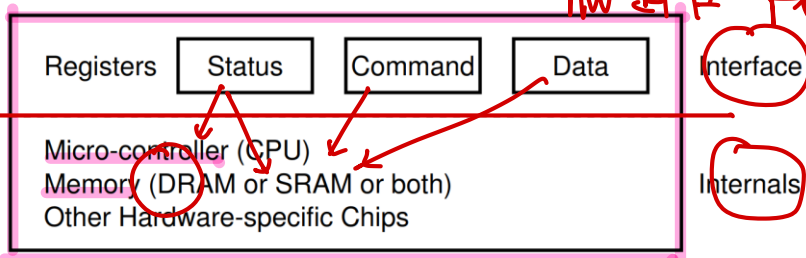
A Canonical Device & Protocol

Device는 많은 종류가 있어 어떤 수 있죠. 커널. 하는 것은 장치 시스템이 제공하는 작업은 추상화

A device has two important component

HW가 어떤 상태인지. (busy/idle,...)

- 1) The **hardware interface** allows the system software to control its operation; this canonical device's hardware interface is comprised of **three registers**:
 - (1) **status register**, which can read to see the current status of the device
 - (2) **command register**, to tell the device to perform a certain task
 - (3) **data register**, to pass data to the device, or get data from device
- 2) Its **internal structure** is implementation specific and responsible for implementing the abstraction the device that presents to the system (e.g. chips, firmware)



```

1) While (STATUS == BUSY)
    ; // wait until device is not busy
2) Write data to DATA register
3) Write command to COMMAND register
   (starts the device and executes the command)
4) While (STATUS == BUSY)
    ; // wait until device is done with your request
   // device의 처리 끝난 때까지 기다림.

```

이 Data를 "send"

The protocol to interact between device and OS has four steps



- 1) OS keeps reading status register until device is ready to get command (**polling**)
→ device의 status를 계속 polling
- 2) OS sends some data down to the data register (**programmed I/O; PIO**)
→ CPU가 높은 것은 → interrupt!
- 3) OS writes a command to the command register to tell the device what to do
spin/ack과 관련된 송/수신 처리가 관여
- 4) OS waits for the device to finish by again polling it in a loop

Lowering CPU Overhead with Interrupts

- The interrupt can improve the OS and device interaction**

인간적으로 CPU가 쉬고 있는 동안 장치를 위해서 불필요하게!
→ but 바쁜 것을 spin - wait를 하게 하면 나쁜 것

 - Instead of polling the device repeatedly, OS can issue a request, put the calling process to sleep, and context switch to another task.
 - When the device is finished with the operation, it will raise a hardware interrupt, causing CPU to jump into OS at a predetermined **ISR** or interrupt handler.
- Interrupts allow for overlap of computation and I/O**
 - Without interrupts: 1) the system simply spins, polling the status of the device repeatedly until the I/O is complete, 2) P_1 can run again
 - With interrupts: 1) OS runs P_2 on CPU while the disk services P_1 's request, 2) an interrupt occurs when I/O finished, 3) OS wakes up P_1 and run it again




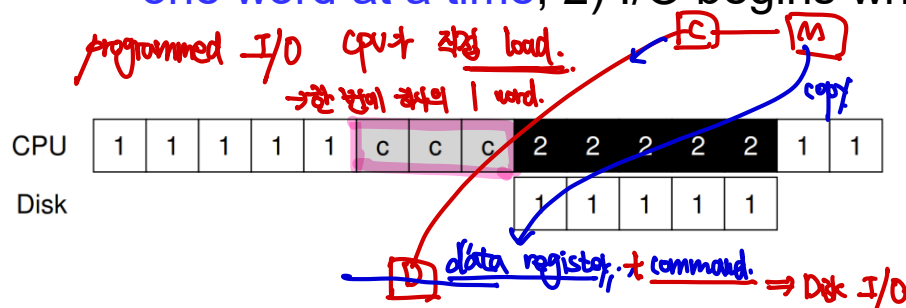
- Using interrupts is not always the best solution**
 - Lots of interrupts will slow down the system: context switch, interrupt handling
 - The polling may be the best if a device is fast, otherwise, interrupt → hybrid
 - Livelock issue due to huge interrupts → polling or interrupt coalescing (interrupt 합치기)

More Efficient Data Movement with DMA

CPU는 바쁘는데, 어쩔 수 없이 할 거를 더 많이 한다.

When using PIO to transfer a large chunk of data to a device:

- CPU is once again overburdened with a rather trivial task, and thus **wastes a lot of CPU time and effort** that could better be spent running other processes
- PIO: 1) P_1 initiates I/O, which must **copy** data from memory to device explicitly, **one word at a time**, 2) I/O begins when the copy is complete, 3) P_1 runs again



The DMA allows us to utilize the CPU more effectively

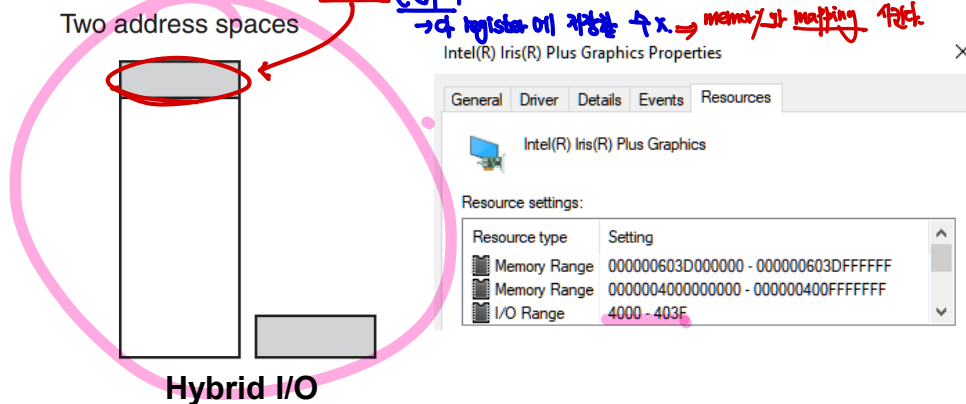


→ CPU의 제법 중요한 일과 더불어 다른 일을 할 수 있는 시스템의 장치.

- A **DMA engine** is essentially a specific device in system that can orchestrate transfers between devices and main memory **without much CPU intervention**
- DMA: 1) OS (P_1) programs the DMA engine by **telling it the transfer information** (i.e. **source and target address of data, size of data, etc**),
2) OS is done with the transfer and can proceed with other work P_2
3) **DMA controller** raises an **interrupt** upon the DMA transfer,
4) P_1 can run again

Method of Device Interaction



- The oldest method is to have ^①explicit I/O instructions**
 - This is used by IBM mainframes for many years and these instructions specify a way for OS to send data to specific device registers (e.g. in, out on x86) *명시적인 I/O 인스트럭션.*
 - Such instructions are usually privileged and OS is the only entity to use them *특권 있는 코드 (OS만 가능)*
- The second one to interact with devices is ^②memory-mapped I/O.**
 - This makes device registers available *장치 레지스터를 실제 메모리처럼* as if they were memory locations
 - To access a register, OS issues a load (to read) or store (to write) the address *load/store 인스트럭션을 장치 어드레스와 함께*
- A hybrid I/O scheme is also available**
 - This scheme uses memory-mapped I/O data buffers and separate I/O ports for the control registers (e.g. graphics controller) *Status command. data (GPU) → data register에 저장할 수 x → memory와 mapping 생김.*



Fitting into the OS: The Device Driver

- How to fit devices, each of which have very specific interfaces, into the OS, which we want to keep as general as possible?
 - **Abstraction**: a **device driver** encapsulates any specifics of device interaction
- Let's examine the Linux file system software stack
 - A file system (and certainly, an application above) is completely oblivious to the specifics of which disk class it is using
 - It simply issues block read and write requests to the generic block layer
 - Then, the device driver handles the details of issuing the specific request
 - This also provides a **raw interface** to devices, which enables special application (e.g. disk defragmentation) to directly read and write without the abstraction

