

CH 1 Introduction

The Computer Revolution

- Unprecedented speed of progress and innovations in computer systems
 - Analogy to transportation: NY to London in seconds for a penny
- The third revolution of human civilization is made possible
- Enabling new applications
 - Improvement in the cost of computing by factor of 10 → Numerous applications and opportunities appear



Computers in Automobiles



Cell phones



What else?



World wide web



Search engine



Human genome project

(Traditional) Class of Computers



Personal Computers

- General purpose, variety of software
- Subject to cost/performance tradeoff




Supercomputers

- High-end scientific and engineering calculations
- Highest capability and performance but represent a small fraction of the overall computer market



Servers

- Network based
- High capacity, performance, reliability



Embedded computers

- Hidden as components of systems
- Stringent power/performance/cost constraints

Supercomputer 2020

■ Top500 from Wikipedia

Rank ↕	Rmax Rpeak (PFLOPS)	Name ↕	Model ↕	Processor ↕	Interconnect	Vendor ↕	Site country, year	Operating system ↕
1 ▲	415.530 513.855	Fugaku	Supercomputer Fugaku	A64FX	Tofu interconnect D	Fujitsu	RIKEN Center for Computational Science ● Japan, 2020	Linux (RHEL)
2 ▼	148.600 200.795	Summit	IBM Power System AC922	POWER9, Tesla V100	InfiniBand EDR	IBM	Oak Ridge National Laboratory United States, 2018	Linux (RHEL)
3 ▼	94.640 125.712	Sierra	IBM Power System S922LC	POWER9, Tesla V100	InfiniBand EDR	IBM	Lawrence Livermore National Laboratory United States, 2018	Linux (RHEL)
4 ▼	93.015 125.436	Sunway TaihuLight	Sunway MPP	SW26010	Sunway ^[25]	NRCPC	National Supercomputing Center in Wuxi China, 2016 ^[25]	Linux (Raise)
5 ▼	61.445 100.679	Tianhe-2A	TH-IVB-FEP	Xeon E5-2692 v2, Matrix-2000 ^[26]	TH Express-2	NUDT	National Supercomputing Center in Guangzhou China, 2013	Linux (Kylin)
6 ▲	35.450 51.721	HPC5	Dell	Xeon Gold 6252, Tesla V100	Mellanox HDR Infiniband	Dell EMC	Eni Italy, 2020	Linux (CentOS)
7 ▲	27.580 34.569	Selene	Nvidia	Epyc 7742, Ampere A100	Mellanox HDR Infiniband	Nvidia	Nvidia United States, 2020	Linux (Ubuntu)
8 ▼	23.516 38.746	Frontera	Dell C6420	Xeon Platinum 8280 (subsystems with e.g. POWER9 CPUs and Nvidia GPUs were added after official benchmarking ^[10])	InfiniBand HDR	Dell EMC	Texas Advanced Computing Center United States, 2019	Linux (CentOS)
9 ▲	21.640 29.354	Marconi- 100	IBM Power System AC922	POWER9, Volta V100	Dual-rail Mellanox EDR Infiniband	IBM	CINECA Italy, 2020	Linux (RHEL)
10 ▼	21.230 27.154	Piz Daint	Cray XC50	Xeon E5-2690 v3, Tesla P100	Aries	Cray	Swiss National Supercomputing Centre Switzerland, 2016	Linux (CLE)

Supercomputer 2021

■ Top500 from Wikipedia

Rank (previous)	Rmax Rpeak (PFLOPS)	Name	Model	CPU cores	Accelerator (e.g. GPU) cores	Interconnect	Manufacturer	Site country	Year	Operating system
1	442.010 537.212	Fugaku	Supercomputer Fugaku	158,976 × 48 A64FX @2.2 GHz	0	Tofu interconnect D	Fujitsu	RIKEN Center for Computational Science 🇯🇵 Japan	2020	Linux (RHEL)
2▼ (1)	148.600 200.795	Summit	IBM Power System AC922	9,216 × 22 POWER9 @3.07 GHz	27,648 × 80 Tesla V100	InfiniBand EDR	IBM	Oak Ridge National Laboratory 🇺🇸 United States	2018	Linux (RHEL)
3▼ (2)	94.640 125.712	Sierra	IBM Power System S922LC	8,640 × 22 POWER9 @3.1 GHz	17,280 × 80 Tesla V100	InfiniBand EDR	IBM	Lawrence Livermore National Laboratory 🇺🇸 United States	2018	Linux (RHEL)
4▼ (3)	93.015 125.436	Sunway TaihuLight	Sunway MPP	40,960 × 260 SW26010 @1.45 GHz	0	Sunway ^[34]	NRCPC	National Supercomputing Center in Wuxi 🇨🇳 China ^[34]	2016	Linux (Raise)
5▲ (new)	64.590 89.795	Perlmutter	HP	? × 64 Epyc 7763 @2.45 GHz	? × 108 Ampere A100	Slingshot-10	Nvidia	NERSC 🇺🇸 United States	2021	Linux (Cray Linux Environment)
6▼ (5)	63.460 79.215	Selene	Nvidia	1,120 × 64 Epyc 7742 @2.25 GHz	4,480 × 108 Ampere A100	Mellanox HDR Infiniband	Nvidia	Nvidia 🇺🇸 United States	2020	Linux (Ubuntu)
7▼ (6)	61.445 100.679	Tianhe-2A	TH-IVB-FEP	35,584 × 12 Xeon E5-2692 v2 @2.2 GHz	35,584 × 128 Matrix- 2000 ^[35]	TH Express-2	NUDT	National Supercomputer Center in Guangzhou 🇨🇳 China	2013	Linux (Kylin)
8▼ (7)	44.120 70.980	JUWELS (booster module) ^{[36][37]}	BullSequana XH2000	1,872 × 24 AMD Epyc 7402 @2.8 GHz	3,744 × 108 Ampere A100	Mellanox HDR Infiniband	Atos	Forschungszentrum Jülich 🇩🇪 Germany	2020	Linux (CentOS)
9▼ (8)	35.450 51.721	HPC5	Dell	3,640 × 24 Xeon Gold 6252 @2.1 GHz	7,280 × 80 Tesla V100	Mellanox HDR Infiniband	Dell EMC	Eni 🇮🇹 Italy	2020	Linux (CentOS)
10▼ (9)	23.516 38.746	Frontera	Dell C6420	16,016 × 28 Xeon Platinum 8280 @2.7 GHz (subsystems with e.g. POWER9 CPUs and Nvidia GPUs were added after official benchmarking ^[10])	0	InfiniBand HDR	Dell EMC	Texas Advanced Computing Center 🇺🇸 United States	2019	Linux (CentOS)

The Post-PC Era


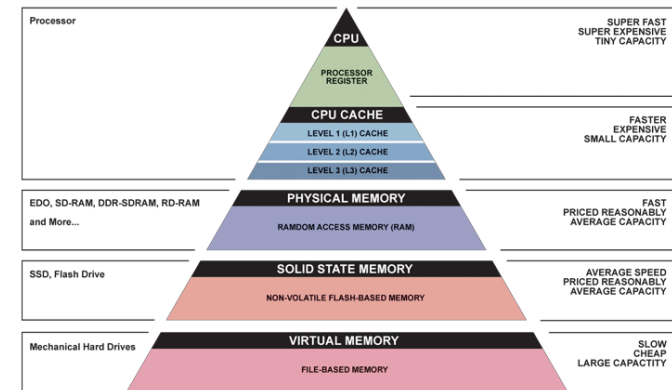
- Personal Mobile Device (PMD)
 - Battery operated
 - Connects to the Internet
 - Smart phones, Tablets, Electronic glasses
- Cloud computing
 - Warehouse Scale Computers (WSCs): 100,000 and more servers
 - Software-as-a-Service (SaaS) model
 - Portions of SW run on PMD and another portion in the Cloud
 - ~~Fog/Edge computing~~

What You Will Learn

- How does high-level language programs run on the hardware?
- What's the interface between S/W and H/W?
- What determines the performance?
- How can we improve the performance?
- What are the reasons for and the consequences of the recent switch from sequential processing to parallel processing?

Important Concepts In Computer Architecture

- Design for Moore's Law
- Use abstraction to simplify design
 - Layers of Representation
- Make the common case fast
 - Related idea: Amdahl's Law
- Performance via:
 - parallelism
 - pipelining
 - prediction
- Hierarchy of memories
- ~~Dependability via Redundancy~~




Below Your Program

- Word processor, database, ... etc.
 - Millions of lines of code



Translation of high-level operations into simple computer instructions


- Hardware can run only low-level, simple instructions
 - ADD, LOAD, STORE, JUMP ... etc
- Several layers between this gap



High-level Language → Instructions

High-level
language
program
(in C)


```
swap(int v[], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```



Compiler

Assembly
language
program
(for MIPS)

Low level programming language



Assembler

swap:

```
muli $2, $5, 4
add $2, $4, $2
lw $15, 0($2)
lw $16, 4($2)
sw $16, 0($2)
sw $15, 4($2)
jr $31
```


Binary machine
language
program
(for MIPS)

```

10001100101010000000001000110000
00000000101000010000000000011000
0000000000001100000001100000100001
10001100011000100000000000000000
100011001111001000000000000000100
10101100111100100000000000000000
101011000110001000000000000000100
0000001111100000000000000000001000
    
```


Components of a Computer

- Basic functions of a Computer
 - Input/Output data
 - Process the data
 - Store the data
- Same components for all kinds of computer
 - Independent of H/W technology
- Processor: Datapath+Control
- Input/output includes
 - User-interface devices
 - Display, keyboard, mouse
 - Storage devices
 - Hard disk, CD/DVD, flash
 - Network adapters
 - For communicating with other computers



Five classic components of a computer


Opening the Box



- I/O dominate the Apple iPad 2 tablet


- I/O: a capacitive multitouch LCD, front-/rear-facing cameras, microphone, speakers, Wi-Fi and Bluetooth networks, etc
- The datapath, control, and memory are a tiny portion

Integrated circuits (chips)



- A5 package includes two memory chips, each with 2 Gb and total 512MB

ARM processor A5



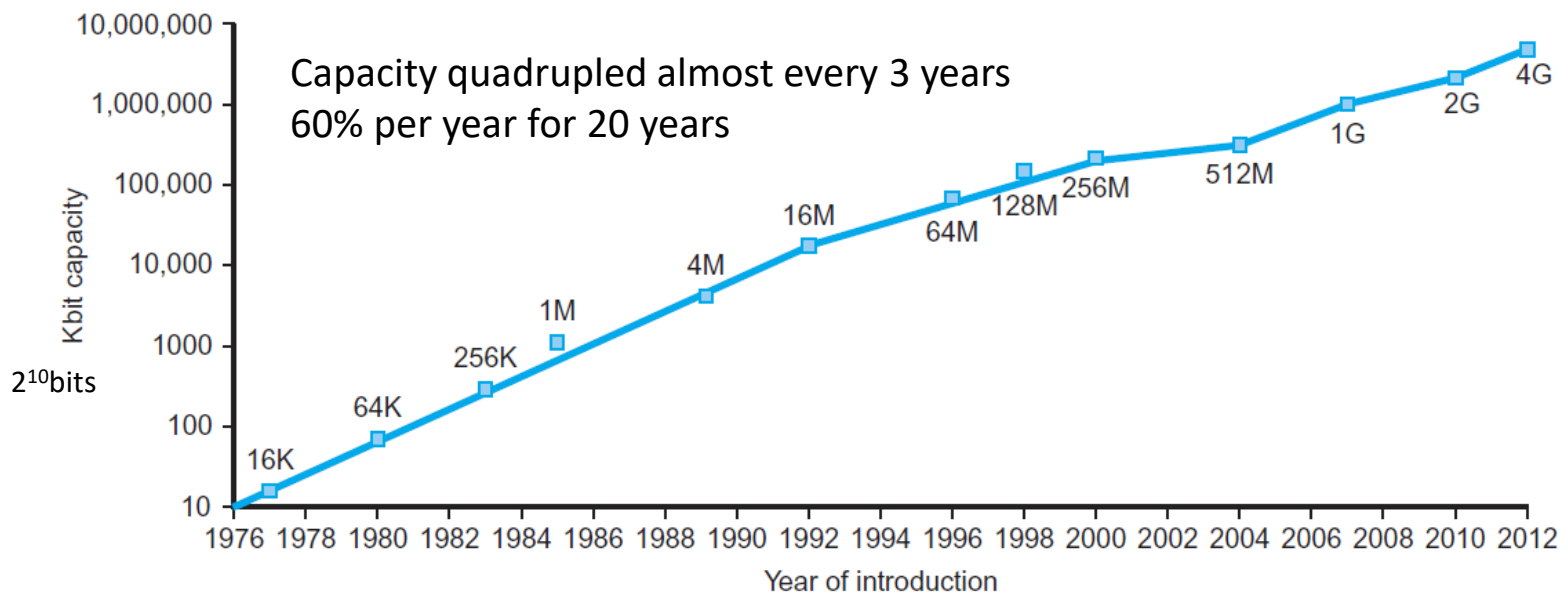
Technology Trends

- Processors and memory have improved at an incredible rate


Year	Technology used in computers	Relative performance/unit cost
1951	Vacuum tube	1
1965	Transistor	35
1975	Integrated circuit	900
1995	Very large-scale integrated circuit	2,400,000
2013	Ultra large-scale integrated circuit	250,000,000,000

* A transistor is simply an on/off switch controlled by electricity.

- Growth of DRAM Capacity



Moore's Law Revisited



Performance

- Measuring the performance is challenging!!
 - Why?
 - Are you comparing apple to orange?
 - When money is involved ...
- Two key metrics in computer performance
 - Execution time (Response time)
 - The time from the start to end
 - Throughput
 - Total amount of work done in a given time

Defining The Performance

- Focus on the response time for now
- Performance = 1/Execution Time

$$P_X > P_Y$$

$$1/E_X > 1/E_Y$$

$$E_X < E_Y$$

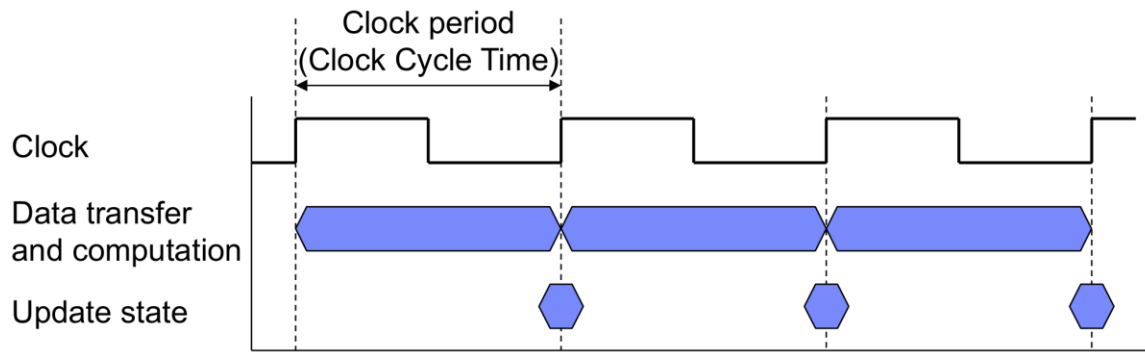
- X is n times faster than Y

$$P_X/P_Y = n$$

$$P_X/P_Y = E_Y/E_X = n$$

Measuring Time

- Wall clock time, Elapsed time
 - Total time it took to complete the job (or task)
- CPU time vs. I/O time
- Clock cycle (ticks, clock ticks, cycles ...)
 - Operation of digital hardware governed by a constant rate clock



- Clock frequency (rate): cycles per second
 - e.g., $4.0\text{GHz} = 4000\text{MHz} = 4.0 \times 10^9 \text{ Hz} = 1/(250 \times 10^{-12})$
- Clock period (Clock Cycle Time): duration of a clock cycle = $1/\text{clock frequency}$
 - e.g., $250\text{ps} = 0.25\text{ns} = 250 \times 10^{-12} \text{ s}$

Classic CPU Performance Equation

- Clock Cycle Time = $\frac{1}{\text{Clock Rate}}$
 - CPU time = **CPU clock cycles** x Clock Cycle Time
 - CPU time = **CPU clock cycles** / Clock Rate
- **IC** (Instruction count)
Number of instructions executed by the program
 - **CPI** (Clock Cycles Per Instruction)
Average clock cycles per instructions for a program (fragment)
CPI = Clock Cycles / Instruction Count
- **CPU clock cycles** = **IC** x **CPI**
 - CPU time = IC x CPI x Clock Cycle Time
 - CPU time = $\frac{IC \times CPI}{\text{Clock Rate}}$

Using CPI (Clock cycles Per Instruction)

■ Example

- Computer A has clock cycle time of 250 ps and CPI of 2.0
- Computer B has clock cycle time of 500 ps and a CPI of 1.2

Which computer is faster?

N: total # of instructions in a program

$$\text{CPU time}_A = N \times 2.0 \times 250 = 500N \text{ ps}$$

$$\text{CPU time}_B = N \times 1.2 \times 500 = 600N \text{ ps}$$

$$\frac{\text{CPU performance}_A}{\text{CPU performance}_B} = \frac{\text{Exec Time}_B}{\text{Exec Time}_A} = \frac{600N}{500N} = 1.2$$

Computer A is 1.2 times faster than Computer B

Comparing Code Segments

■ CPI Information

	CPI for each instruction class		
	A	B	C
CPI	1	2	3

■ Code segment Information

Code Sequence	Instruction counts for each instruction class		
	A	B	C
1	200	100	200
2	400	100	100

■ Which code sequence executes the most instructions?

- Segment1: $200 + 100 + 200 = 500$ instructions
- Segment2: $400 + 100 + 100 = 600$ instructions

■ Which will be faster?

- Segment1: $200 \times 1 + 100 \times 2 + 200 \times 3 = 1000$ cycles
- Segment2: $400 \times 1 + 100 \times 2 + 100 \times 3 = 900$ cycles

■ CPI

- Segment1: $1000/500 = 2.0$ CPI
- Segment2: $900/600 = 1.5$ CPI

Performance Summary


CPU time = IC x CPI x Clock Cycle Time

$$\text{CPU time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock Cycle}}$$

■ Performance difference can arise from:

- Algorithm
 - affects instruction count, and maybe CPI
- Programming Language
 - affects instruction count and CPI
- Compiler
 - affects instruction count and CPI
- ISA (Instruction Set Architecture)
 - affects instruction count, clock rate, CPI

Power Trend




$$\text{Power} \propto \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency}$$

- Faster clock rate means better performance
- Clock rate and power increase slowed
 - Too much power creates cooling problem
- Reducing required voltage allows the increase of the clock rate
 - → not able to reduce the voltage any more
- What can we do to improve the performance?

From Uniprocessor to Multiprocessor

Growth of processor performance



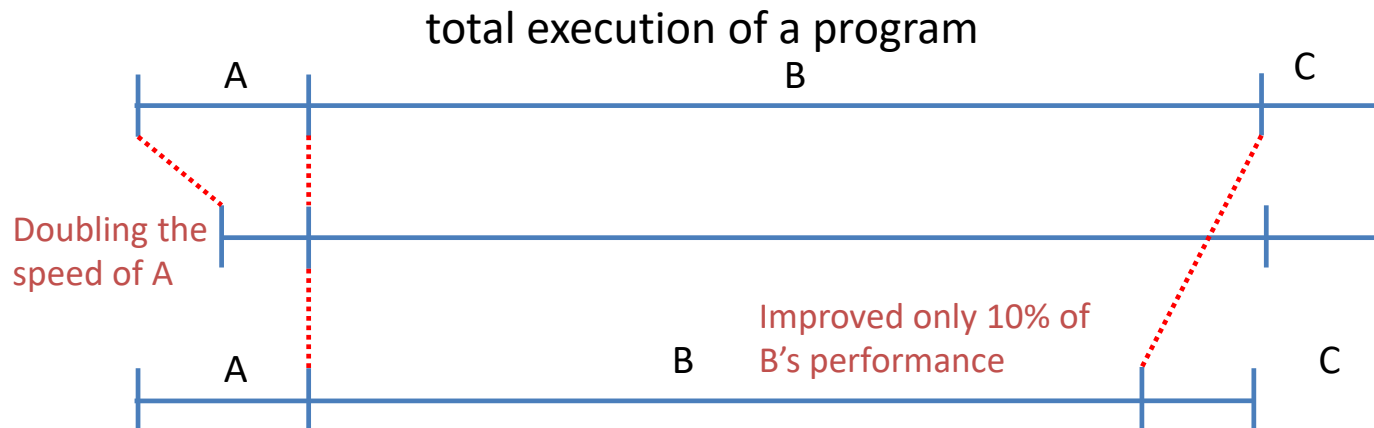
Measured by SPECint benchmark

Implication

- Previously
 - Since H/W performance increased, S/W performance *automatically* increased without any change
- Multi-core (=processor) H/W
 - dual-core, quad-core ... etc
 - S/W has to be re-written to utilize the multi-core.
 - Parallel programming is challenging
 - load balancing
 - communication overhead
 - synchronization

Amdahl's Law

■ Making the common case faster



■ Common misbelief

- Expecting the improvement of one aspect to return the same amount of increase in overall performance

■ Amdahl's Law

- The speedup of a program using multiple processors in parallel computing is limited by the time needed for the sequential fraction of the program.
- Execution time after improvement

$$= \frac{\text{Execution time affected by improvement}}{\text{Amount of improvement}} + \text{Execution time unaffected}$$

Amdahl's Law (continued)

■ Example

- A program runs in 100 seconds.
- Multiply operation consumes 80 seconds.
- We want to improve the performance 5 times.
- How much do we need to improve the performance of the multiply operation?

$$T = 80/n + (100-80)$$

5 times faster = 20 seconds

$$20 = 80/n + 20$$

$$0 = 80/n \quad \leftarrow \text{Not possible to get 5 times improvement!!!}$$

Pitfalls

- ~~Computers at low utilization use little power.~~
 - At 10% utilization, power consumption is 33% of peak power consumption.
- Designing for performance and designing for energy efficiency are ~~unrelated~~ goals.
- Using a subset of the performance equation as a performance metric.
 - MIPS (million instructions per second)

$$\begin{aligned} \text{MIPS} &= \frac{\text{Instruction count}}{\text{Execution time} \times 10^6} \\ &= \frac{\text{Instruction count}}{\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}} \times 10^6} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6} \end{aligned}$$

- Cannot compare different ISA using MIPS
- MIPS varies by program