

Chapter 2 Instructions: Language of the Computer

Instructions

■ Instruction

- Basic command describing an operation of a processor
- The interface between hardware and software

High-level
language
program
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
}
```



Assembly
language
program
(for MIPS)

```
swap:
  muli $2, $5, 4
  add $2, $4, $2
  lw $15, 0($2)
  lw $16, 4($2)
  sw $16, 0($2)
  sw $15, 4($2)
  jr $31
```

Opcode (operation code)
Operands: data for the operation

opcode	operands
--------	----------

Assembly language representation

add \$2, \$4, \$2
opcode operand

Binary language representation

00000000000110000001100000100001

Binary machine
language
program
(for MIPS)

```
0000000001010000100000000000011000
000000000000110000001100000100001
1000110001100010000000000000000000
1000110011110010000000000000000100
1010110011110010000000000000000000
1010110001100010000000000000000100
00000011111000000000000000000001000
```



Instruction Set

- Instruction Set Architecture (ISA)
 - Set of instructions that can run on a given H/W (Processor)
 - Language of a computer
 - Interface between hardware and software
- Instruction sets are similar across different processors
 - Similar hardware technology
 - Similar underlying principles
 - Common set of functionalities all computers must provide
- Popular instruction sets
 - ARMv7, ARMv8, Intel x86, Intel IA-64, IA-32, Sun SPARC, IBM POWER, DEC Alpha, AMD64 ...

Microarchitecture

- Implementation of the ISA under specific design constraints and goals
- Implementations (μ arch) can be many as long as it satisfies the specification (ISA)
 - x86 ISA has many implementations: 286, 386, 486, Pentium, Pentium Pro, Pentium 4, Core, ...
- Anything done in hardware without exposure to software
 - Pipelining, Speculative execution, Memory access scheduling,

CISC vs. RISC

- CISC (Complex Instruction Set Computer)
 - Older design idea (x86 instruction set is CISC)
 - (Multi-clock) Complex and variable length instructions
 - Many (powerful) instructions supported within the ISA
 - Ex) LOAD and STORE incorporated into instructions
 - Pros
 - Makes assembly programming much easier (lots of assembly programming in 60-70's)
 - Compiler is also simpler
 - Reduced instruction memory usage
 - Cons
 - Designing CPU is much harder
 - Approach in the early days
 - To add new instructions to meet new requirements/needs
- **New Philosophy: Let's keep instructions small and simple. Let's handle complex tasks in software.**

CISC vs. RISC

- RISC (Reduced Instruction Set Computer)
 - Newer concept than CISC
 - Simple, standardized instructions
 - ARM, MIPS, POWER, RISC-V
 - Small instruction set, CISC type operation becomes a chain of RISC operations
- Pros
 - Easier to design CPU
 - Smaller instruction set => higher clock speed
- Cons
 - Assembly language typically longer (compiler design issue)
 - Heavy use of memory
- Most modern x86 processors are implemented using RISC techniques

CISC Approach: Expecting smaller inst count, but CPI increases.

$$\text{CPU time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock Cycle}}$$

RISC Approach: Trying to reduce CPI, but inst count increases.

RISC-V

- RISC-V is an open standard ISA based on RISC principles
 - A completely open ISA that is freely available to academia and industry
 - RISC-V is simpler and more elegant than 80x86 ISA
 - Rapid adoption
- Fifth RISC ISA design developed at UC Berkeley
 - Now managed by the RISC-V foundation (riscv.org)

Arithmetic Operation

■ Notation

- **add a,b,c** $\# a = b + c$

■ RISC-V arithmetic instruction must have 3 operands. Why? Design simplicity.

■ If we want to add 4 values; b, c, d and e into a

- **add a, b, c**
- **add a, a, d**
- **add a, a, e**

→ Three instructions are required

■ Design principle

- Simplicity favors regularity.
 - Implementation becomes easier
 - Potential performance benefits

Arithmetic Example

- C code:

```
f = (g + h) - (i + j);
```

- Compiled Assembly code:

```
add t0, g, h  # temp t0 = g + h
```

```
add t1, i, j  # temp t1 = i + j
```


```
sub f, t0, t1 # f = t0 - t1
```

Register Operands

- Operands of arithmetic instructions must be registers
 - Register: storage inside CPU
- Register size of RISC-V: 64 bits (*doubleword*)
 - Number of registers: 32
- RISC-V notation convention for expressing registers
 - (x + register number): x0, x1, x2, ...

C code:

```
f = (g + h) - (i + j);
```



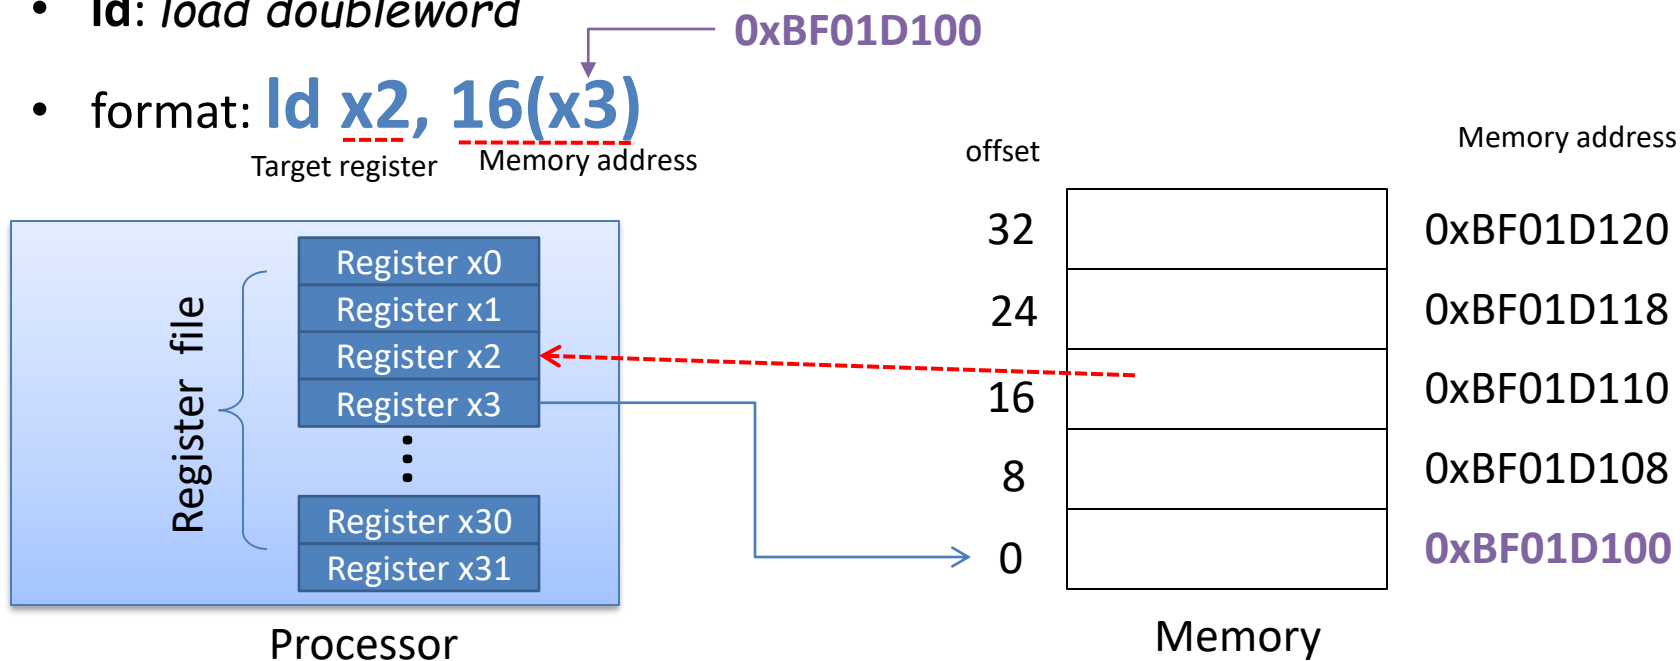
Compiled RISC-V code:

```
add x5, x20, x21 # temp x5 = g + h
add x6, x22, x23 # temp x6 = i + j
sub x19, x5, x6  # f = x5 - x6
```

f, g, h, i, j are assigned to the registers x19, x20, x21, x22, x23, respectively


Memory Operands

- Registers are limited, and data is usually larger.
 - Arithmetic must be done on registers.
 - Therefore, we need '*data transfer instructions*'
 - To/From memory and registers
- Copying data from memory to register: load
 - Address: Location of data element within a memory array
 - **ld**: *load doubleword*
 - format: **ld x2, 16(x3)**
 - Target register
 - Memory address



Endian

- Big-endian vs. little-endian
 - In Big-endian, the most significant byte of data is placed at the lowest address.
 - In Little-endian, the least significant byte of data is placed at the lowest address.
- Number: 0x1234ABCD




- RISC-V uses “little-endian”

Memory Operands

- Store instruction
 - sd: store doubleword
 - format: sd x9, 96(x22)
- Base address of the array A is in x22
- h is associated with register x21

C code:

$A[12] = h + A[8];$

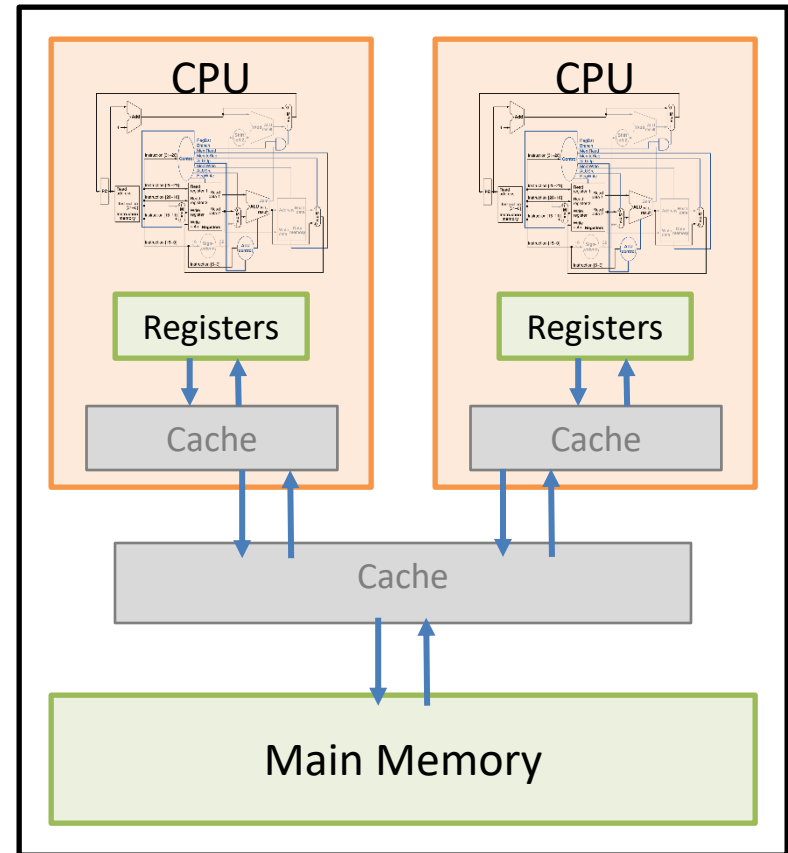


Compiled RISC-V code:

```
ld x9, 64(x22)
add x9, x21, x9
sd x9, 96(x22)
```

Register vs. Memory

- Access speed
 - Register > Memory
- Capacity
 - Register < Memory
- Energy consumption
 - Register < Memory
- Programs most likely have more variables than the number of registers.
- Frequently used variables:
 - Better to be in register than memory
 - Compilers try to optimize for this
 - *Spilling*: process of putting less commonly used variables into memory



Constant or Immediate Operands

- Arithmetic operation that uses a constant number

Ex) Adding 4 to a variable

- Option 1: Memory has 4 stored somewhere. It must be loaded first.

```
var1: .dword 4
...
ld x9, var1(x3)    # x3 + var1's offset
add x22, x22, x9
```

- Option 2: Use immediate operands

- “add immediate”, `addi`

```
addi x22, x22, 4
```

- Immediate operand instructions are faster and use less energy
- 4 is included inside the instruction bits

Constant Zero

■ Register $x0$

- Special register hard-wired to have value 0
- Cannot be overwritten (modified)

- Useful for common operations

- E.g., move data between registers

```
add x22, x21, x0          # x22 ← x21
```

- E.g., negate the value

```
sub x22, x0, x22          # x22 = -x22
```


Signed and Unsigned Numbers

- Numbers in base 2 (binary numbers)
 - ex) 123 base 10 = 1111011 base 2
- Value of i th digit d : $d \times \text{Base}^i$
- $1101_{\text{two}} = (1 \times 2^3) + (1 \times 2^2) + (0 \times 2^1) + (1 \times 2^0)$
 $= 8 + 4 + 0 + 1 = 13_{\text{ten}}$
- 1101 in a doubleword

63		55		47		39		32
0000	0000	0000	0000	0000	0000	0000	0000	
31		23		15		8		0
0000	0000	0000	0000	0000	0000	0000	1101	

- LSB (least significant bit) the right most bit, bit position 0
- MSB (most significant bit) the left most bit, bit position 63

Unsigned Numbers

- RISC-V word can represent 2^{64} bit patterns

00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000_{two} = 0_{ten}
00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000001_{two} = 1_{ten}
00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000010_{two} = 2_{ten}
... ..
11111111 11111111 11111111 11111111 11111111 11111111 11111111 11111101_{two} = 18,446,774,073,709,551,613_{ten}
11111111 11111111 11111111 11111111 11111111 11111111 11111111 11111110_{two} = 18,446,744,073,709,551,614_{ten}
11111111 11111111 11111111 11111111 11111111 11111111 11111111 11111111_{two} = 18,446,744,073,709,551,615_{ten}

- Decimal values calculated using:

$$(x_{63} \times 2^{63}) + (x_{62} \times 2^{62}) + (x_{61} \times 2^{61}) + \dots + (x_1 \times 2^1) + (x_0 \times 2^0)$$

→ representation of the unsigned numbers

- Representation limit

- Computer can express the numbers only within the given bit range
- If arithmetic operation results in the larger number than representable, **overflow** happens.

Signed Numbers

- Expressing the sign of a number
 - One bit to indicate positive or negative sign
 - Where should the sign be? How does it affect the arithmetic operation?
- Naïve solution
 - Simply use the MSB as a sign → *“sign and magnitude”* representation

$$0000_2 = 0_{10}$$

$$0001_2 = 1_{10}$$

$$0010_2 = 2_{10}$$

$$0011_2 = 3_{10}$$

$$0100_2 = 4_{10}$$

$$0101_2 = 5_{10}$$

$$0110_2 = 6_{10}$$

$$0111_2 = 7_{10}$$

$$1000_2 = -0_{10}$$

$$1001_2 = -1_{10}$$

$$1010_2 = -2_{10}$$

$$1011_2 = -3_{10}$$

$$1100_2 = -4_{10}$$

$$1101_2 = -5_{10}$$

$$1110_2 = -6_{10}$$

$$1111_2 = -7_{10}$$

$$6 - 3 = 0110 + 1011 = ???$$

Signed Numbers

■ 1's complement

- Invert positive number bits to use as a negative number

$$0000_2 = 0_{10}$$

$$0001_2 = 1_{10}$$

$$0010_2 = 2_{10}$$

$$0011_2 = 3_{10}$$

$$0100_2 = 4_{10}$$

$$0101_2 = 5_{10}$$

$$0110_2 = 6_{10}$$

$$0111_2 = 7_{10}$$

$$1111_2 = -0_{10}$$

$$1110_2 = -1_{10}$$

$$1101_2 = -2_{10}$$

$$1100_2 = -3_{10}$$

$$1011_2 = -4_{10}$$

$$1010_2 = -5_{10}$$

$$1001_2 = -6_{10}$$

$$1000_2 = -7_{10}$$

$$6 - 3 = 0110 + 1100 = 0010 = 2 \text{ ???}$$

$$6 - 3 = 0110 + 1100 + \text{carry } 1 = 0011 = 3$$

■ 2's complement

$$0000_2 = 0_{10}$$

$$0001_2 = 1_{10}$$

$$0010_2 = 2_{10}$$

$$0011_2 = 3_{10}$$

$$0100_2 = 4_{10}$$

$$0101_2 = 5_{10}$$

$$0110_2 = 6_{10}$$

$$0111_2 = 7_{10}$$

$$1111_2 = -1_{10}$$

$$1110_2 = -2_{10}$$

$$1101_2 = -3_{10}$$

$$1100_2 = -4_{10}$$

$$1011_2 = -5_{10}$$

$$1010_2 = -6_{10}$$

$$1001_2 = -7_{10}$$

$$1000_2 = -8_{10}$$

$$6 - 3 = 0110 + 1101 = 0011 = 3$$

Signed and Unsigned Numbers

■ 2's Complement

- All negative numbers have 1 in MSB
 - To find out if number is positive or negative, H/W needs to check only one bit
- Computing the decimal value

$$x = -x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

ex) 11111111 11111111 11111111 11111111 11111111 11111111 11111111 11111100

$$= (-1 \times 2^{63}) + (1 \times 2^{62}) + (1 \times 2^{61}) + \dots + (1 \times 2^2) + (0 \times 2^1) + (0 \times 2^0)$$
$$= -9,223,372,036,854,775,808 + 9,223,372,036,854,775,804 = -4$$

■ Negating 2's complement binary number

- Negate all bits and add 1

$$0000_2 = 0_{10}$$

$$0001_2 = 1_{10}$$

$$0010_2 = 2_{10}$$

$$\mathbf{0011}_2 = 3_{10}$$

$$1111_2 = -1_{10}$$

$$1110_2 = -2_{10}$$

$$\mathbf{1101}_2 = -3_{10}$$

$$1100_2 = -4_{10}$$

$$2 \rightarrow -2: 0010 \rightarrow 1101 \rightarrow \text{add } 1 \rightarrow 1110$$

$$-2 \rightarrow 2: 1110 \rightarrow 0001 \rightarrow \text{add } 1 \rightarrow 0010$$

Signed and Unsigned Numbers

- Sign extension
 - Converting a binary number in n bits to a binary number in m bits, where $m > n$.
 - Take the MSB and replicate it to fill up all the bits
 - ex) 0000 0010 = 2
 - \rightarrow 0000 0000 0000 0010 = 2
 - ex) 1111 1110 = -2
 - \rightarrow 1111 1111 1111 1110 = -2
- Signed load performs sign extension to correctly load value
 - loading 16 bit number into 32 bit register
 - instructions: **lb** (load byte), **lh** (load halfword), **lw** (load word)
 - If you don't want the sign extension, use **lbu** (load byte unsigned)

Representing Instructions

■ RISC-V **R-type** instruction format (32bit)

funct7	rs2	rs1	funct3	rd	opcode
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

- opcode: Basic operation
- rd: destination register
- funct3: additional opcode
- rs1: 1st source register
- rs2: 2nd source register
- funct7: additional opcode

Instruction	Format	funct7	rs2	rs1	funct3	rd	opcode
add (add)	R	0000000	reg	reg	000	reg	0110011
sub (sub)	R	0100000	reg	reg	000	reg	0110011

add x9, x20, x21

0	21	20	0	9	51
---	----	----	---	---	----

0000000	10101	10100	000	01001	0110011
---------	-------	-------	-----	-------	---------

Representing Instructions

■ RISC-V **I-type** instruction format

Interpreted as 2's
complement number

immediate	rs1	funct3	rd	opcode
12 bits	5 bits	3bits	5bits	7bits

Instruction encoding information

Instruction	Format	immediate	rs1	funct3	rd	opcode
addi (add immediate)	I	constant	reg	000	reg	0010011
ld (load doubleword)	I	address	reg	011	reg	0000011

ld x14, 8(x2)

000000001000	00010	010	01111	0000011
--------------	-------	-----	-------	---------

← Find the error!

Representing Instructions

■ RISC-V **S-type** instruction format

- Two registers and one immediate field
- Two immediate fields: 7bits+5bits
- rs1 and rs2 remain in the same position for all instructions

immediate	rs2	rs1	funct3	immediate	opcode
7 bits	5 bits	5 bits	3bits	5bits	7bits

immediate	rs1	funct3	rd	opcode
12 bits	5 bits	3bits	5bits	7bits

Compare with R-type

sd x14, 8(x2)

0000000		00010	011	01000	
---------	--	-------	-----	-------	--

← Fill in the blank

Instruction encoding information

Instruction	Format	immed-iate	rs2	rs1	funct3	immed-iate	opcode
sd (store doubleword)	S	address	reg	reg	011	address	0100011

Logical Operations

Logical operations	C operators	Java operators	RISC-V instructions
Shift left	<<	<<	sll, slli
Shift right	>>	>>>	srl, srli
Shift right arithmetic	>>	>>	sra, srai
Bit-by-bit AND	&	&	and, andi
Bit-by-bit OR			or, ori
Bit-by-bit XOR	^	^	xor, xori
Bit-by-bit NOT	~	~	xori

■ Shift operation

00000000 00001001 = 9

shift left 5 00000001 00100000 = ?

■ Shift instructions (I-type)

- slli, srli – shift left (right) logical immediate
- sra, srai – shift right arithmetic (immediate): fill with sign-bit

1100 0011 → 1110 0001

AND/OR Instructions

- and, andi

```
and x9, x10, x11    // reg x9 = reg x10 & reg x11
```

[illegible][illegible][illegible]

- or, ori

```
or x9, x10, x11 // reg x9 = reg x10 | reg x11
```

[illegible][illegible][illegible]

XOR Instructions

■ xor

```
xor x9, x10, x12 // reg x9 = reg x10 ^ reg x12
```

x10	00000000 00000000 00000000 00000000 00000000 00000000 00001101 11000000
x12	11111111 11111111 11111111 11111111 11111111 11111111 11111111 11111111
x9	11111111 11111111 11111111 11111111 11111111 11111111 11110010 00111111

■ not

- xor with 111...1111

Input 1	Input 2	output
0	1	1
1	1	0
0	0	0
1	0	1

Instructions for Making Decisions

- Branch instructions

- *if* statement

```
beq rs1, rs2, L1
```

- beq: "branch if equal"
- Go to the statement labeled L1 if the value in register rs1 equals the value in register rs2

```
bne rs1, rs2, L1
```

- bne: "branch if not equal"

- beq, bne: conditional branches

Compiling if-then-else into Conditional Branches

- C code segment

```
if (i==j)
    f = g + h;
else
    f = g - h;
```

- Variable to register assignments
 - f:x19, g:x20, h:x21, i:x22, j:x23

- RISC-V assembly code

Text after # is comment.

```
    bne x22, x23, Else    # go to Else if i ≠ j
    add x19, x20, x21     # f = g + h (skipped if i ≠ j)
    beq x0, x0, Exit      # if 0 == 0, go to Exit
Else: sub x19, x20, x21   # f = g - h (skipped if i = j)
Exit:
```

Compiling a while loop

- C code segment

```
While (save[i]==k)
    i += 1;
```

- Variable to register assignments

- i:x22, k:x24, base of save: x25

- RISC-V assembly code

Text after # is comment.

```
Loop: slli x10, x22, 3      # x10 = i x 8
      add x10, x10, x25    # x10 = address of save[i]
      ld x9, 0(x10)        # x9 = save[i]
      bne x9, x24, Exit    # go to Exit if save[i] ≠ k
      addi x22, x22, 1     # i = i+1
      beq x0, x0, Loop     # go to Loop
```

Exit:

Other Conditional Branches

■ blt: “branch if less than”

```
blt rs1, rs2, L1
```

- Compare the values (2's complement number) in rs1 and rs2
- Branch if rs1 is smaller

■ bge: “branch if greater than or equal”

```
bge rs1, rs2, L1
```

- Branch if rs2 is smaller

■ bltu, bgeu

- Numbers in registers are treated as unsigned

Array Index Out-of-bound check

- Array index check

```
int[] age = new int[5];  
age[-2] = 12;  
age[2] = 4;  
age[12] = 4;
```

- Need to check if $\text{index} > 0$ and $\text{index} \leq \text{max_length}$

- bgeu can be used to check both conditions

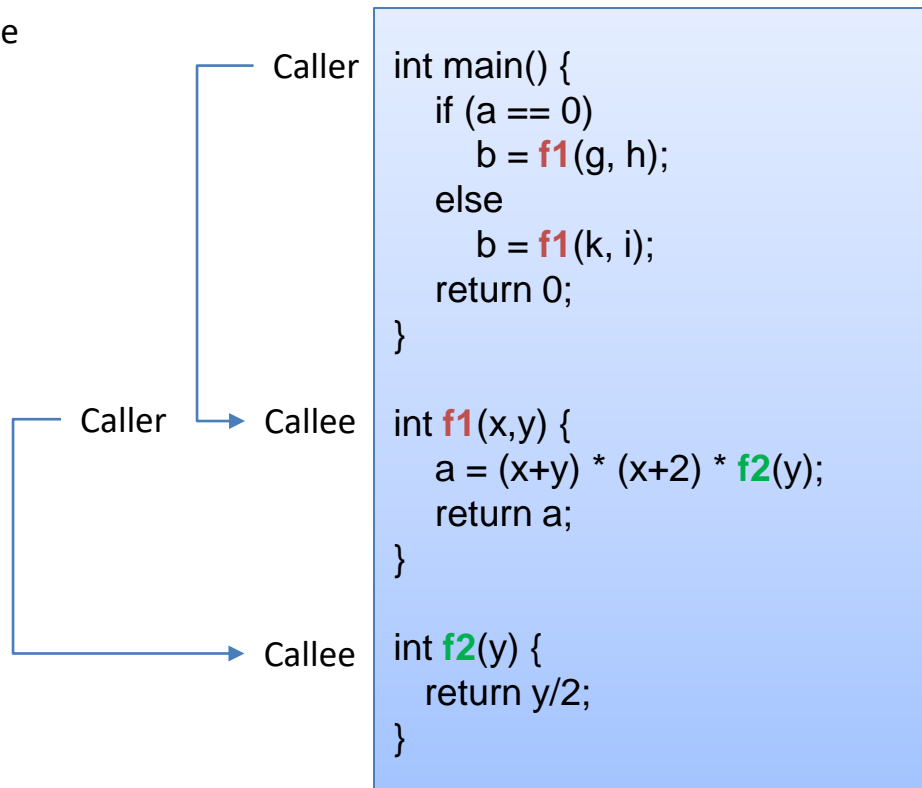
Text after # is comment.

```
# if x20 >= x11 or x20 < 0, goto IndexOutOfBounds  
bgeu x20, x11, IndexOutOfBounds
```

Supporting Procedures

■ Procedure (function) calls

- Pass function parameters
- Transfer control
 - Acquire necessary memory space
- Run procedure
- Store return value (result) to predefined location
- Return control



Supporting Procedures

- RISC-V conventions
 - x10~x17: eight parameter registers in which to pass parameters or return values
 - x1: one return address register
- RISC-V instruction for procedure handling

```
jal x1, ProcedureAddress
```

- jal: “jump-and-link”
 - Branch unconditionally and save the address of the next instruction (return address) to the designated register
- Program counter (PC) register
 - Contains the address of current instruction being executed
 - jal saves PC+4 to the rd (usually x1)
- Unconditional branch without saving the return address

```
jal x0, Label
```



Supporting Procedures

- Returning from a procedure
 - jalr: “jump and link register” instruction (I-type)

```
jalr x0, 0(x1)
```

- Branches to the address stored in register x1 (i.e., $0 + \text{address in } x1$)
- ** x0 is hard-wired to zero, writing to x0 will have the effect of discarding the value (e.g., return address)


Using More Registers

- Procedure may need more than 8 registers for execution

- Compiler chooses to use other additional registers
 - Save the value of registers to memory → spill
 - Restore the value after procedure finishes

- Ideal place for spilling registers: stack

- Last-in-first-out queue in memory
- Memory area for spilling registers
- Stack pointer, x2 (sp)
- Stack grows from high address to low address
- stack pointer is adjusted by one doubleword for each register that is saved (i.e., push) and restored (i.e., pop)



Compiling a Leaf Procedure

- Leaf procedure
 - Procedure that does not call other procedure

Text after `//` is comment.

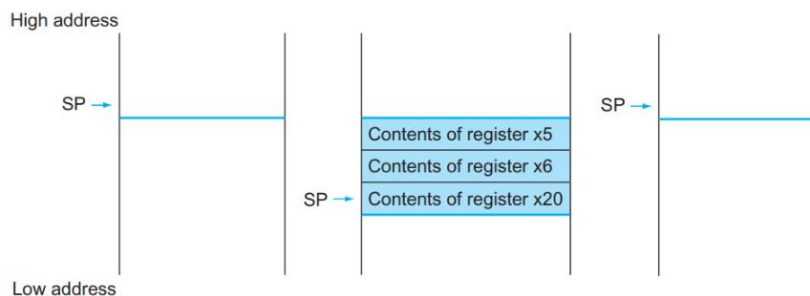
```
long long int
leaf_example(
    long long int g,
    long long int h,
    long long int i,
    long long int j)
{
    long long int f
    f = (g + h) - (i + j);
    return f;
}

// g → x10, h → x11
// i → x12, j → x13
// f → x20
```

leaf_example:

```
addi sp, sp, -24 // adjust stack to make room for 3 items
sd    x5, 16(sp) // save register x5 for use afterwards
sd    x6, 8(sp)  // save register x6 for use afterwards
sd    x20, 0(sp) // save register x20 for use afterwards
add x5, x10, x11 // register x5 contains g + h
add x6, x12, x13 // register x6 contains i + j
sub x20, x5, x6  // f = x5 - x6, which is (g + h) - (i + j)
addi x10, x20, 0 // returns f (x10 = x20 + 0)

ld x20, 0(sp) // restore register x20 for caller
ld x6, 8(sp)  // restore register x6 for caller
ld x5, 16(sp) // restore register x5 for caller
addi sp, sp, 24 // adjust stack to delete 3 items
jalr x0, 0(x1) // branch back to calling routine
```



Register Saving Convention


- Two groups of registers in RISC-V
 - No need for callee to save
 - x5-x7, x28-x31
 - registers for temporary use
 - Callee must preserve (save, use and restore) if used
 - x8,x9, x18-x27

Text after // is comment.

```
addi sp, sp, -24           // adjust stack to make room for 3 items
sd x5, 16(sp)           // save register x5 for use afterwards
sd x6, 8(sp)            // save register x6 for use afterwards
sd x20, 0(sp)              // save register x20 for use afterwards
add x5, x10, x11           // register x5 contains g + h
add x6, x12, x13           // register x6 contains i + j
sub x20, x5, x6            // f = x5 - x6, which is (g + h) - (i + j)
addi x10, x20, 0           // returns f (x10 = x20 + 0)
ld x20, 0(sp)              // restore register x20 for caller
ld x6, 8(sp)            // restore register x6 for caller
ld x5, 16(sp)          // restore register x5 for caller
addi sp, sp, 24            // adjust stack to delete 3 items
jalr x0, 0(x1)             // branch back to calling routine
```

Nested Procedure

Issues with nested procedures



Register Preservation Responsibility

Callee saves these

Caller saves these

Preserved	Not preserved
Saved registers: x8-x9, x18-x27	Temporary registers: x5-x7, x28-x31
Stack pointer register: x2(sp)	Argument/result registers: x10-x17
Frame pointer: x8(fp)	
Return address: x1(ra)	
Stack above the stack pointer	Stack below the stack pointer

Recursive Procedure

Factorial

```
long long int fact(long long int n) {
    if (n < 1) return (1);
    else return (n * fact(n - 1));
}
```

Preserved	Not preserved
Saved registers: x8-x9, x18-x27	Temporary registers: x5-x7, x28-x31
Stack pointer register: x2(sp)	Argument/result registers: x10-x17
Frame pointer: x8(fp)	
Return address: x1(ra)	
Stack above the stack pointer	Stack below the stack pointer

to the caller

fact:

```
addi sp, sp, -16 // adjust stack for 2 items
sd x1, 8(sp) // save the return address
sd x10, 0(sp) // save the argument n
addi x5, x10, -1 // x5 = n - 1
bge x5, x0, L1 // if(n-1) >= 0, goto L1
addi x10, x0, 1 // return 1
addi sp, sp, 16 // pop 2 items off stack
jalr x0, 0(x1) // return to caller
```

L1:

```
addi x10, x10, -1 // n>=1: argument gets (n-1)
jal x1, fact // call fact with (n-1)
addi x6, x10, 0 // return from jal: move result of fact(n-1) to x6:
ld x10, 0(sp) // restore argument n
ld x1, 8(sp) // restore the return address
addi sp, sp, 16 // adjust stack pointer to pop 2 items
mul x10, x10, x6 // return n*fact(n-1)
jalr x0, 0(x1) // return to the caller
```

where to?

Text after // is comment.

Procedure's Local Data on the Stack

- Local variables within the procedure

- Placed into stack

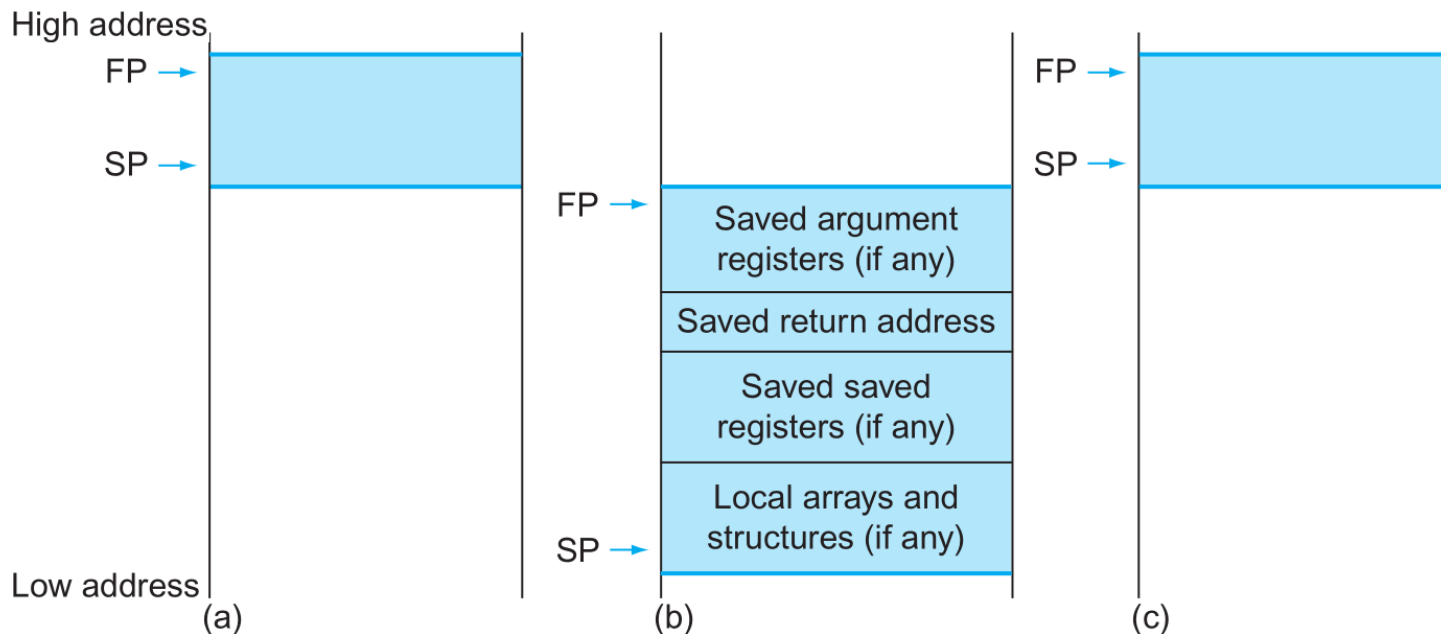
- Procedure frame (= activation record)

- Segment of stack that contains *saved registers, return address and local variables*

- Frame pointer, FP (x8) – point to the first word of the procedure frame

- SP may change during execution

```
int fact (int n) {
    int m = 3;
    if (n<1) return (1);
    else return (m*n*fact(n-1));
}
```



Allocating Memory on the Heap

■ Address space convention

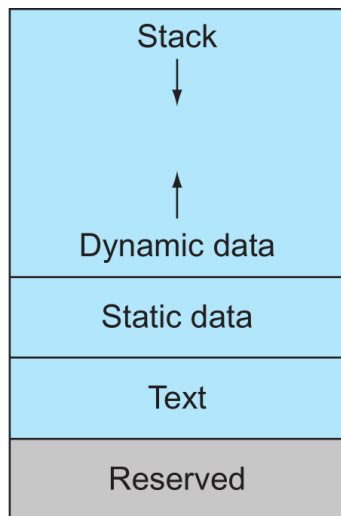
- Address space: range of addresses that can be used/allocated/referred
 - It has nothing to do with the real physical memory

SP → 0000 003f ffff fff0_{hex}

0000 0000 1000 0000_{hex}

PC → 0000 0000 0040 0000_{hex}

0



■ Text segment

- text = instructions, program binaries

■ Static data segment

- constants and static variables


■ Heap (for dynamic data)

- Place where dynamically allocated memory goes to
 - linked list, tree ...
- Grows upward
- malloc() and free() in C

Wide Immediate Operands

- There are times when constants are too big to fit into 12 bits
- RISC-V provides **lui** (*Load Upper Immediate*) instruction

- U-type** instruction



- Load 20 bit constant into bit position 31~12 from a register
 - Left 32 bits are extended with 31th bit value
 - Rightmost 12 bits are filled with zeros
- How do we make this value in a register?

```
00000000 00000000 00000000 00000000 00000000 00111101 00000101 00000000
```

- `lui x19, 976 // 976 = 0000 0000 0011 1101 0000`


```
00000000 00000000 00000000 00000000 00000000 00111101 00000000 00000000
```

- `addi x19, x19, 1280 // 1280 = 0101 00000000`

```
00000000 00000000 00000000 00000000 00000000 00111101 00000101 00000000
```


Branch Instruction Format

■ SB-type




- Can represent branch addresses from -4096 to 4094 in multiple of 2 (i.e. even address)

bne x10, x11, 2000 // if x10 != x11, go to location 2000 ten = 0111 1101 0000 **0**



- imm[0] is assumed to be always 0, and thus not encoded
- immediate field is 13 bits

■ UJ-type



- For unconditional branch instruction **jal** (the only instruction in this type)
- imm[20:1] : 20-bit address immediate (signed number)
- jal uses 21-bit address
- imm[0] is assumed to be always zero

PC-relative Addressing


- Implication of immediate field size
 - Conditional branch: 13 bits
 - Unconditional branch: 21 bits
 - Program size = $2^{\text{immediate_field_size}}$
- new PC \leftarrow **register** + branch offset
 - Allow program to be 2^{64} and use branches
- PC-relative addressing enables:
 - Branch within $\pm 2^{10}$ instructions
 - Jumps within $\pm 2^{18}$ instructions
- Branching far away

beq x10, x0, L1



bne x10, x0, L2
jal x0, L1


L2:




RISC-V Addressing Modes

■ Addressing mode: How to locate data needed as operands

1. Immediate addressing




2. Register addressing



Registers

Register

3. Base addressing



Memory

Register




Byte

Halfword

Word

Doubleword

4. PC-relative addressing



Memory

PC



Word

Jumping Very Far Away

- RISC-V allows very long jumps to any 32-bit address
 - lui write bits 12~31 of the address to a temporary register
 - jalr adds the lower 12 bits of the address to the register and jumps to the sum

```
beq x10, x0, L1
```



```
lui x9, L1_U20  
addi x9,x9, L1_L12  
bne x10,x0, L2  
jalr x0, 0(x9)
```

```
L2:
```