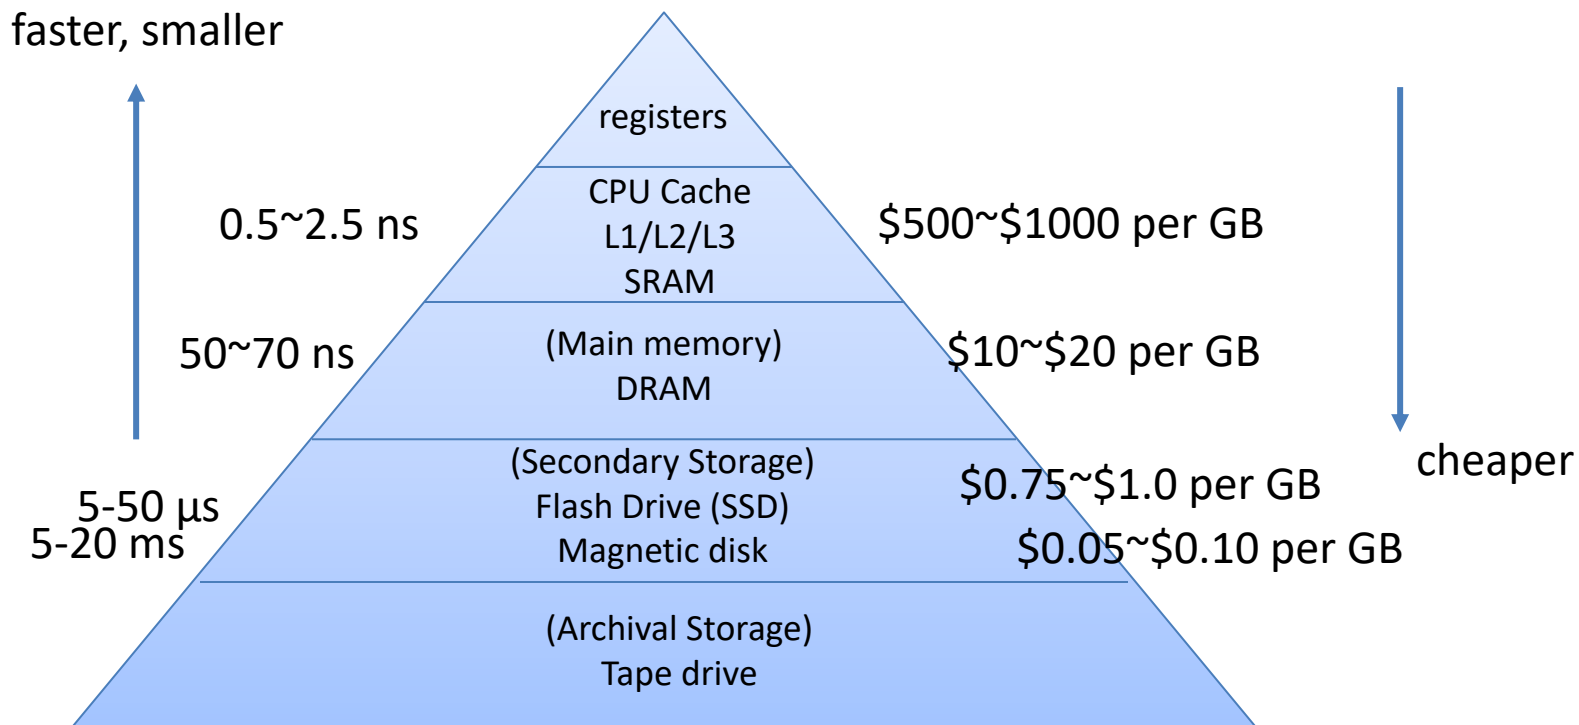


Memory Hierarchy

Memory Hierarchy




- Multiple levels of memory with different speeds and sizes
 - memory: any medium that can store data (e.g., DRAM, disk ...)



Cache Memory Concept

■ Cache memory


- Faster but smaller memory that holds data temporarily to improve the access time of a slower memory
- Transparent



Cache Memory

- Underlying principle of Cache
 - Principle of Locality
 - Programs tend to access the same data in a near future
 - Temporal locality
 - If data is accessed, it is likely that it will be accessed again soon
 - Spatial locality
 - If data is accessed, it is likely that near-by data will also be accessed soon
- Cache memory is smaller than the lower level memory
 - Select only a small set of data to store in the cache

Cache Terminology



- Block (or line)
 - Minimum unit of information
- Hit
 - When data is present in the upper level
- Hit rate (hit ratio)
 - Fraction of memory access that is 'hit'
- Miss, miss rate
- Hit time
 - Time to access the upper level (including the time to determine hit or miss)
- Miss penalty
 - Time to copy a block from the lower level
 - Replace and transfer data

Cache Access

- Assume that:
 - CPU requests one word at a time
 - Block size is one word
- Memory access: X_1, \dots, X_{n-1}, X_n

X_4
X_1
X_{n-2}
X_{n-1}
X_2
X_3

a. Before the reference to X_n

X_4
X_1
X_{n-2}
X_{n-1}
X_2
X_n
X_3

b. After the reference to X_n

- Cache placement issue

Direct-mapped Cache

■ In direct-mapped cache:

- Cache location is determined as a function of the memory address
- A block (word) can go into only one place in the cache

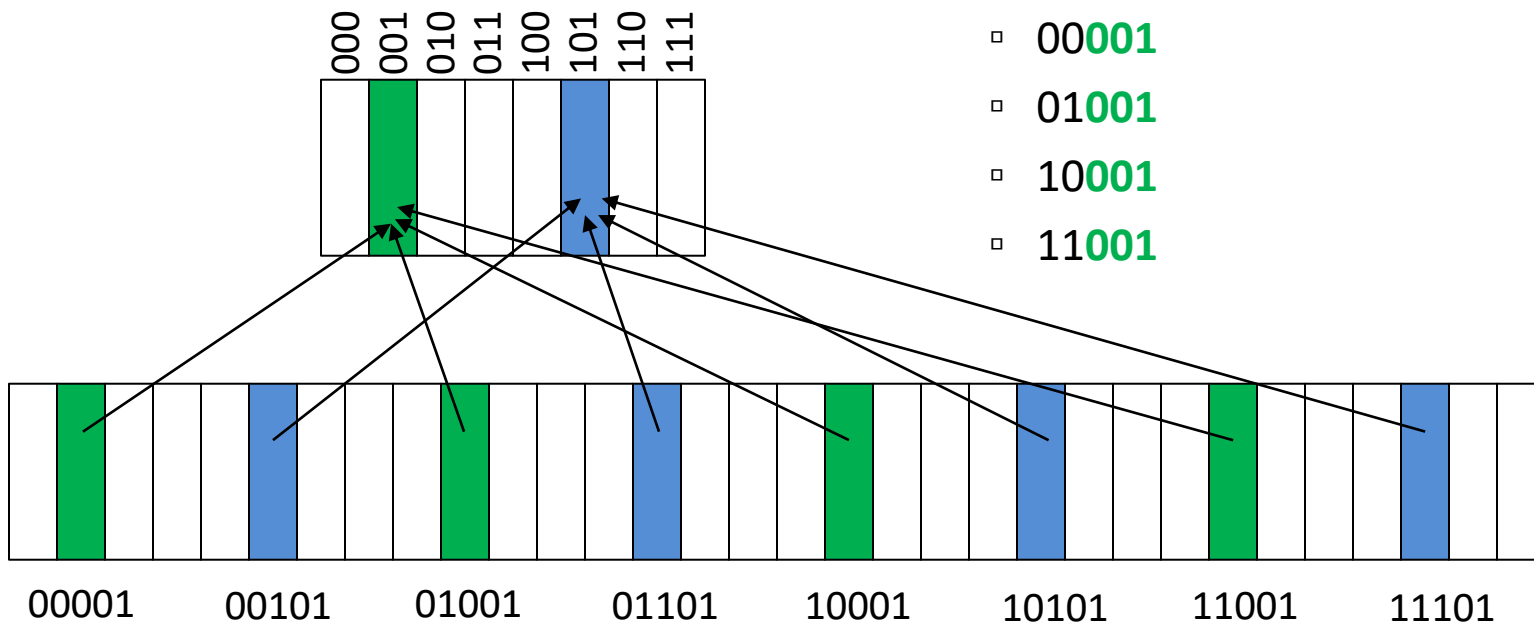
↔ fully associative cache

- Cache slot = (word address) mod (# of cache blocks)

■ # of cache blocks are in power of 2

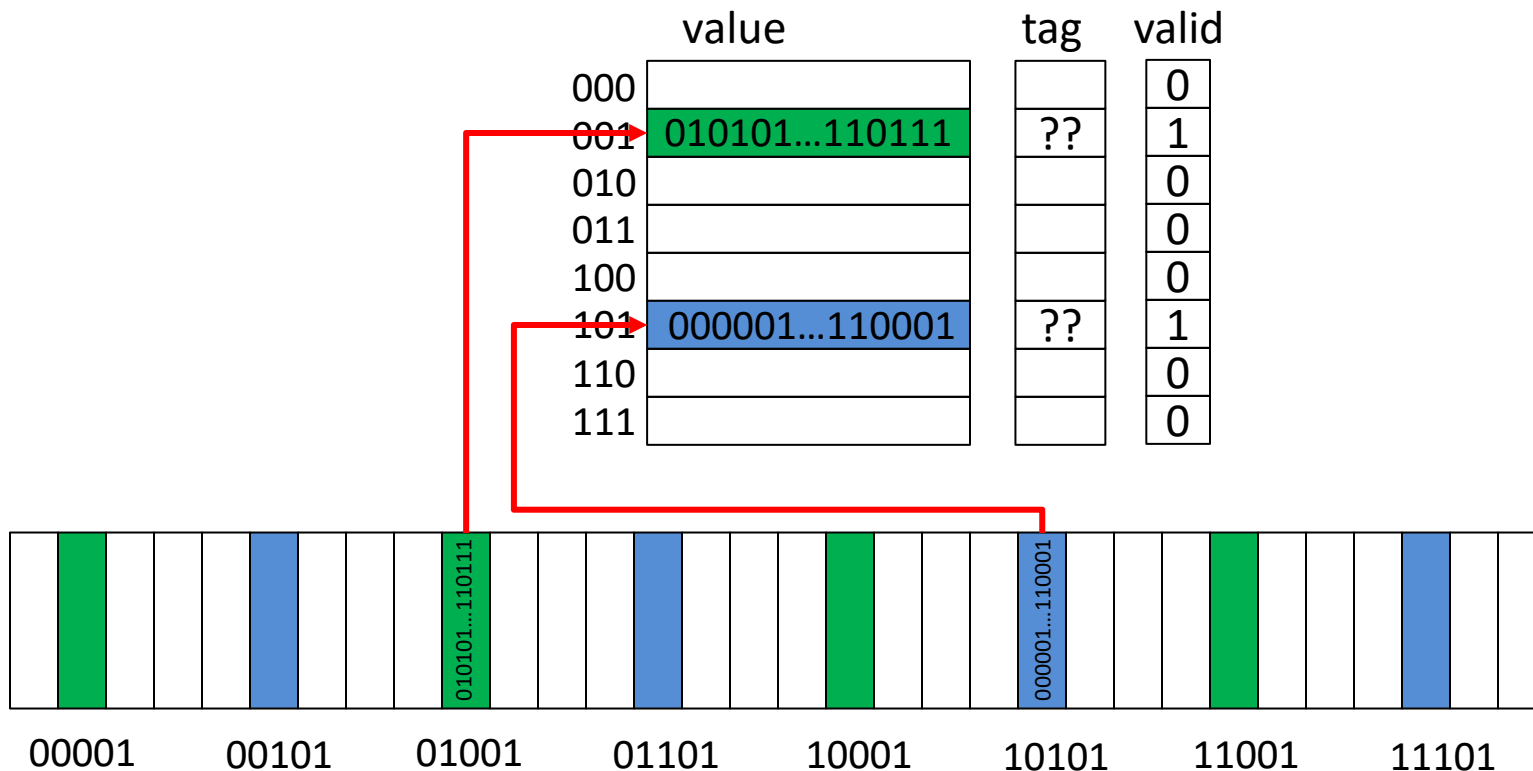
- Blocks that go into location **001**

- 00**001**
- 01**001**
- 10**001**
- 11**001**



Direct-mapped Cache

- Multiple block maps to one location in the cache
 - How do we know which address the block came from?
 - Block in 001 comes from 00001, 01001, 10001, 11001
- Tags
 - High-order bits of the address – **00**001, **01**001, **10**001, **11**001



Accessing a Cache

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

Index	V	Tag	Data
000	N		
001	N		
010	Y	11_{two}	Memory (11010_{two})
011	N		
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

Index	V	Tag	Data
000	Y	10_{two}	Memory (10000_{two})
001	N		
010	Y	11_{two}	Memory (11010_{two})
011	N		
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110_{two}	miss (5.9b)	$(10110_{\text{two}} \bmod 8) = 110_{\text{two}}$
26	11010_{two}	miss (5.9c)	$(11010_{\text{two}} \bmod 8) = 010_{\text{two}}$
22	10110_{two}	hit	$(10110_{\text{two}} \bmod 8) = 110_{\text{two}}$
26	11010_{two}	hit	$(11010_{\text{two}} \bmod 8) = 010_{\text{two}}$
16	10000_{two}	miss (5.9d)	$(10000_{\text{two}} \bmod 8) = 000_{\text{two}}$
3	00011_{two}	miss (5.9e)	$(00011_{\text{two}} \bmod 8) = 011_{\text{two}}$
16	10000_{two}	hit	$(10000_{\text{two}} \bmod 8) = 000_{\text{two}}$
18	10010_{two}	miss (5.9f)	$(10010_{\text{two}} \bmod 8) = 010_{\text{two}}$
16	10000_{two}	hit	$(10000_{\text{two}} \bmod 8) = 000_{\text{two}}$

Index	V	Tag	Data
000	Y	10_{two}	Memory (10000_{two})
001	N		
010	Y	11_{two}	Memory (11010_{two})
011	Y	00_{two}	Memory (00011_{two})
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

Index	V	Tag	Data
000	Y	10_{two}	Memory (10000_{two})
001	N		
010	Y	10_{two}	Memory (10010_{two})
011	Y	00_{two}	Memory (00011_{two})
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

Address Subdivision

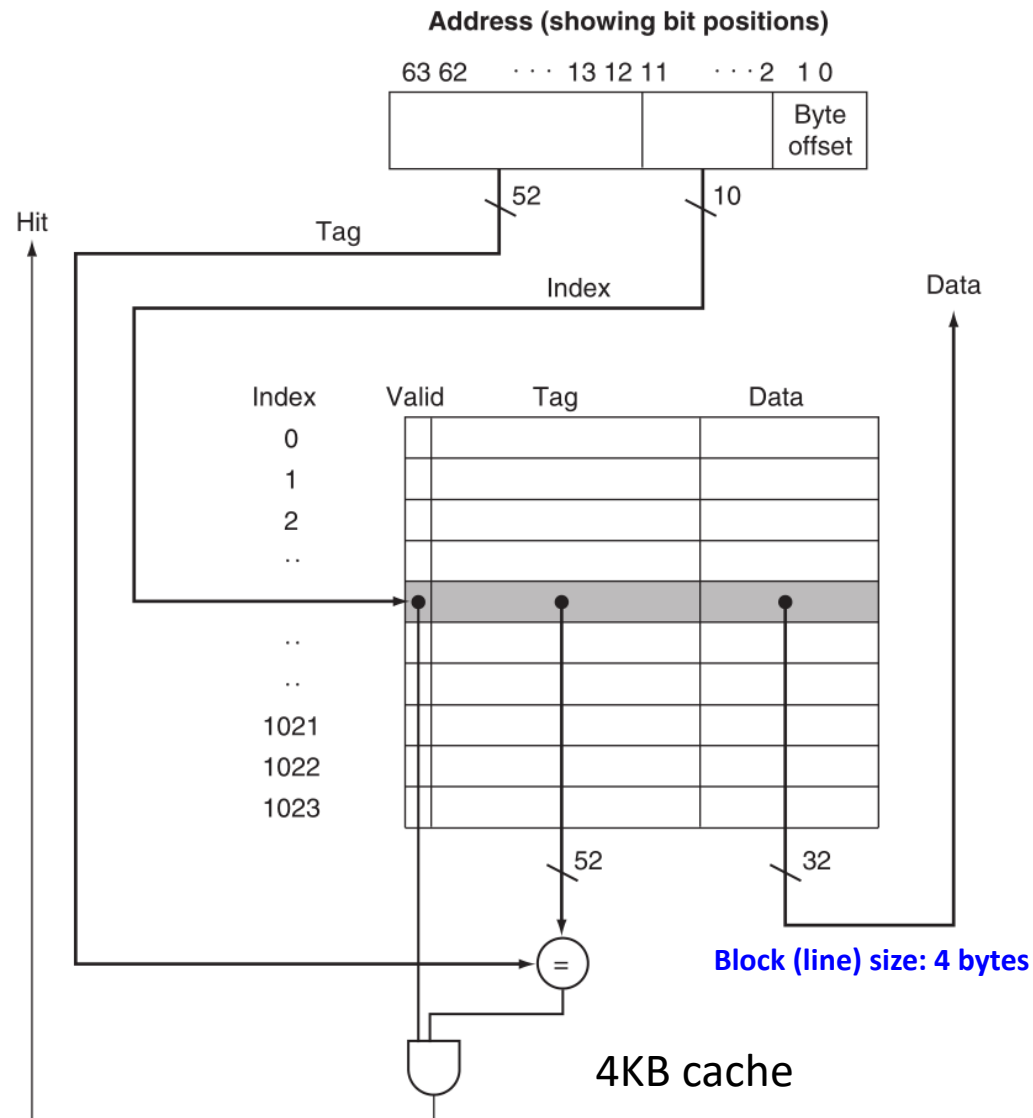
- Address is divided into tag part and index part

- tag part: 52 bits
- index part: 10 bits
- byte offset: 2 bits

- Total number of bits in the cache
 n: number of index bits
 m: 2^m is the # of words
 (m=0 if 1 word in a line)


$$2^n \times \sum \begin{array}{l} \text{valid bit: 1} \\ \text{tag bits: } 64 - (n + m + 2) \\ \text{data bits: } 2^m \times 32 \end{array}$$

*Cache size usually includes only the data size



Block Size and Miss Rate

■ Larger block size



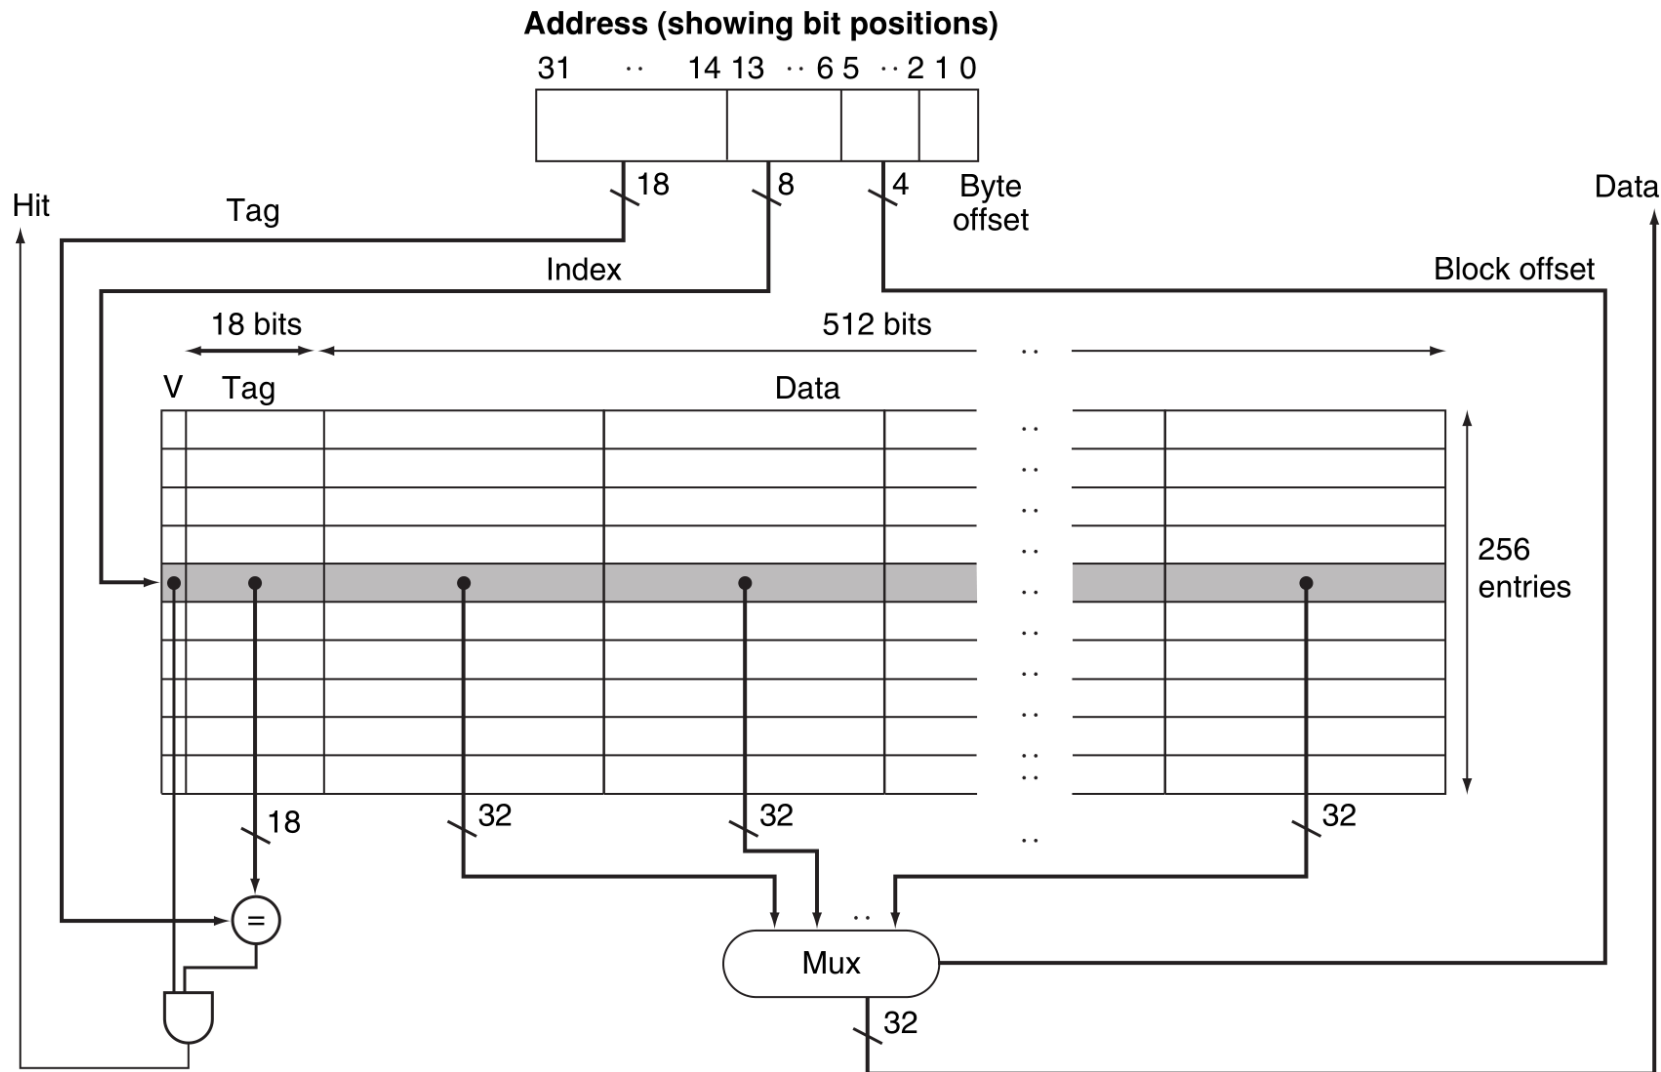
* What about larger cache size?

■ Miss penalty

- Time to load a block from MM(Main memory) to Cache
 - Set-up time + transfer time
- Reducing the miss penalty (by hiding the transfer time)
 - “early-restart” → more effective for instruction than data
 - “critical word first”

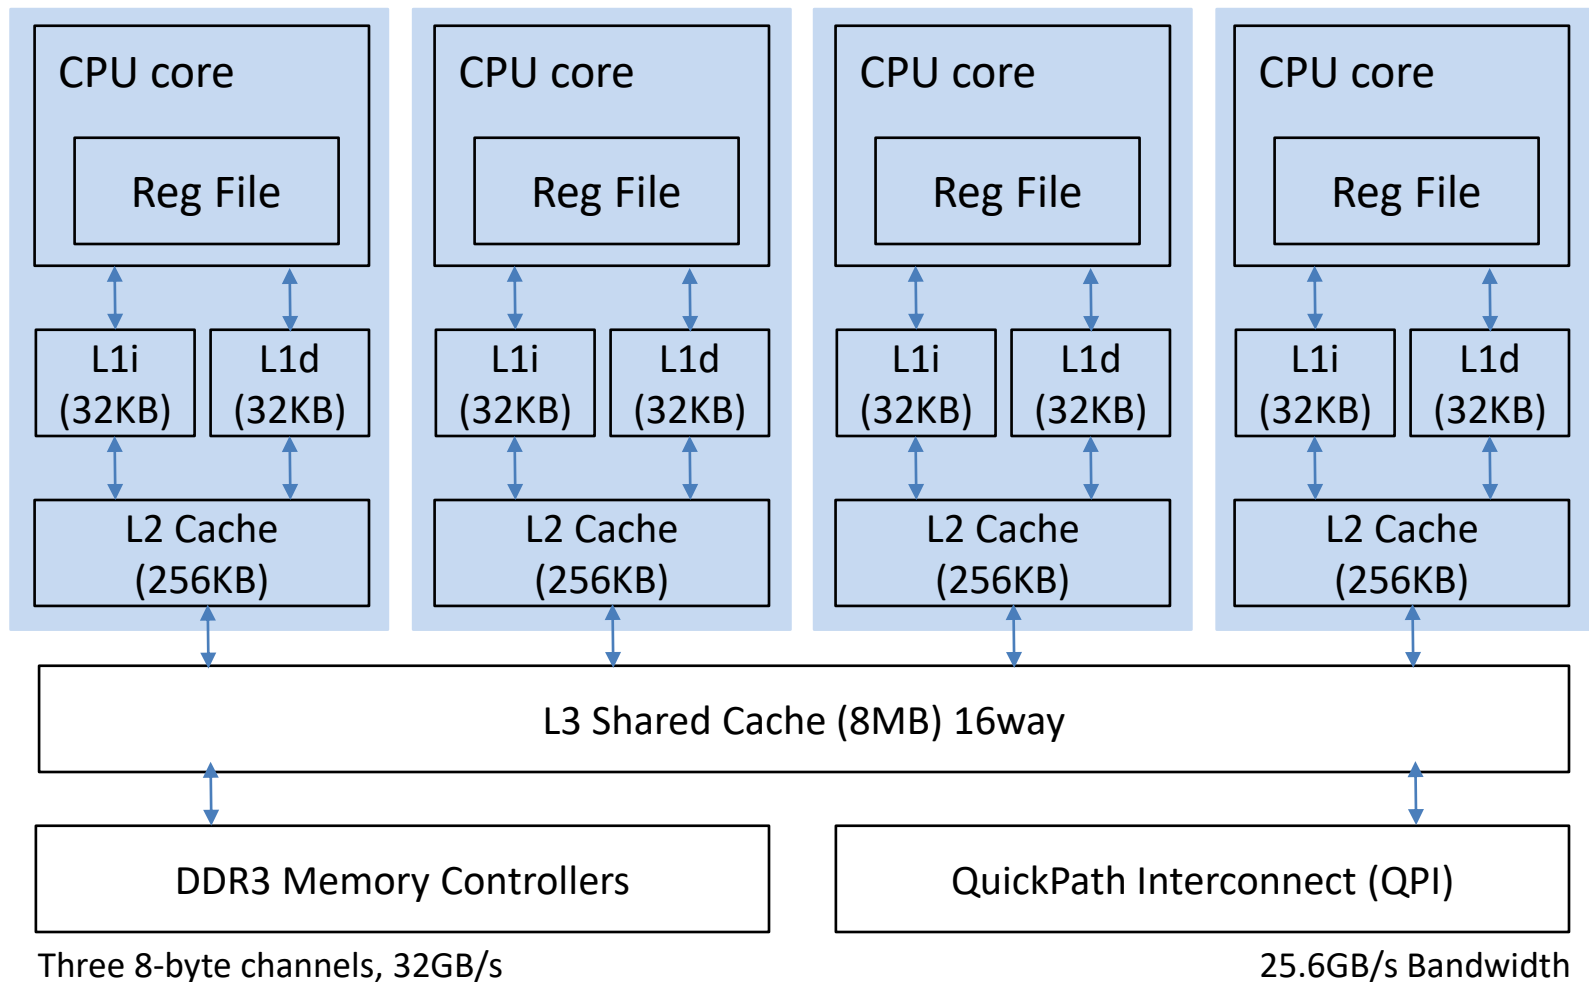

Example Cache

■ Intrinsity FastMATH embedded microprocessor



Intel core i7 Block Diagram

Four x86 SMT(Simultaneous Multithreading) Processors
Dedicated L1, L2 Cache
Shared L3 Cache



Cache Miss Handling in Processor

- Separate cache for instruction and data stream

- Instruction and data cache at L1

Processors	MIPS32 74K Atheros AR9344	Intel i7-4770	Apple A10	ARM Cortex-A57
Instruction cache	32KB, 32B line, 4-way	32KB, 64B line, 8-way	64KB, 64B line, 4-way	48KB, 64B line, 3-way
Data cache	32KB, 32B line, 4-way	32KB, 64B line, 8-way	64KB, 64B line, 4-way	48KB, 64B line, 2-way

- Cache miss on instruction (identical for data miss)

- Send PC-4 to memory unit
- Instruct MM (main memory) to perform read
- Wait until data is ready
- Write data to cache entry, fill up the tag and valid fields
- Re-send instruction address to the memory – hit!


- ★(advanced) Improving the cache performance

- In In-order execution, pipeline needs to stall on cache miss!
- Out-of-order execution: non-blocking cache
- Victim cache – small (fully associative) cache holding the replace blocks

Write Handling in Cache


Write-through cache

- Write operation completes after updating the main memory and the cache
- Data consistency maintained



Write-back cache

- Write operation completes after only updating the cache
- Main memory is synced only when a block is replaced from cache



Cache Performance

■ CPU time

- Normal execution: includes cache hit time
- Memory stall cycle: cache miss penalty

$$\begin{aligned} CPUtime &= total\ clock\ cycles \times clock\ cycle\ time \\ &= (CPU\ cycles + \textbf{Memory stall cycles}) \times clock\ cycle\ time \end{aligned}$$

■ Assume the read and write miss penalties are the same

- Write-through: complication from *write buffer*
- Write-back: potential stall from sync when a block is replaced

$$\textbf{Memory stall cycles} = Read\ stall\ cycles + Write\ stall\ cycles$$

$$= inst\ count \times \frac{memory\ access}{inst\ count} \times miss\ rate \times miss\ penalty$$

Cache Performance

- Instruction cache miss rate: 2%
- Data cache miss rate: 4%
- Ideal case CPI: 2
- Question: How much faster is the processor with perfect cache?
 - Number of instructions: N
 - Instruction miss and data miss are separate
 - total penalty cycles = inst miss cycles + data miss cycles
 - inst miss cycles = $N \times 2\% \times 100 = 2N$
 - data miss cycles = $N \times 36\% \times 4\% \times 100 = 1.44N$
 - Total cycles = $2N + 2N + 1.44N = 5.44N$
- Total cycles for perfect cache = $2N$
- Performance is better by $5.44N/2N = 2.72$


계산 전반적으로 재검토

Data miss시 cache hit 사이클도 추가해야함.

Associative Cache

- Direct-mapped cache: a block can be placed in one location
- Fully associative cache: a block can be placed anywhere in the cache
- Set associative cache: more than one location, but not anywhere (e.g., n-way set associative cache)


Direct mapped



Set associative



Fully associative



Associative Cache

- Direct-mapped cache and fully associative cache are *special cases* of set associative cache
 - direct-mapped: one-way set associative
 - fully associative: n-way set associative

- Advantage of associative cache

- Decreased miss rate

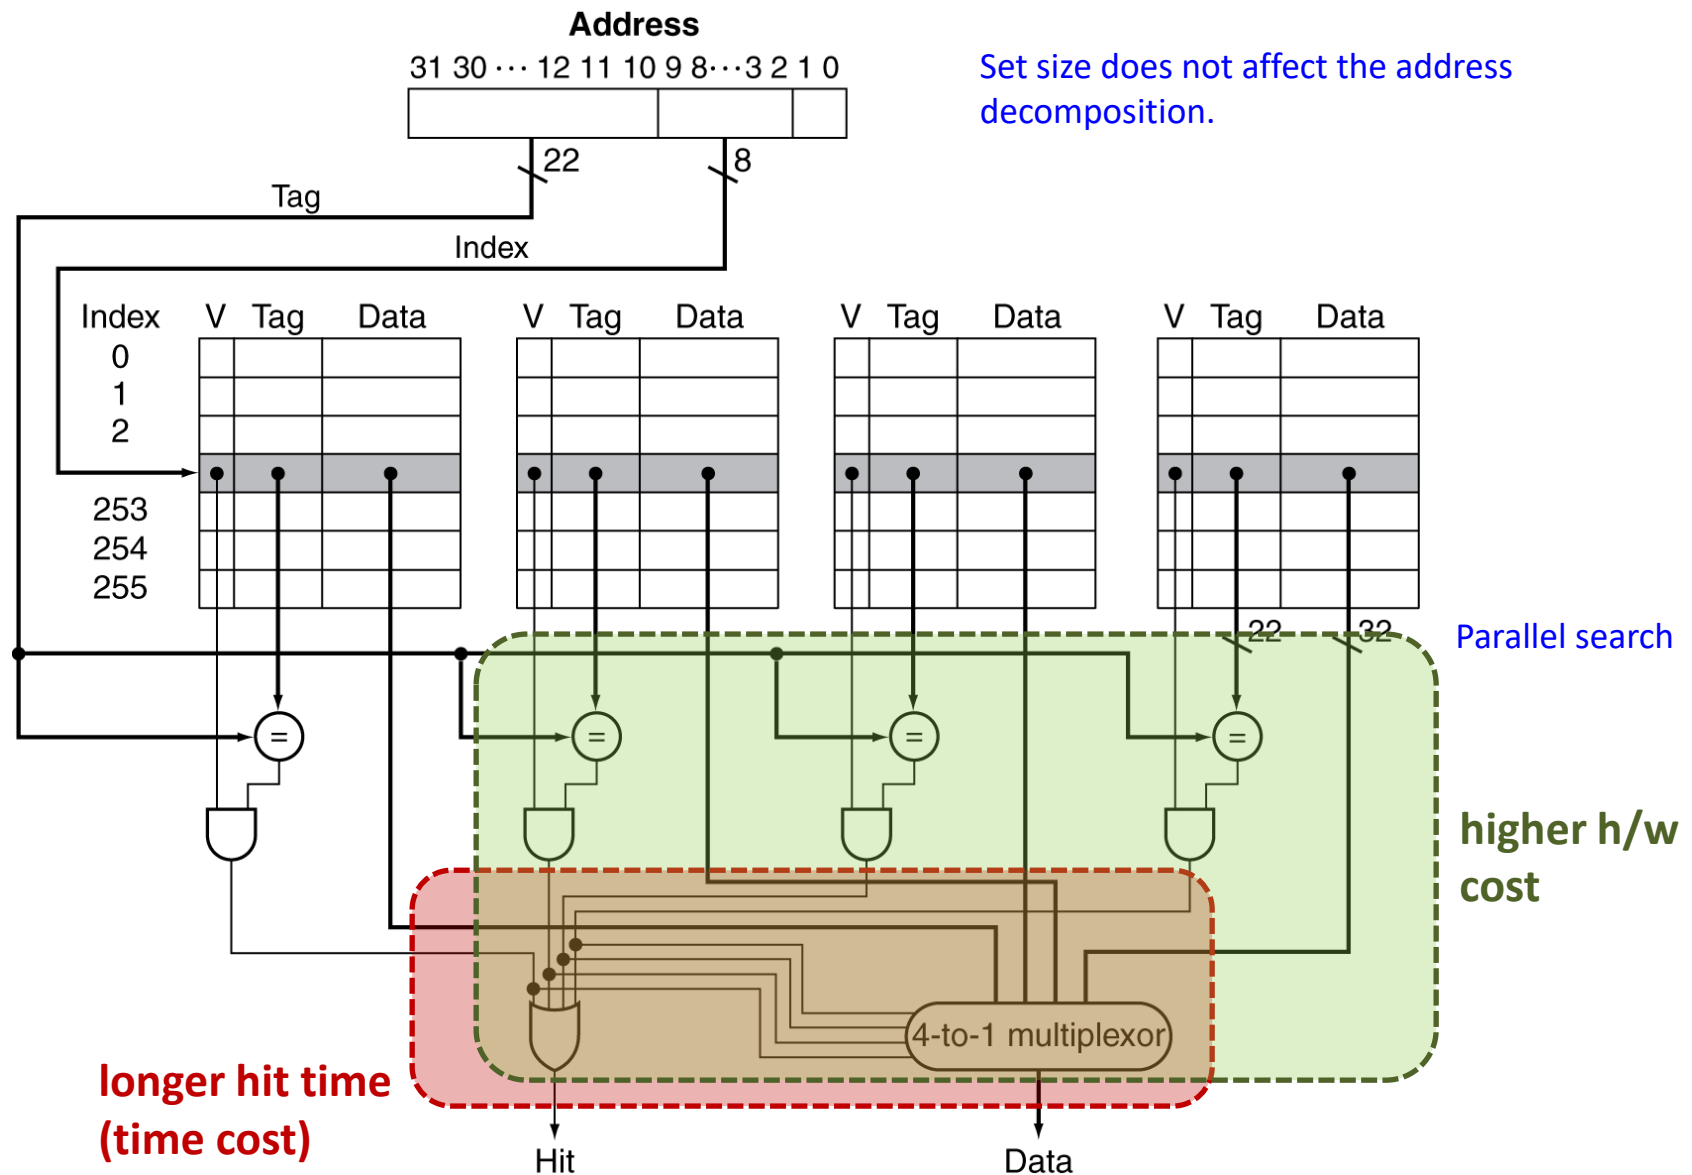
Associativity	Data miss rate
1	10.3%
2	8.6%
4	8.3%
8	8.1%

64KB data cache
16-word line
Intrinsity FastMATH
processor for SPEC
CPU2000 benchmarks

- Disadvantage

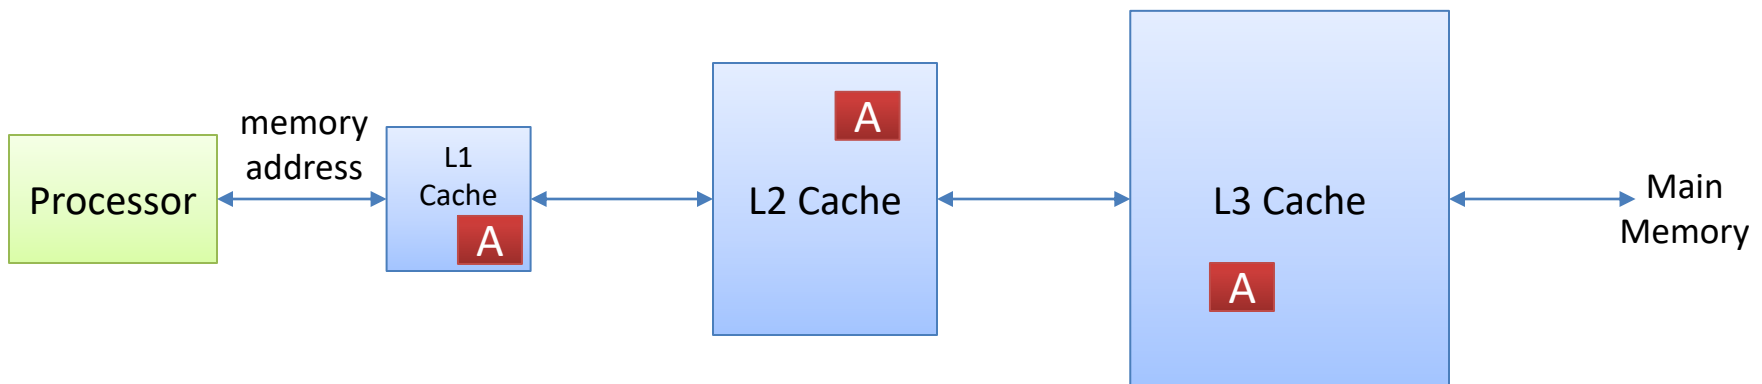
- Longer hit time
 - All tags within the set must be compared and selected
 - Higher H/W cost

4-way Set Associative Cache Schematic




Multi-level Cache

- Modern processors have more than 1 level of cache
 - L1, L2, shared L3 cache on chip
- L1: small, fast
- L2: larger, slower
 - higher associativity, larger block size
 - Access time is less critical than L1
- L1 cache miss → served by L2 cache
 - L2 cache miss → served by L3 cache



Multilevel Cache Performance


■ Performance with L1 Cache only



Base CPI	1.0
Clock rate	4 Ghz
L1 Miss rate	2%
Main memory access latency	100 ns

- i) Avg CPI = 1 cycle x 1.0 + 400 cycles x 0.02 = 9
- ii) Avg CPI = 1 cycle x 0.98 + (400+1) cycles x 0.02 = 0.98 + 8.02 = 9

■ Performance with L2 cache




- i) Avg CPI = 1 cycle x 1.0 + 20 cycles x 0.02 + 400 cycles x 0.005 = 3.4
- ii) Avg CPI = 1 cycle x 0.98 + (20+1) cycles x 0.015 + (400+20+1) cycles x 0.005 = 3.4

■ Performance comparison

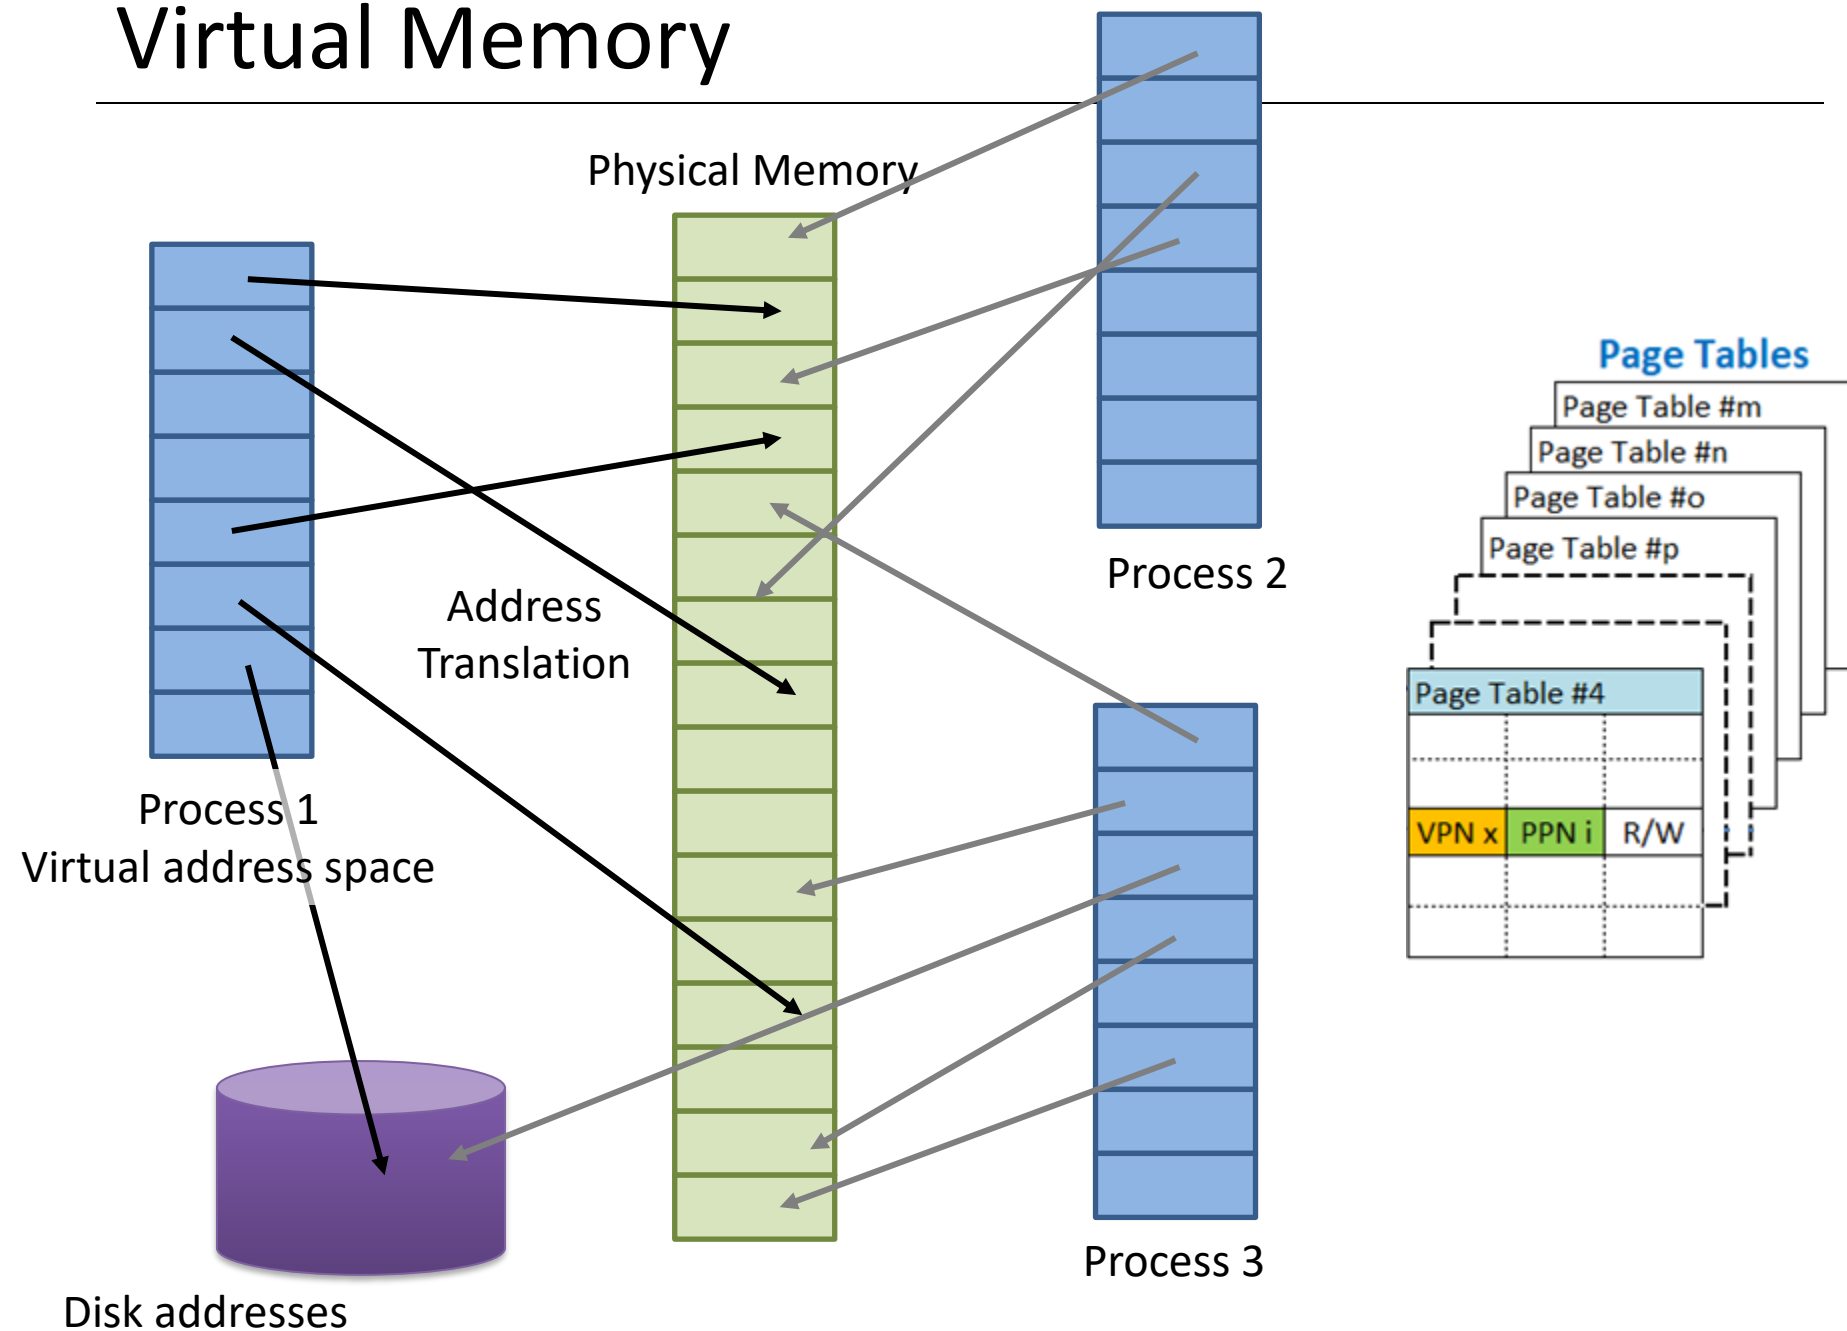
- $\frac{9}{3.4} = 2.6$

Virtual Memory

- Virtual address space
 - Each process has 4GB address space (in 32 bit architecture)
 - Broken up into 4KB pages
 - Only small subset of these pages are in active use.

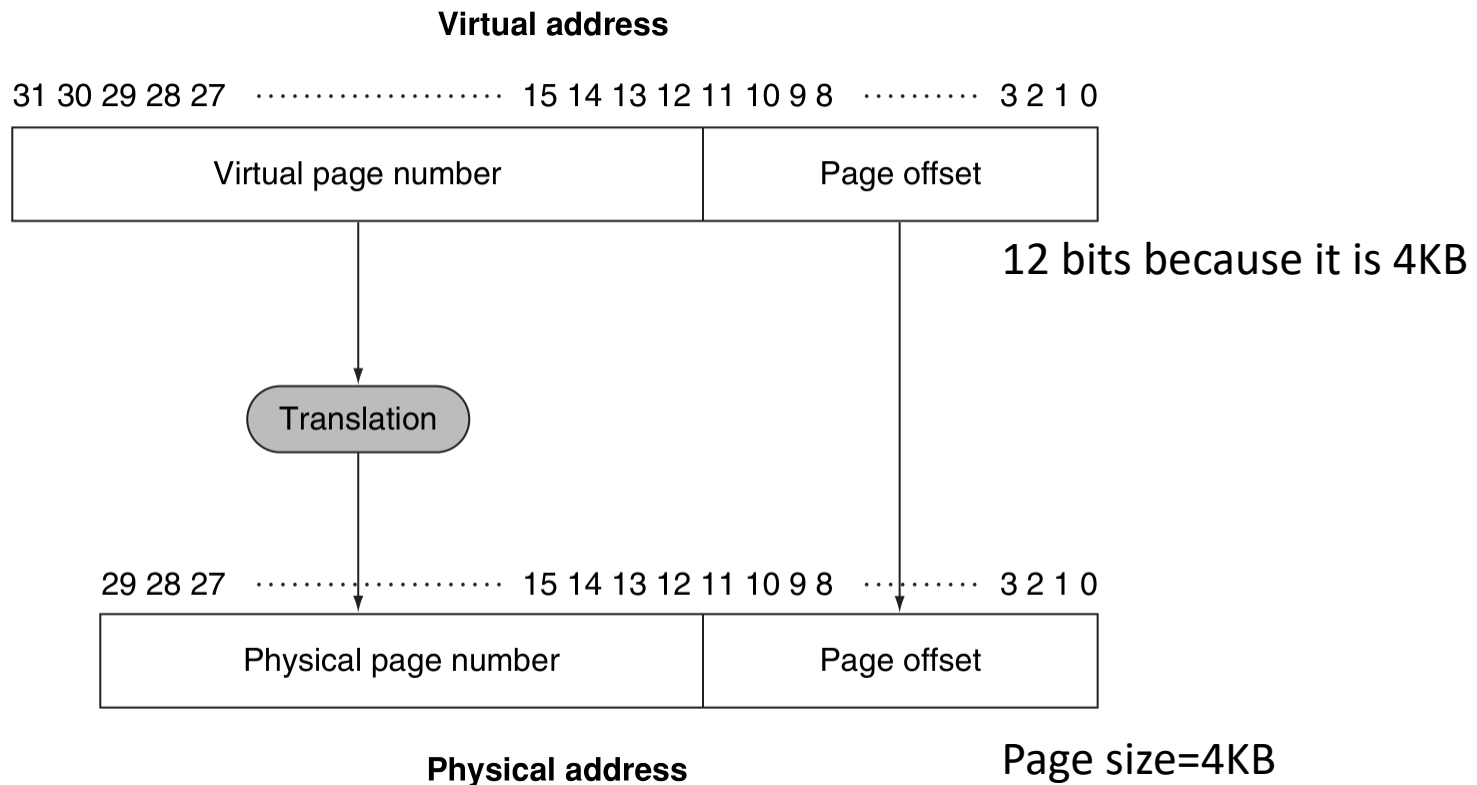


Virtual Memory



Virtual to Physical Address Mapping

- Virtual address consists of:
 - Virtual page number, Page offset

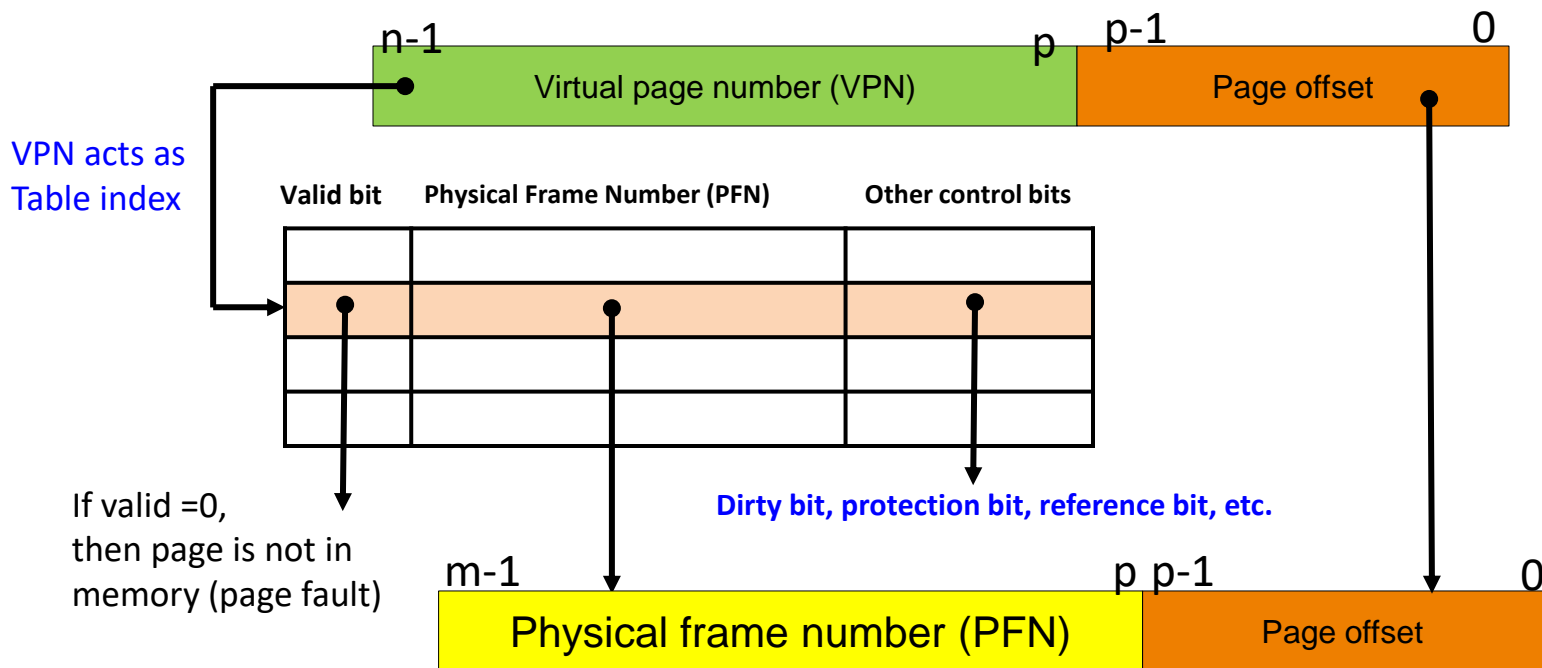


Number of physical pages (frames) = $2^{18} = 256K$

Total Physical Mem: $256 \times 4KB = 1GB$

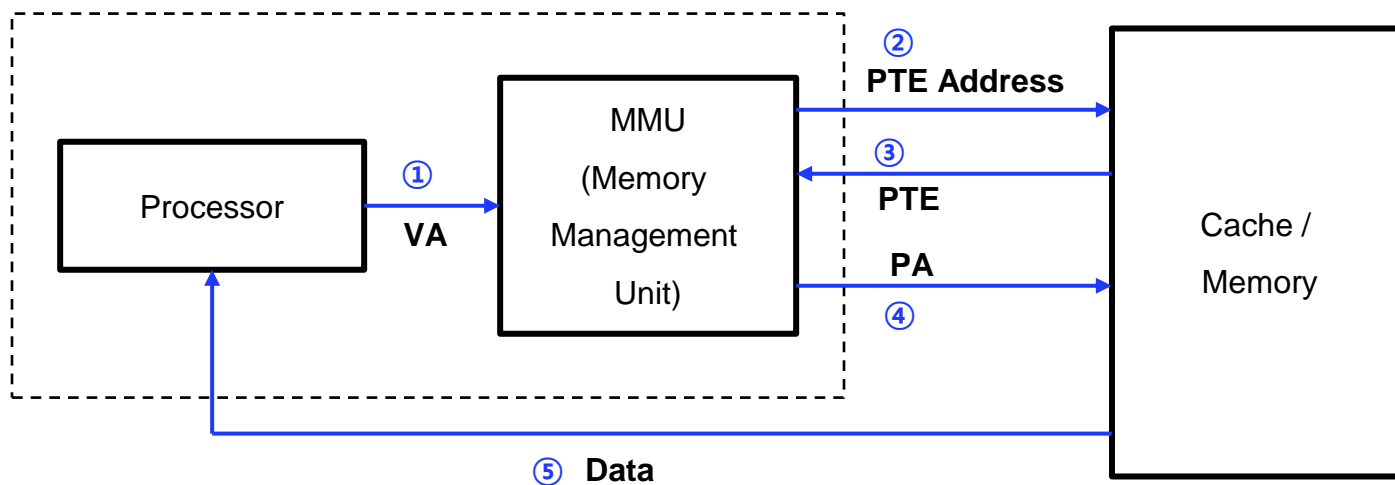
Address Translation

- How to obtain the physical frame number (PFN) for the virtual page number (VPN)?
 - Page Table** contains the mapping information in each entry called **Page Table Entry (PTE)**
 - VPN is used as an index for the PTE in a page table



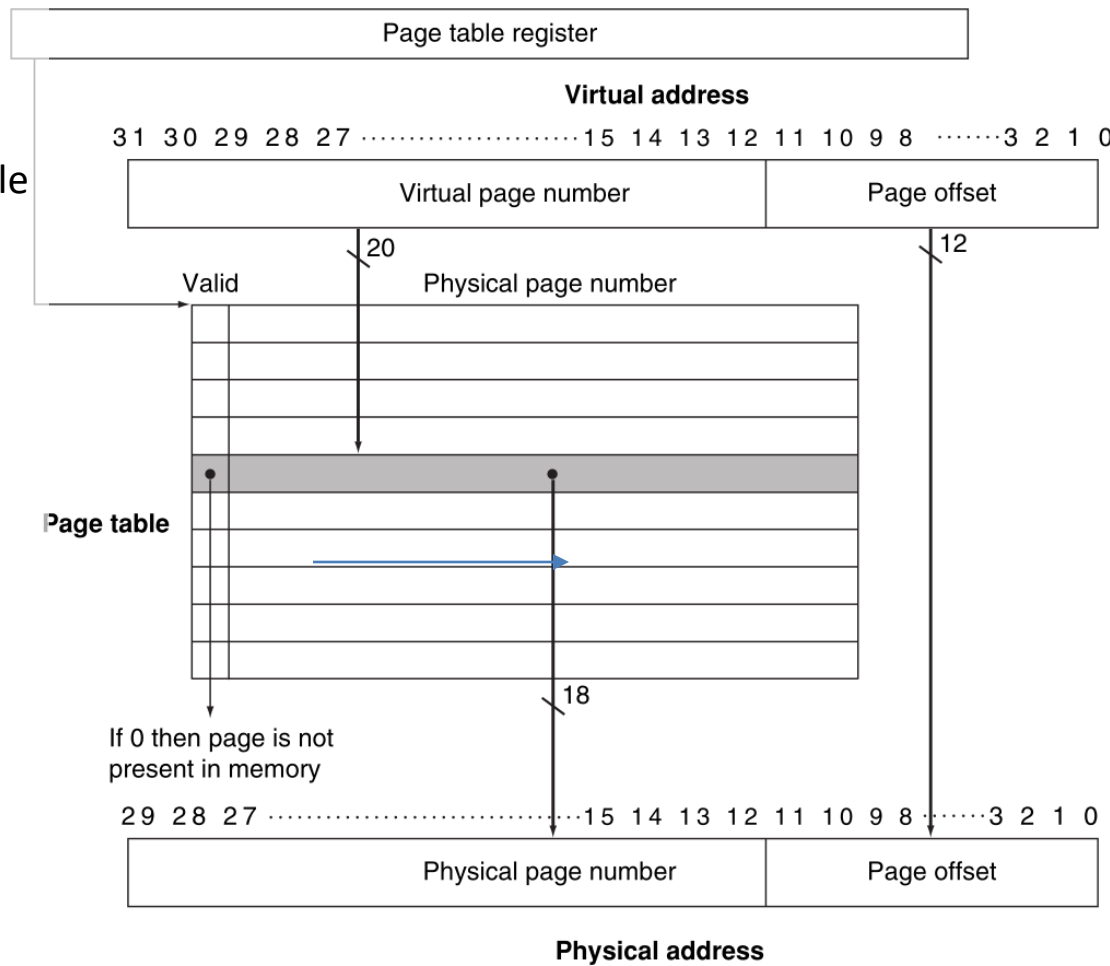
Address Translation in hardware

- When page is in the memory : **Page Hit**
 - 1) Processor sends virtual address to MMU (Memory Management Unit)
 - 2,3) MMU fetches PTE from page table in memory
 - 4) MMU sends physical address to L1 cache
 - 5) L1 cache sends a data word to the processor



Page Table Look-up

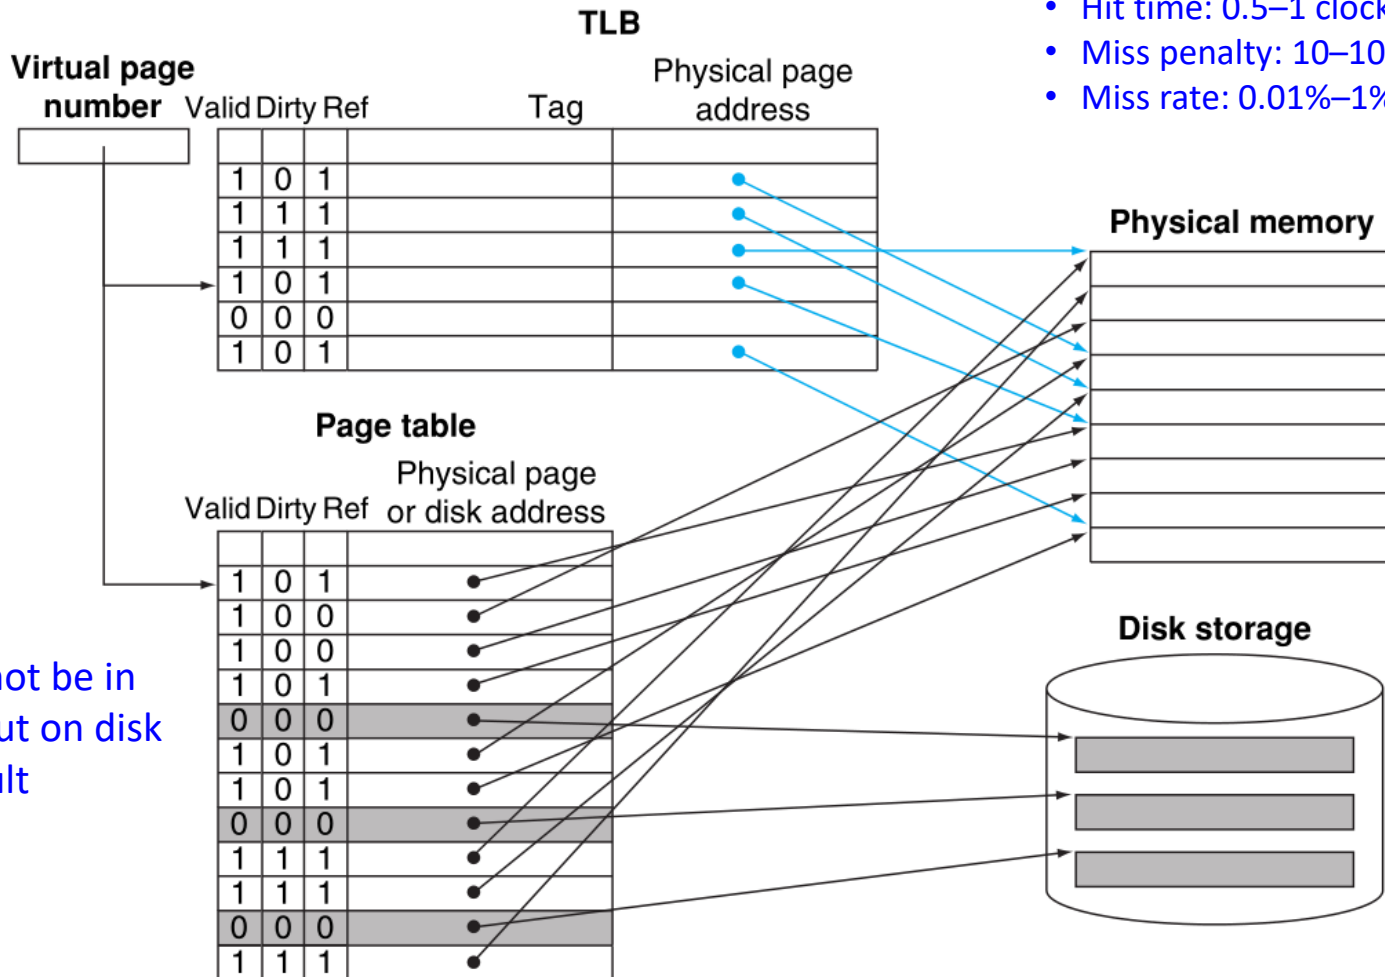
- Page table look-up is costly
 - Page table is stored in memory
 - Each memory access requires page table look-up.
 - One page table look-up may require more than one memory access
 - Leverage the reference locality!!
 - Translation of page will probably be needed again.
- Translation-lookaside buffer (TLB)
 - Special address translation cache
 - Contains subset of v-to-p page mappings



TLB (Table Lookaside Buffer)

- Since TLB is a cache, it has the tag field.
- If no entry is found in TLB, page table is searched.

- TLB size: 16–512 entries
- Block size: 1–2 page table entries (typically 4–8 bytes each)
- Hit time: 0.5–1 clock cycle
- Miss penalty: 10–100 clock cycles
- Miss rate: 0.01%–1%

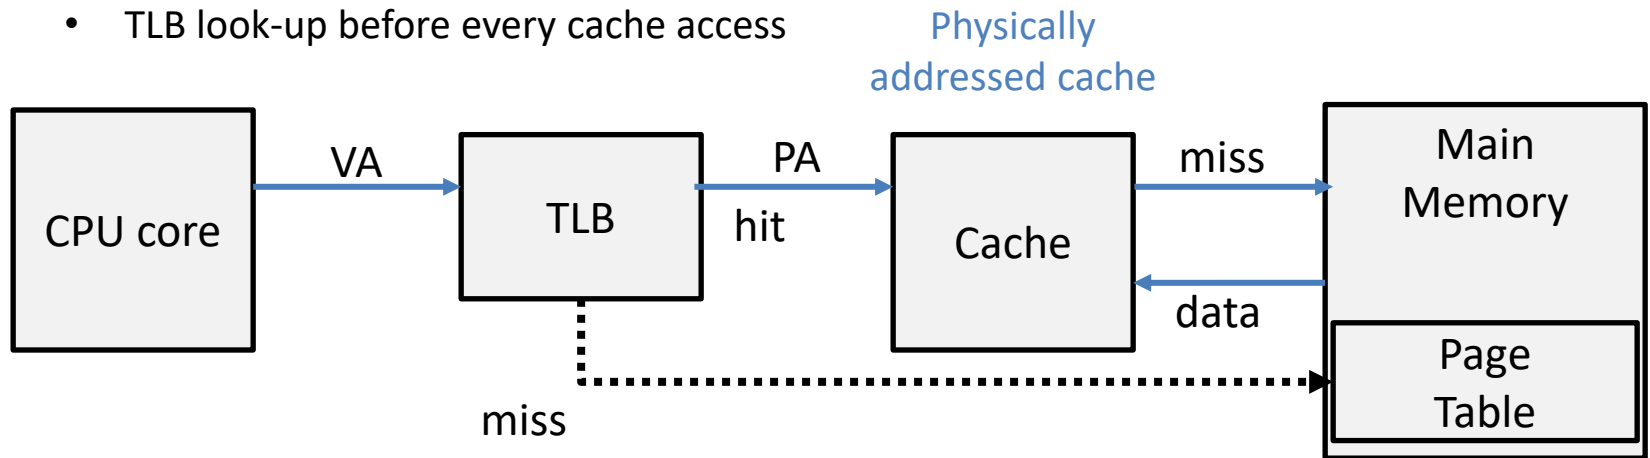


Page may not be in memory, but on disk
→ Page fault

TLB, CPU and Memory


■ Physically addressed cache

- TLB look-up before every cache access



■ Virtually addressed cache

- Cache must be flushed frequently



TLB, CPU and Memory

- Virtually-indexed physically tagged cache (VIPT)
 - Cache index bit should not overlap with tag part
 - TLB look-up and cache index search runs in parallel
 - If cache becomes large, aliasing problem arises.
 - Most processors use VIPT for L1 cache

